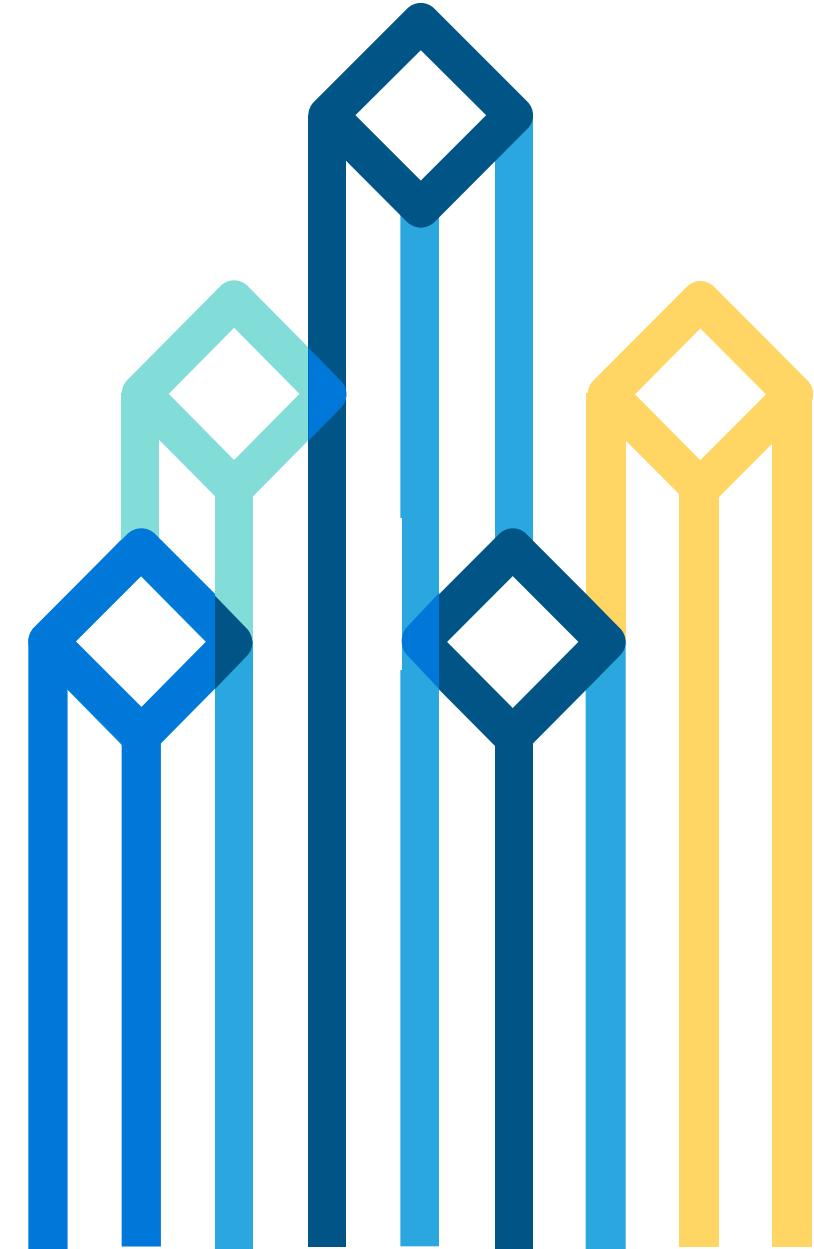




Apache HBase: Overview, Hands-On, and Use Cases

Apekshit Sharma
Dima Spivak



Who we are

Apekshit Sharma

- Distributed Software Engineer, Cloudera
 - Software Engineer, Google
 - Apache HBase contributor
 - Performance improvements and configuration framework

Dima Spivak (@dimaspivak)

- Distributed Software Engineer, Cloudera
 - Research Assistant (Physics), University of Minnesota
 - Apache HBase contributor
 - Test frameworks and automation

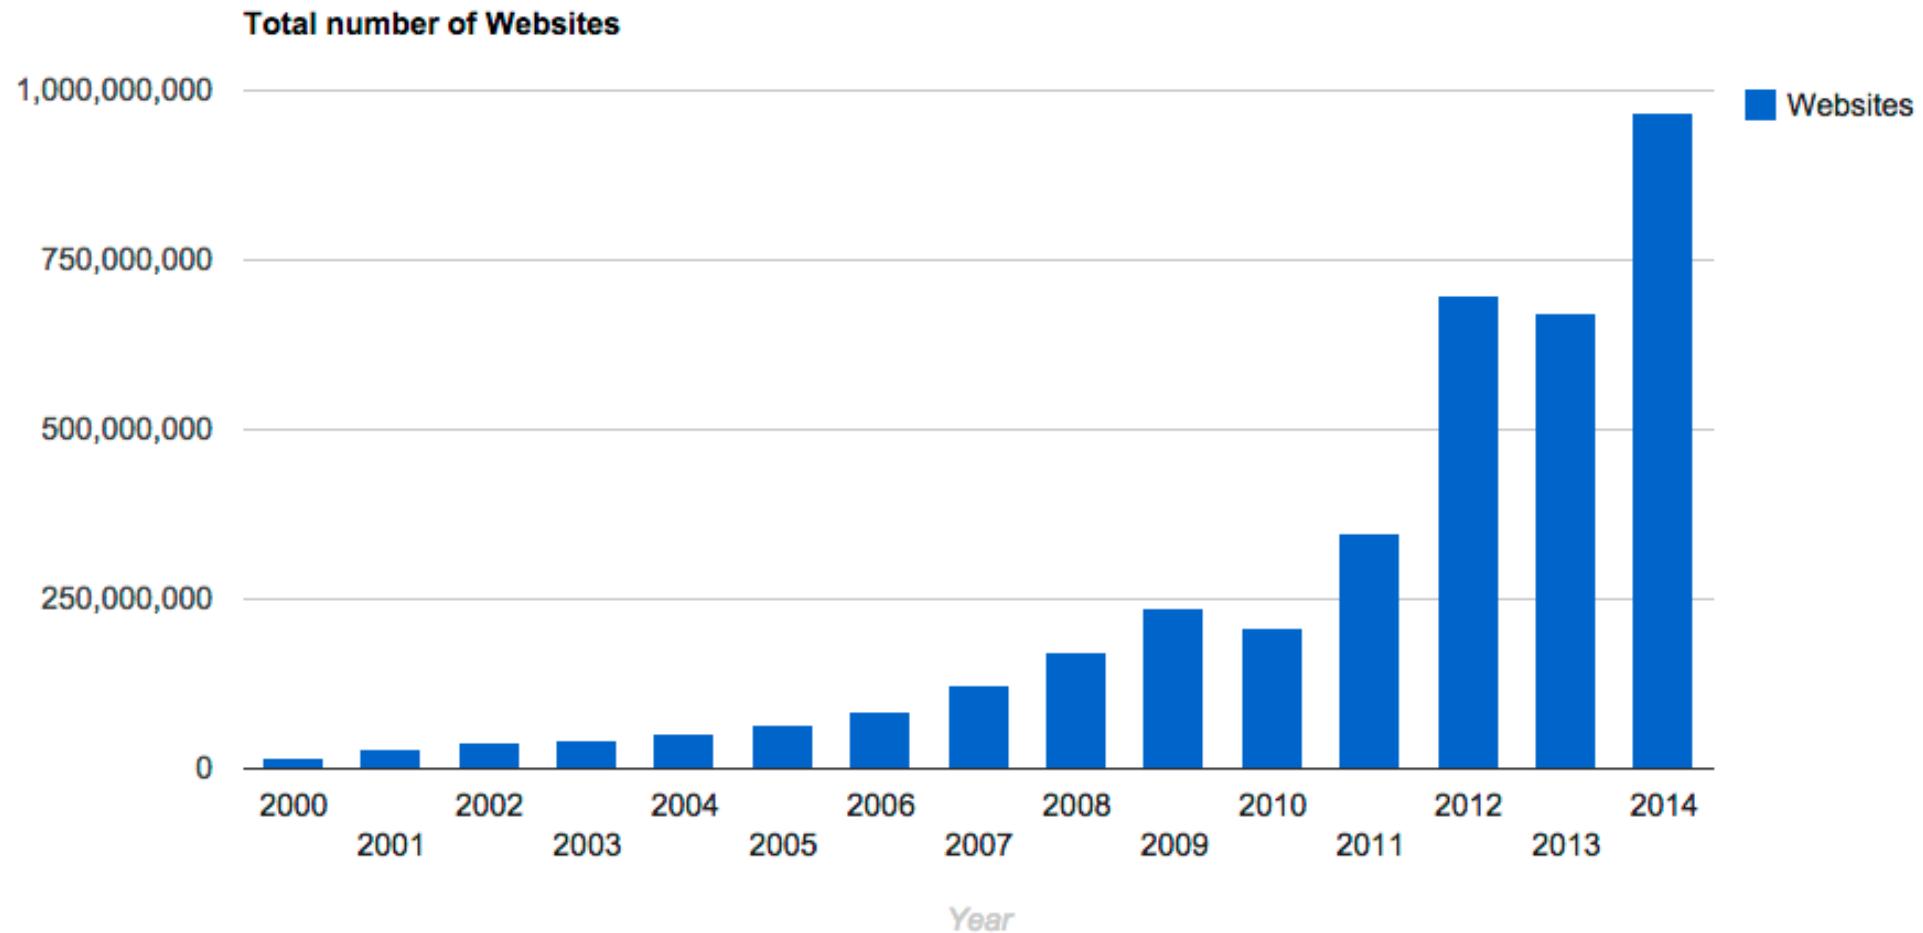
Contents

- Motivation
- Introduction to Apache HBase
 - Data model
- [Hands-On: Installation, HBase shell](#)
- **Break**
- A slightly more in-depth introduction to Apache HBase
 - Apache Hadoop
 - System internals
 - APIs
- **Break**

Contents

- Industry use cases & patterns
- Augmenting HBase
 - OpenTSDB
 - Apache Phoenix

Motivation



<http://www.internetlivestats.com/total-number-of-websites/>

Motivation

“We've known it for a long time: the web is big.”
– Jesse Alpert & Nissan Hajaj, Google

<http://googleblog.blogspot.com/2008/07/we-knew-web-was-big.html>

Motivation

- Indexing the internet has challenges:
 - Scale
 - Volume
 - Rate
 - Diversity of content
 - URLs
 - High-resolution images
 - Video
- Access

Motivation

Bigtable: A Distributed Storage System for Structured Data

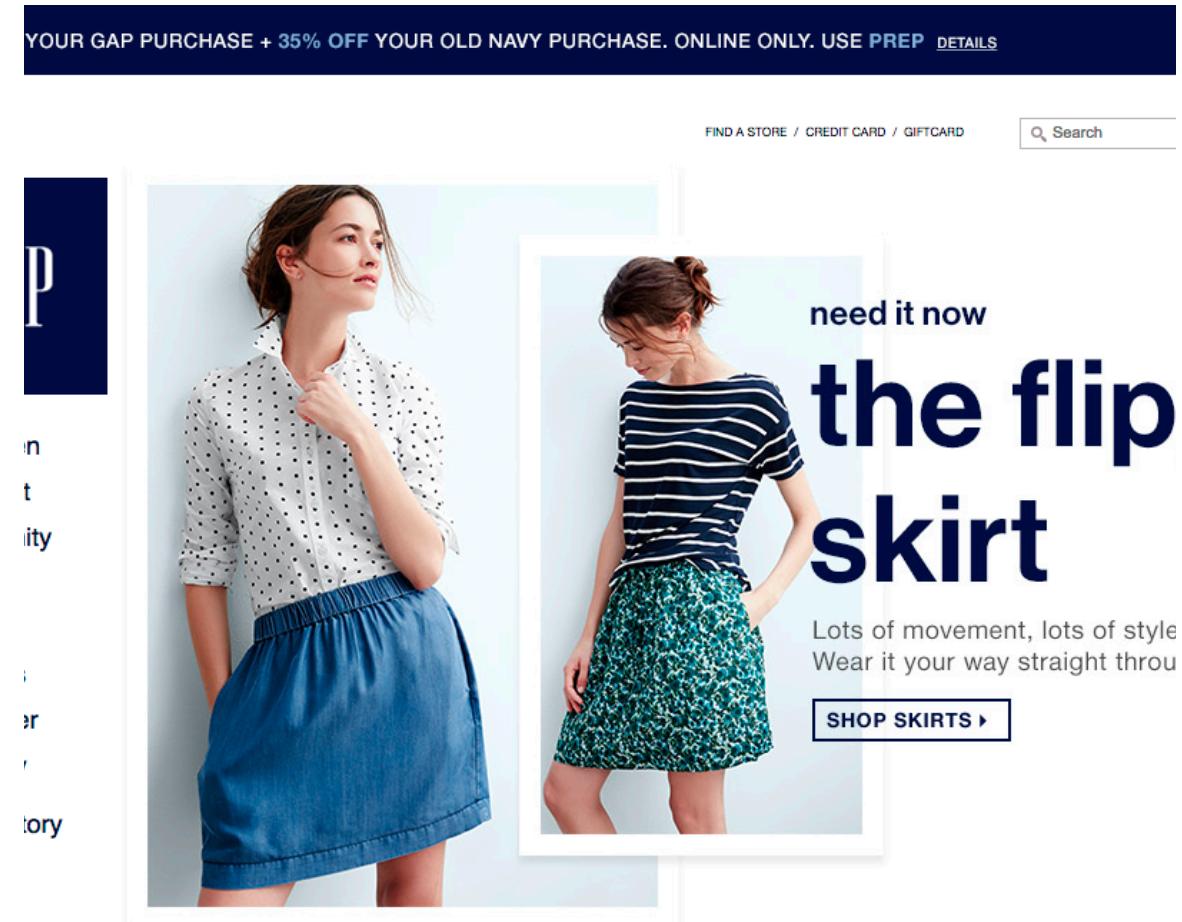
Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach
Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber

{fay,jeff,sanjay,wilsonh,kerr,m3b,tushar,fikes,gruber}@google.com

Google, Inc.

Motivation

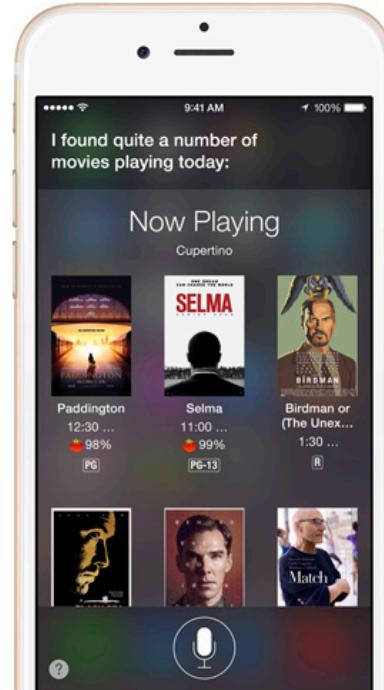
- What if you're not trying to index the internet?



Motivation

- Data for analytical processing
- User-facing real-time platforms

"What movies are playing today?"



Introduction to Apache HBase

- “Apache HBase™ is the Hadoop database, a distributed, scalable, big data store.”

<http://hbase.apache.org/>

Introduction to Apache HBase

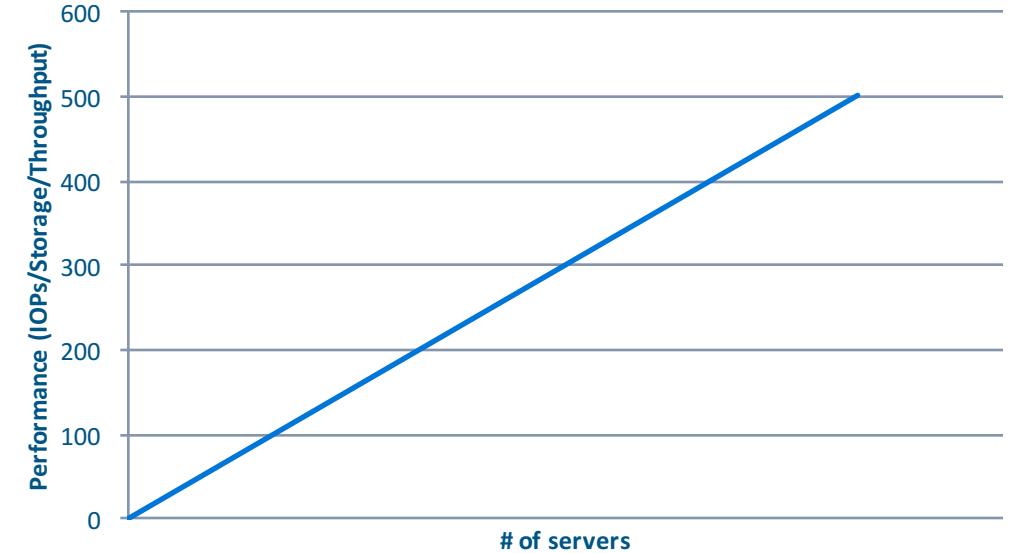
- Apache HBase is an **open source, horizontally scalable, consistent, random access**, low latency data store built on top of Apache Hadoop.

Apache HBase is open source

- Apache 2.0 License
- A community project with committers and contributors from diverse organizations
 - Cloudera, Facebook, Salesforce.com, Huawei, TrendMicro, eBay, HortonWorks, Intel, Twitter, ...
 - Code license means anyone can modify and use the code.

Apache HBase is horizontally scalable

- Adding more servers **linearly increases** performance and capacity
 - Storage capacity
 - Input/output operations
- Store and access data on **commodity servers**



- **Largest cluster:** > 3000 nodes, > 100 PB
- **Average cluster:** 10-40 nodes, 100-400TB

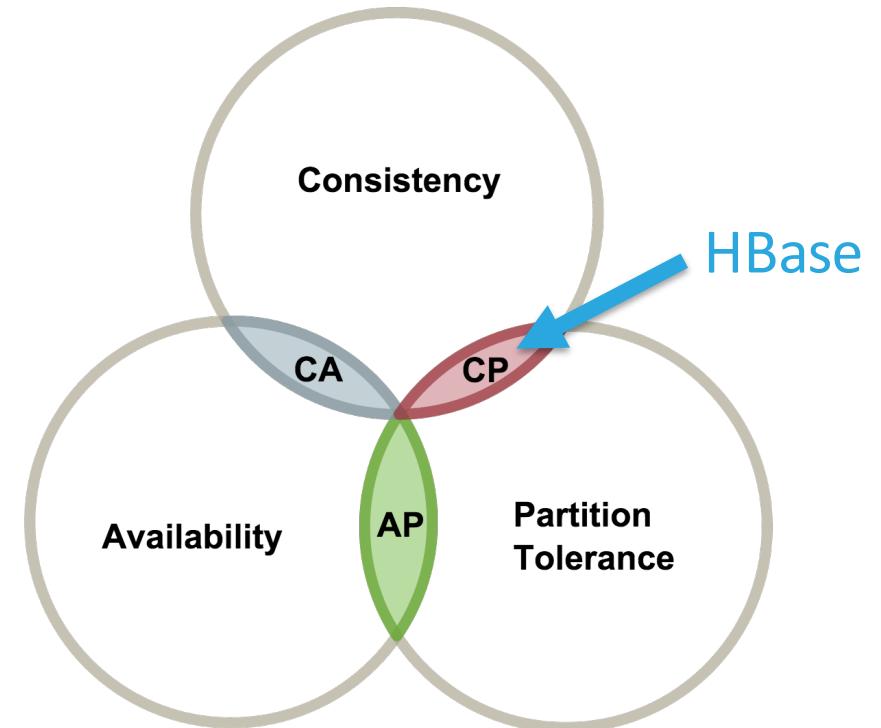
Apache HBase is horizontally scalable

- Commodity servers (crica 2015)
 - 12-24 2-4TB hard disks
 - 2 octa-core CPUs, 2-3 GHz
 - 64 - 512 GBs of RAM
 - 10 Gbps ethernet
 - \$5,000 - \$10,000 / machine



Apache HBase is consistent

- Brewer's theorem
 - Consistency
 - Availability
 - Partition tolerance



Data model

- Data is stored... in a big table
 - Sorted map datastore
- Tables consist of sorted **rows**, each of which has a primary **row key**
- Each row has a set of **columns**
- A column is specified as a **column family** and **column qualifier** pair
- A given **cell** (row, column family:column qualifier) can have different time-stamped values

Data model

Row key	info:height	info:state	roles:hadoop	roles:hbase
cutting	'9ft'	'CA'	'Founder'	
todd	'5ft7'	'CA'	'PMC' (ts=2011) 'Committer' (ts=2010)	'Committer'

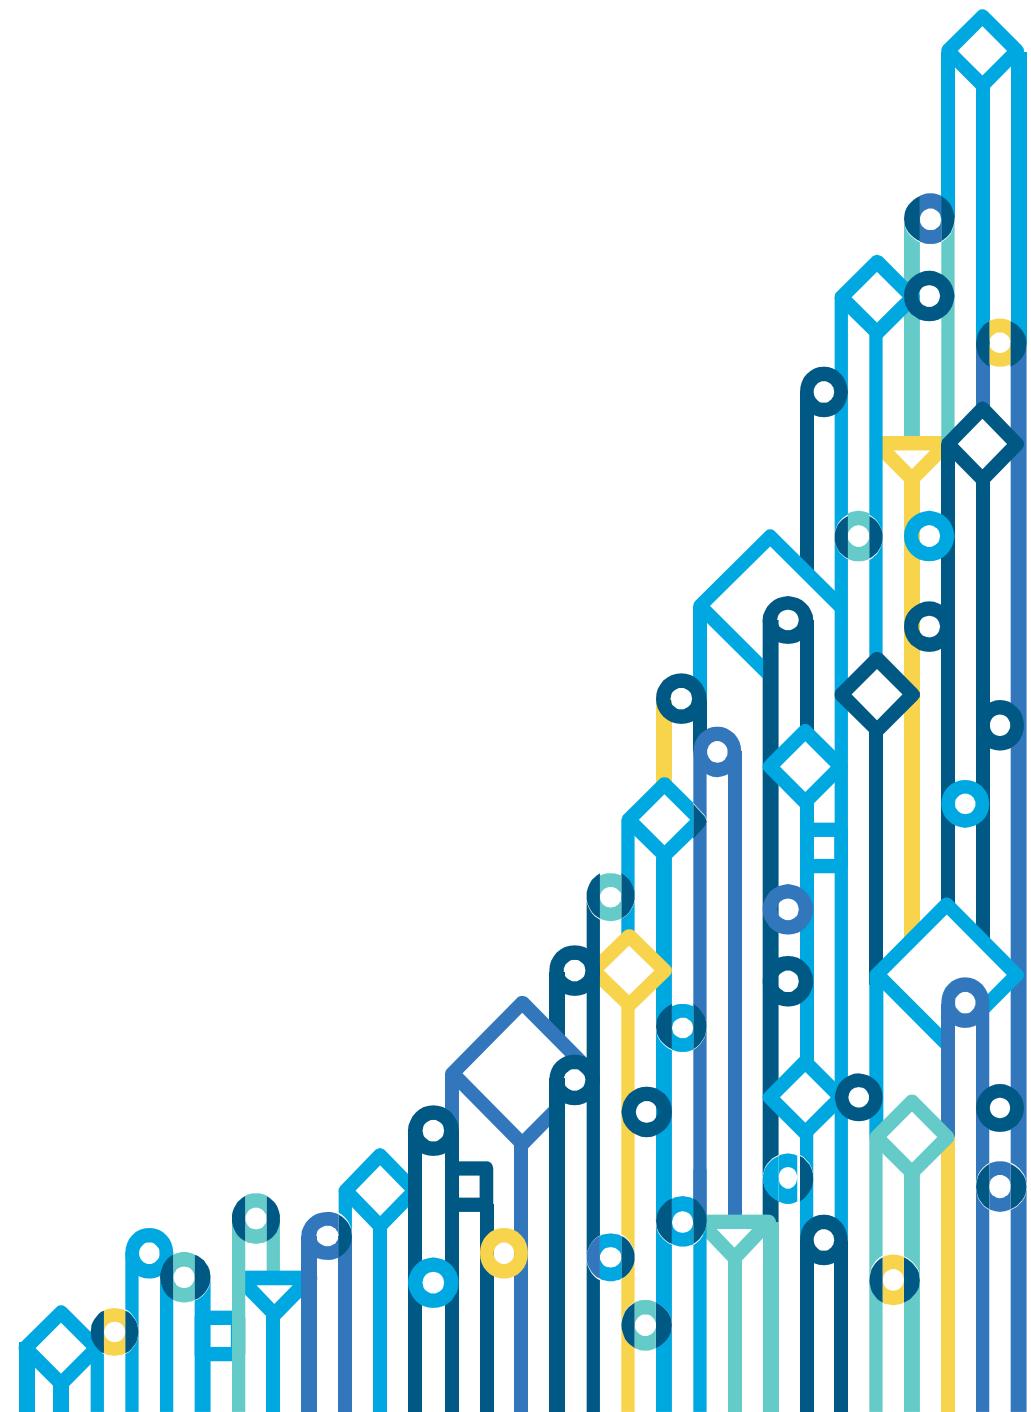


Hands On

Apache HBase installation

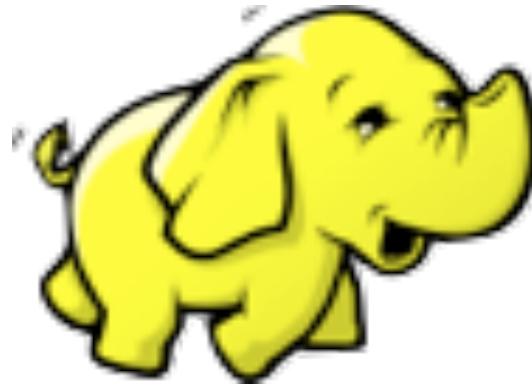
The HBase shell

<http://pastebin.com/nMkZeq5S>



Whats up for the next 1 hour?

Understanding basic architecture of



HDFS
(Apache Hadoop)

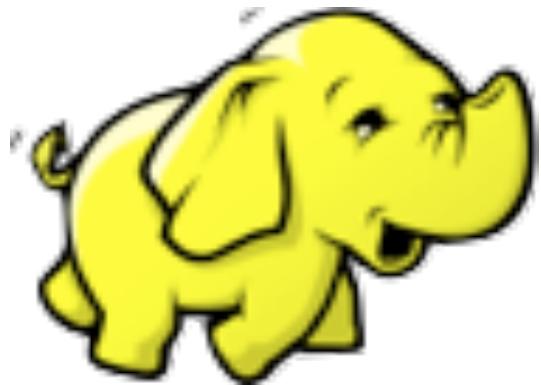


And, more Hands-On with Apache HBase.



cloudera
Break

Understanding basic architecture of



Hadoop (HDFS)

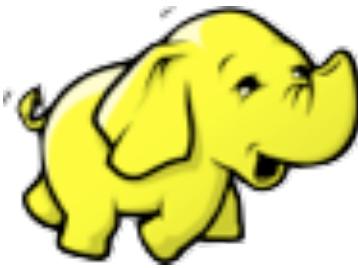
APACHE
HBASE

Apache Hadoop



open source
commodity servers
horizontally scalable
highly fault-tolerant
massive processing power

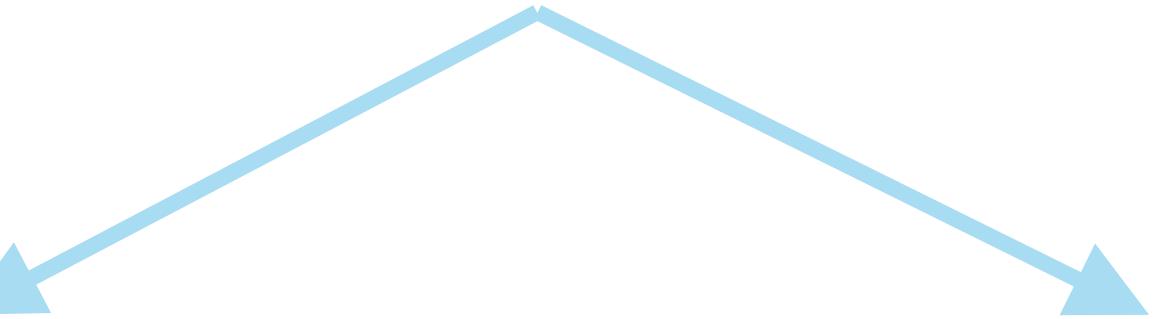
2 Core Components



Apache Hadoop

HDFS
(Hadoop Distributed File System)

MapReduce +
YARN



The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung

Google*

2003

GFS

- distributed file system
- commodity servers
- horizontally scalable
- highly fault-tolerant
- proprietary

HDFS

- **distributed file system**
- **commodity servers**
- **horizontally scalable**
- **highly fault-tolerant**
- **open source**

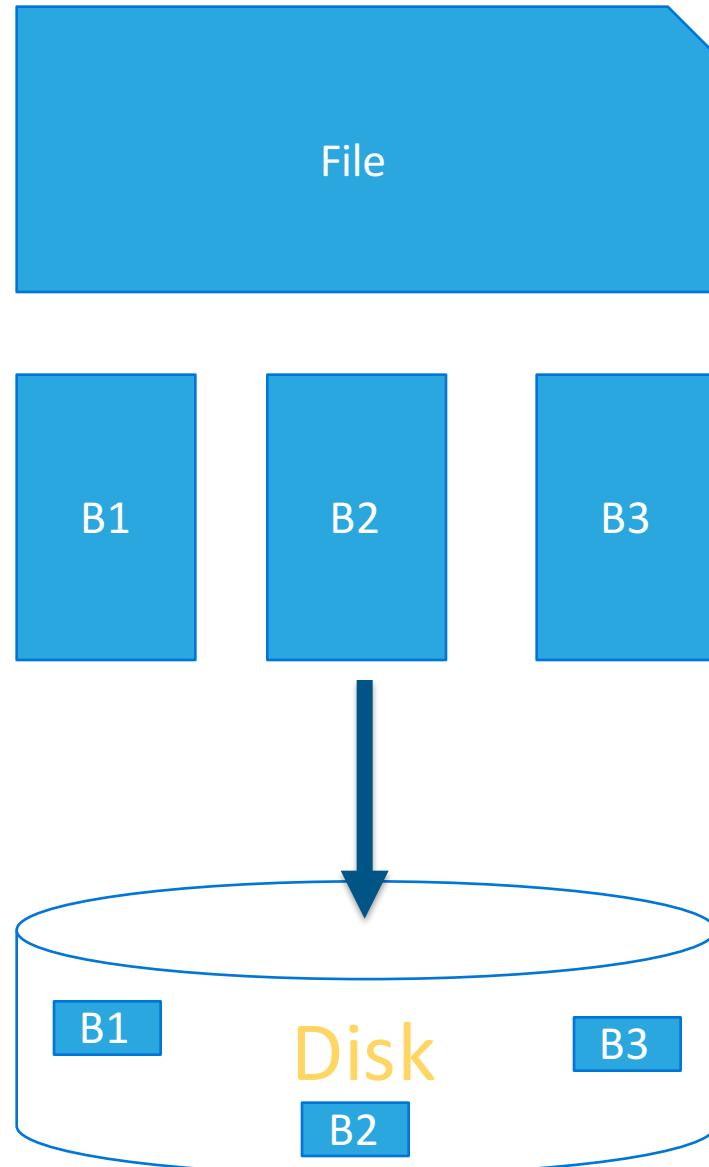
HDFS API

- File
 - Open, Close, Read, Write, Move, etc
- Directories
 - Create, Delete, etc
- Permissions
 - Owners, Groups, rwx permissions

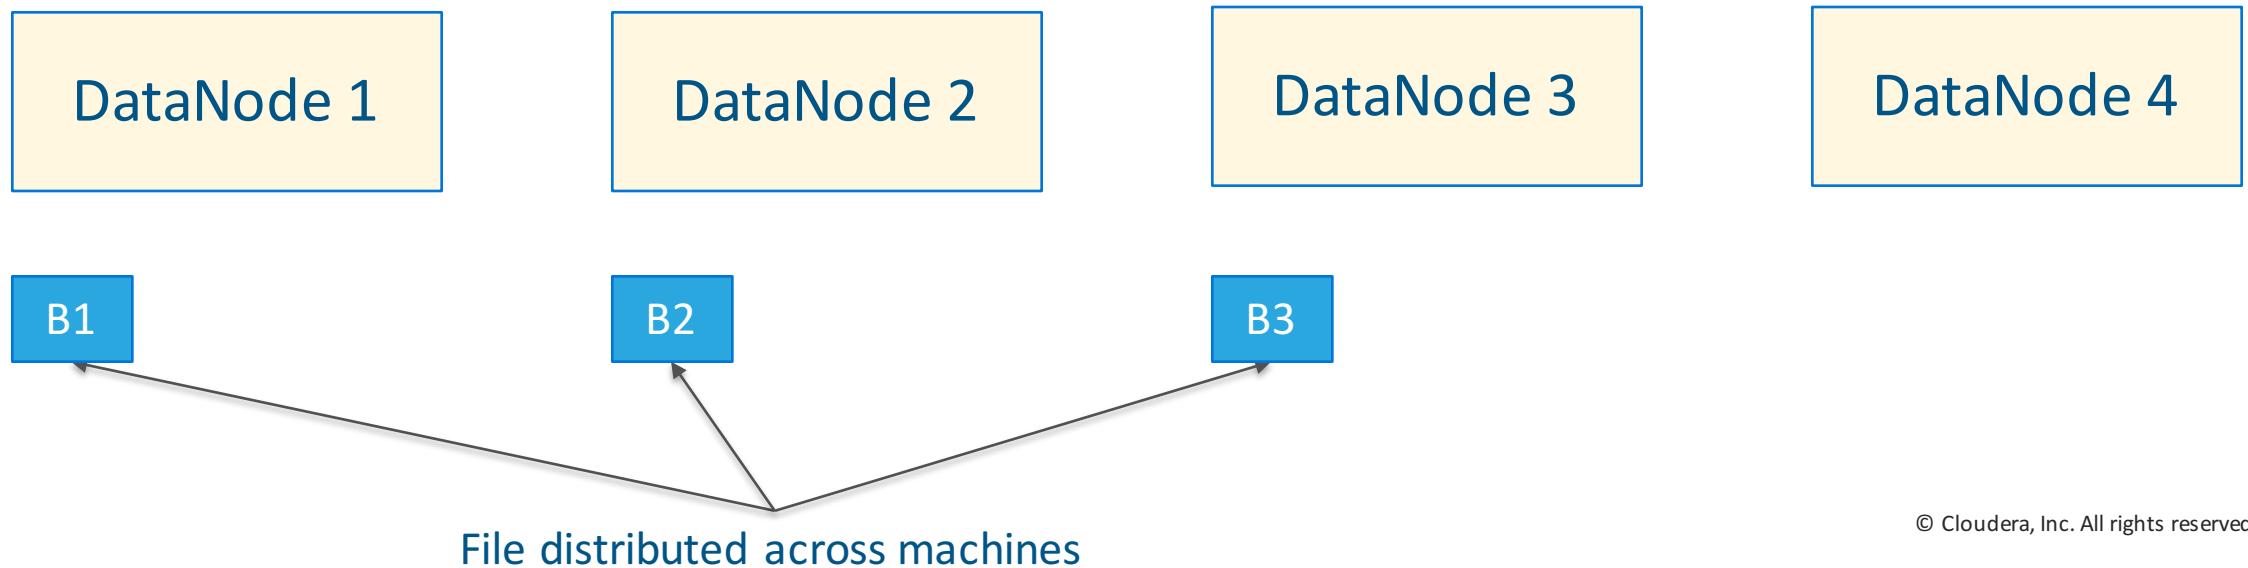
Basic Architecture of HDFS

Local file system

File system will split
the file into blocks



HDFS



HDFS DataNode

DataNode 1

DataNode 2

DataNode 3

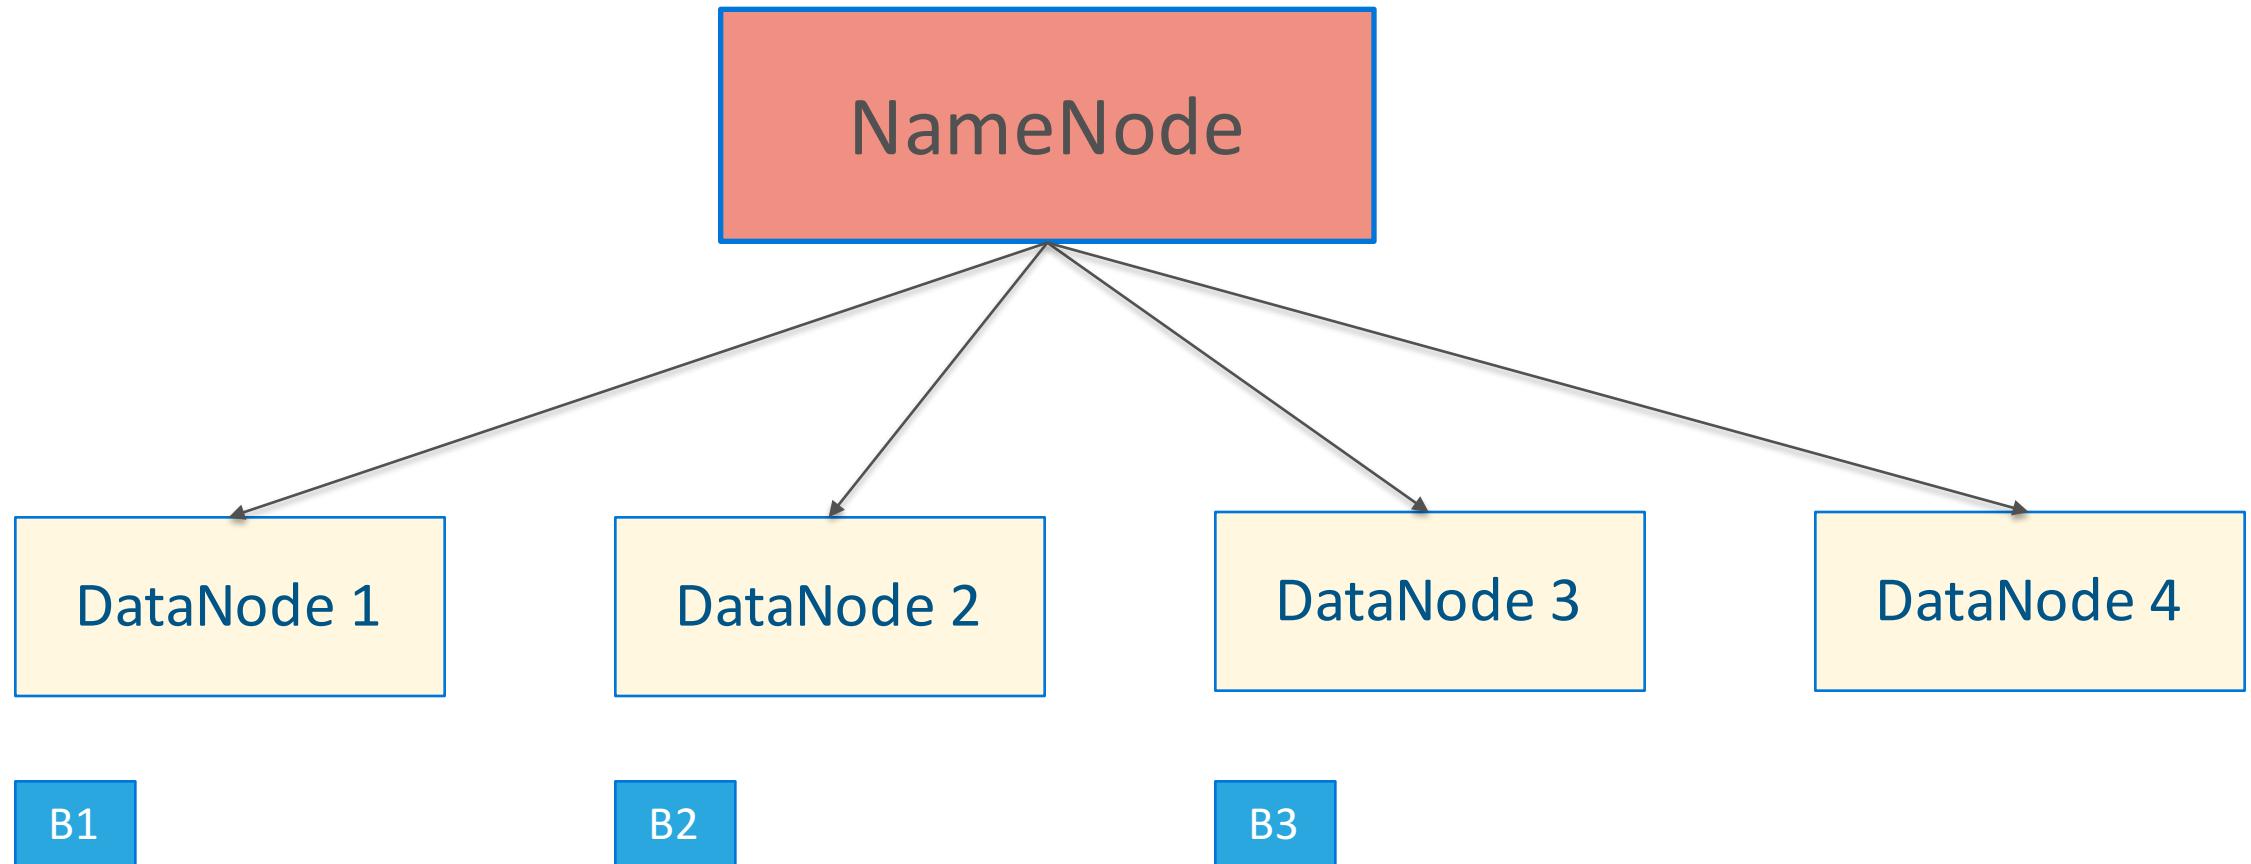
DataNode 4

B1

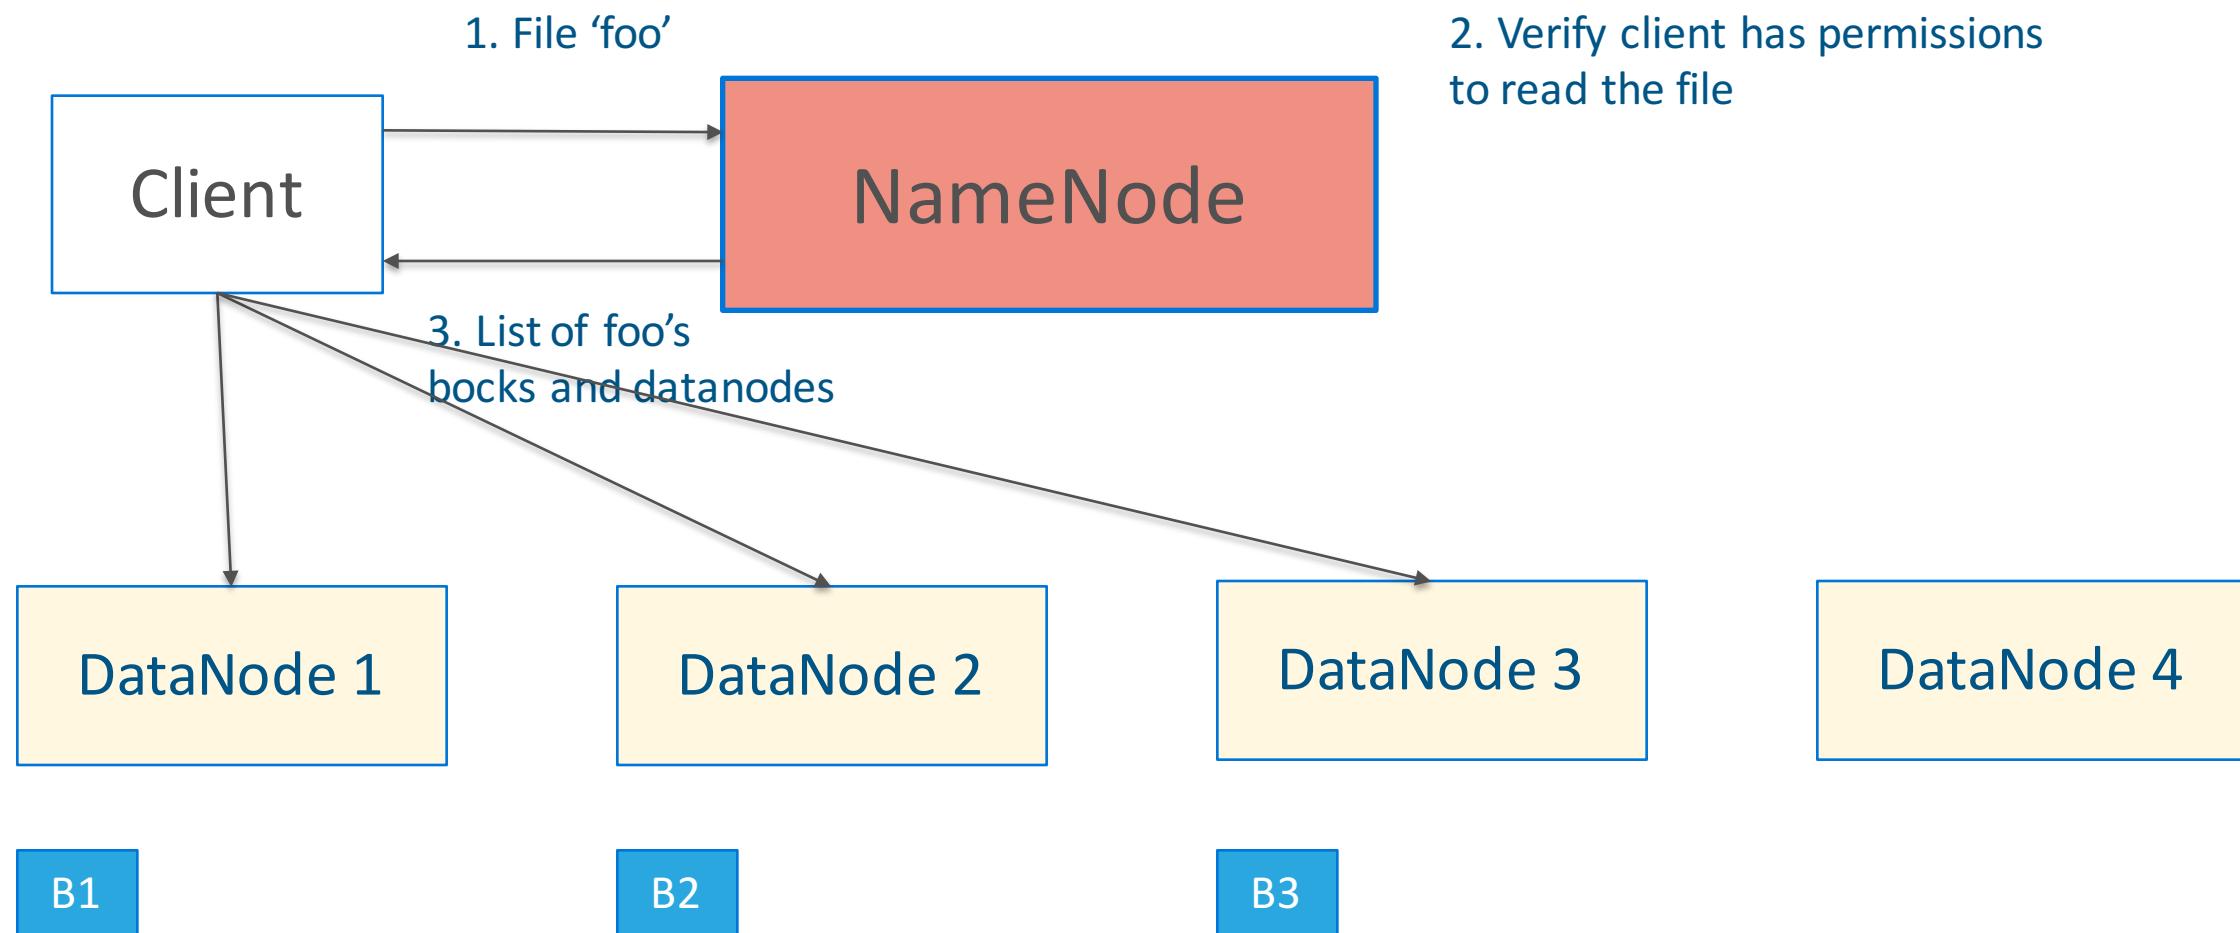
B2

B3

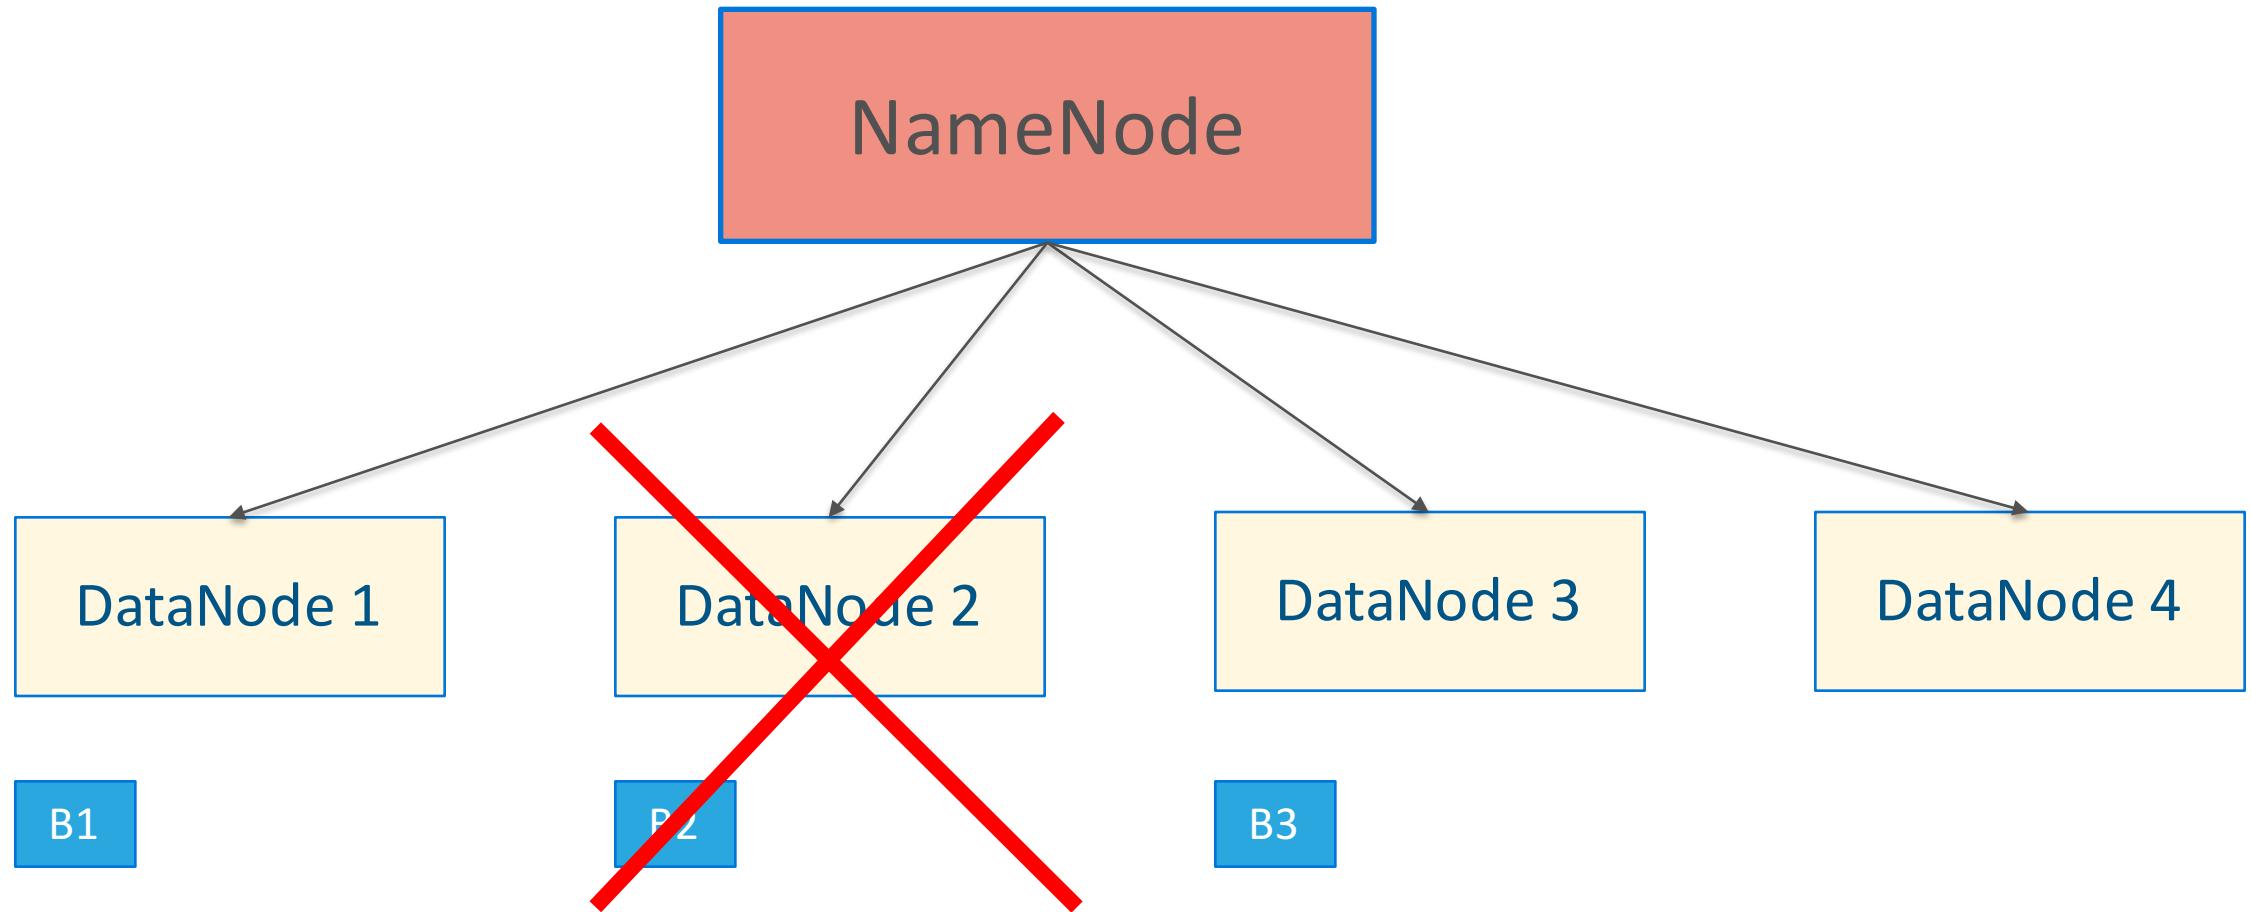
HDFS NameNode



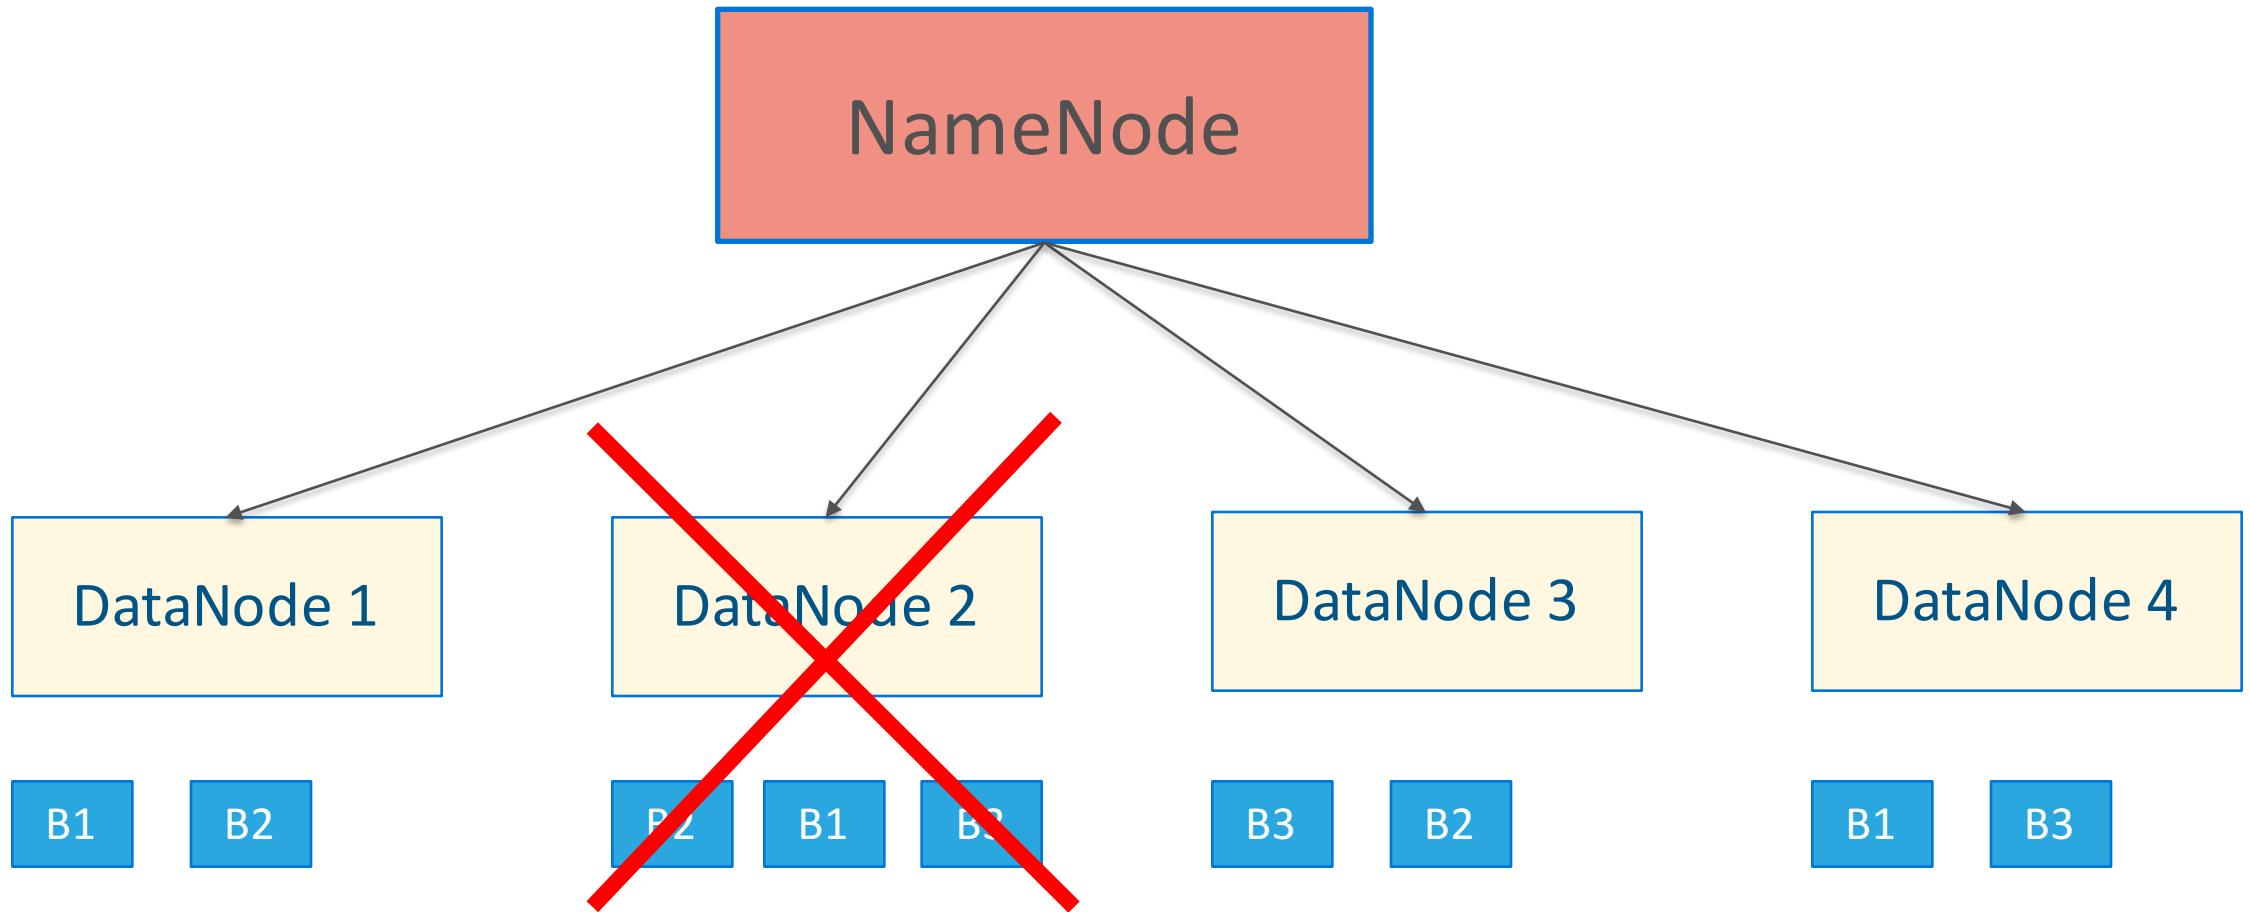
HDFS Reading a file



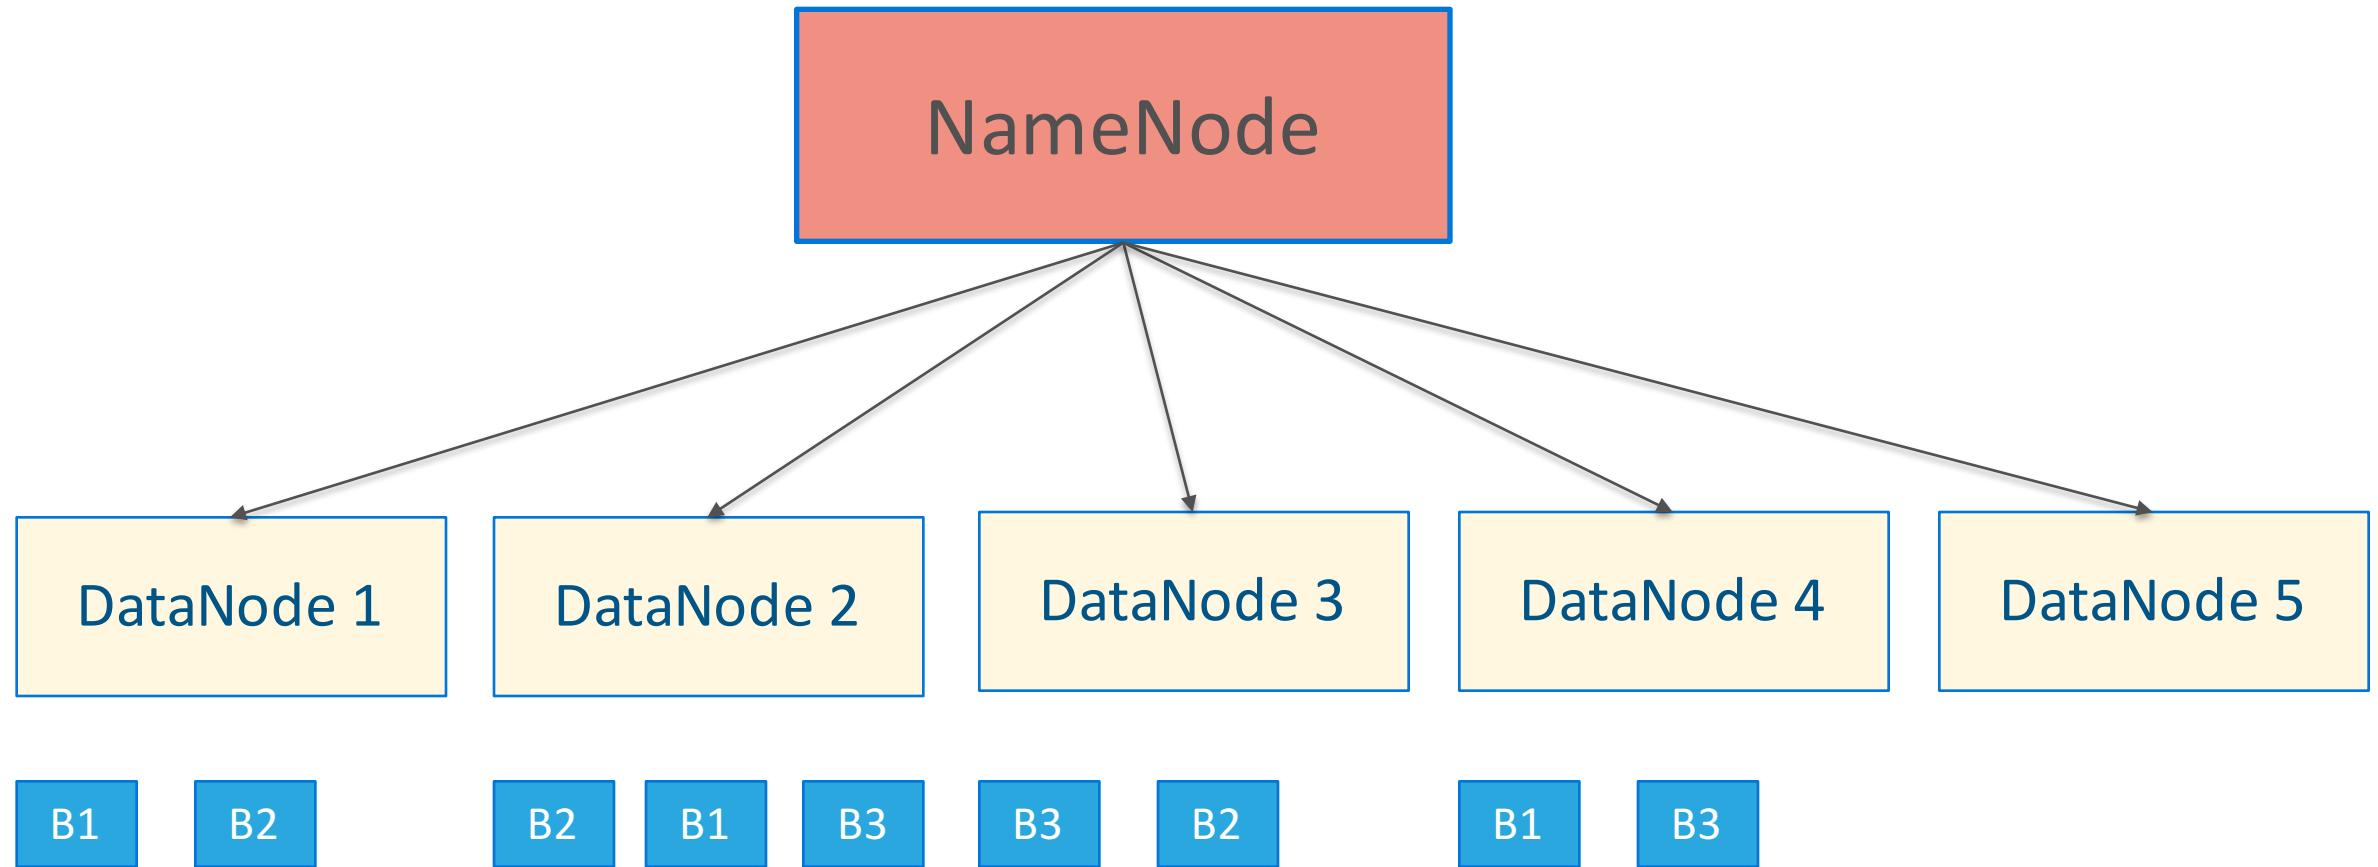
HDFS Fault tolerance



HDFS Redundancy



HDFS Horizontal Scalability



Let's look at some existing HDFS systems...

- **Yahoo! HDFS Clusters**
40k+ servers, 100k+ CPUs, 450PB data
- **Facebook HDFS Cluster**
15TB new data per day
1200+ machines, 30PB in one cluster
- *Lots of 5-40 node clusters at companies without petabytes of data (web, retail, finance, telecom, research, government)*

But.... there are restrictions!

It's not a magic wand!

Files are **append only**

- Access Model : Write-once-read-many
- Can not change existing contents

Not designed for small files

- Block sizes are in MB (default 128MB)
 - Designed for typical GBs / TBs of file sizes
 - Normal files system have 4kb block size!

Summary

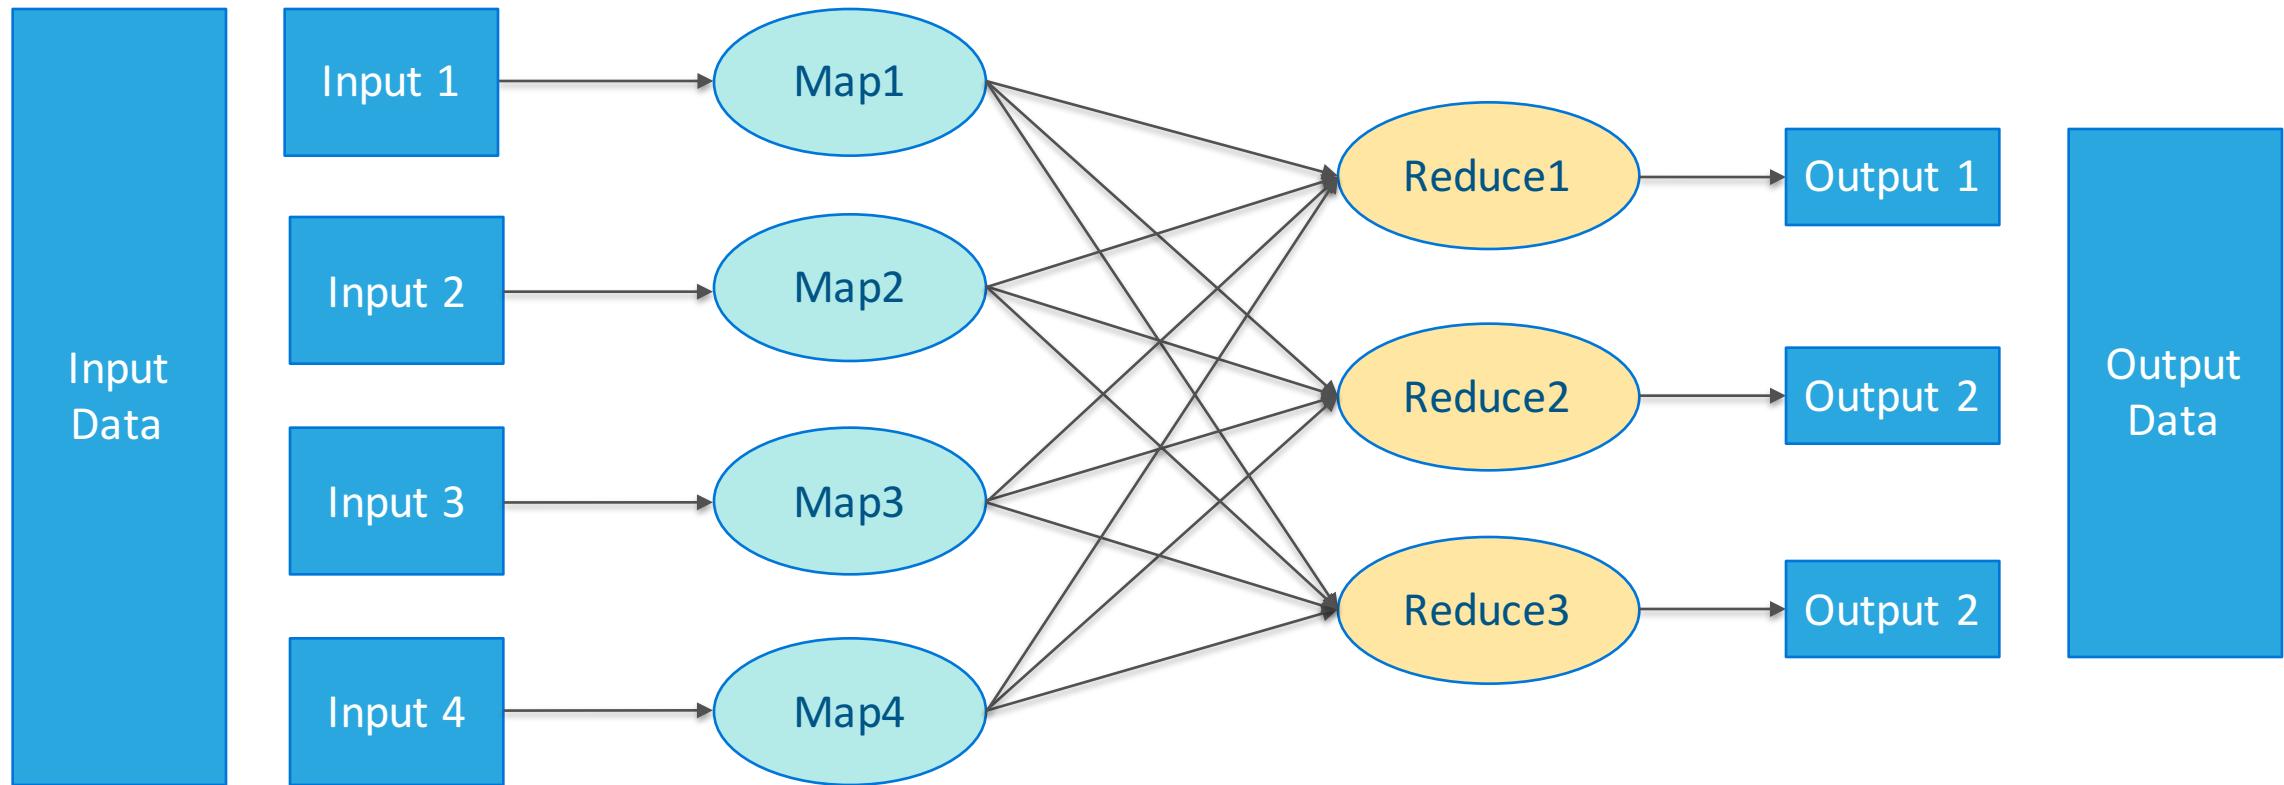
HDFS is a great distributed file system!

- Store massive data
- Scalable
- High throughput
- Fault tolerance

MapReduce

- Distributed processing framework
- Commodity machines
- Fault tolerance

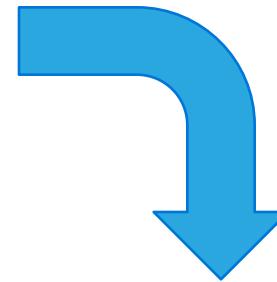
MapReduce





HBase Architecture

Name	Weight	UPC	Price
Prego Tomato Sauce	67 Oz	xxxxxxxx	\$4.97
Trumoo Lowfat Chocolate Milk	128 Oz	xxxxxxxx	\$2.99
Gatorade Lemon-Lime	64 Oz	xxxxxxxx	\$3.98



	info:weight	info:upc	info:price
Prego Tomato Sauce	67 Oz	xxxxxxxx	\$4.97
Trumoo Lowfat Chocolate Milk	128 Oz	xxxxxxxx	\$2.99
Gatorade Lemon-Lime	64 Oz	xxxxxxxx	\$3.98

HBase Architecture

	info:weight	info:upc	info:price
Prego Tomato Sauce	67 Oz	xxxxxxxx	\$4.97
Trumoo Lowfat Chocolate Milk	128 Oz	xxxxxxxx	\$2.99
Gatorade Lemon-Lime	64 Oz	xxxxxxxx	\$3.98

HBase Architecture

	info:weight	info:upc	info:price
Prego Tomato Sauce	67 Oz	xxxxxxxx	\$4.97
Trumoo Lowfat Chocolate Milk	128 Oz	xxxxxxxx	\$2.99
Gatorade Lemon-Lime	64 Oz	xxxxxxxx	\$3.98
A New Product	4 Oz	xxxxxxxx	\$9.99

HBase Architecture

	info:weight	info:upc	info:price
Prego Tomato Sauce	67 Oz	xxxxxxxx	\$4.97
Trumoo Lowfat Chocolate Milk	128 Oz	xxxxxxxx	\$2.99
Gatorade Lemon-Lime	64 Oz	xxxxxxxx	\$3.98
A New Product	4 Oz	xxxxxxxx	\$9.99
Yet Another New Product	8 Oz	xxxxxxxx	\$19.99

HBase Architecture

	info:weight	info:upc	info:price
Prego Tomato Sauce	67 Oz	xxxxxxxx	\$4.97
Trumoo Lowfat Chocolate Milk	128 Oz	xxxxxxxx	\$2.99
Gatorade Lemon-Lime	64 Oz	xxxxxxxx	\$3.98
A New Product	4 Oz	xxxxxxxx	\$9.99
Yet Another New Product	8 Oz	xxxxxxxx	\$19.99
Four More Products (1)	16 Oz	xxxxxxxx	\$9.99
Four More Products (2)	16 Oz	xxxxxxxx	\$9.99
Four More Products (3)	16 Oz	xxxxxxxx	\$9.99
Four More Products (4)	16 Oz	xxxxxxxx	\$9.99

HBase Architecture | Regions

- Tables are chopped up into regions (split).
- A region is only served by a single “region server” at a time.
- RegionServer can serve multiple regions.

HBase Architecture | Regions

	info:weight	info:upc	info:price
A New Product	4 Oz	xxxxxxxx	\$9.99
Four More Products (1)	16 Oz	xxxxxxxx	\$9.99
Four More Products (2)	16 Oz	xxxxxxxx	\$9.99
Four More Products (3)	16 Oz	xxxxxxxx	\$9.99
Four More Products (4)	16 Oz	xxxxxxxx	\$9.99

Served by RegionServer on machine 2

	info:weight	info:upc	info:price
Prego Tomato Sauce	67 Oz	xxxxxxxx	\$4.97
Trumoo Lowfat Chocolate Milk	128 Oz	xxxxxxxx	\$2.99
Gatorade Lemon-Lime	64 Oz	xxxxxxxx	\$3.98
Yet Another New Product	8 Oz	xxxxxxxx	\$19.99

Served by RegionServer on machine 3

HBase Architecture | Regions

	info:price	info:upc	info:weight
A New Product	\$9.99	xxxxxxxx	4 Oz
Four More Products (1)	\$9.99	xxxxxxxx	16 Oz
Four More Products (2)	\$9.99	xxxxxxxx	16 Oz
Four More Products (3)	\$9.99	xxxxxxxx	16 Oz
Four More Products (4)	\$9.99	xxxxxxxx	16 Oz

Served by RegionServer on machine 2

	info:price	info:upc	info:weight
Gatorade Lemon-Lime	\$3.98	xxxxxxxx	64 Oz
Prego Tomato Sauce	\$4.97	xxxxxxxx	67 Oz
Trumoo Lowfat Chocolate Milk	\$2.99	xxxxxxxx	128 Oz
Yet Another New Product	\$19.99	xxxxxxxx	8 Oz

Served by RegionServer on machine 3

HBase Architecture

	info:price	info:upc	info:weight	available:store1	available:store2	available:store3
A New Product	\$9.99	xxxxxxxx	4 Oz	Yes	Yes	Yes
Four More Products (1)	\$9.99	xxxxxxxx	16 Oz	Yes	Yes	Yes
Four More Products (2)	\$9.99	xxxxxxxx	16 Oz	Yes	Yes	No
Four More Products (3)	\$9.99	xxxxxxxx	16 Oz	Yes	Yes	No
Four More Products (4)	\$9.99	xxxxxxxx	16 Oz	No	No	No

	info:price	info:upc	info:weight	available:store1	available:store2	available:store3
Gatorade Lemon-Lime	\$3.98	xxxxxxxx	64 Oz	Yes	Yes	Yes
Prego Tomato Sauce	\$4.97	xxxxxxxx	67 Oz	Yes	No	Yes
Trumoo Lowfat Chocolate Milk	\$2.99	xxxxxxxx	128 Oz	No	No	Yes
Yet Another New Product	\$19.99	xxxxxxxx	8 Oz	Yes	Yes	Yes

HBase Architecture

	info:price	info:upc	info:weight	available:store1	available:store2	available:store3
A New Product	\$9.99	xxxxxxxx	4 Oz	Yes	Yes	Yes
Four More Products (1)	\$9.99	xxxxxxxx	16 Oz	Yes	Yes	Yes
Four More Products (2)	\$9.99	xxxxxxxx	16 Oz	Yes	Yes	No
Four More Products (3)	\$9.99	xxxxxxxx	16 Oz	Yes	Yes	No
Four More Products (4)	\$9.99	xxxxxxxx	16 Oz	No	No	No

	info:price	info:upc	info:weight	available:store1	available:store2	available:store3
Gatorade Lemon-Lime	\$3.98	xxxxxxxx	64 Oz	Yes	Yes	Yes
Prego Tomato Sauce	\$4.97	xxxxxxxx	67 Oz	Yes	No	Yes
Trumoo Lowfat Chocolate Milk	\$2.99	xxxxxxxx	128 Oz	No	No	Yes
Yet Another New Product	\$19.99	xxxxxxxx	8 Oz	Yes	Yes	Yes

HBase Architecture | Column family

- A column family is a set of related columns.
 - Group sets of columns that have similar access patterns.
- Tune read performance.
 - Compression
 - Version retention policies
 - Cache priority

HBase Architecture

Store				Region
	info: price	info: upc	info: weight	
Gatorade Lemon-Lime	\$3.98	xxxxxxxx x	64 Oz	
Prego Tomato Sauce	\$4.97	xxxxxxxx x	67 Oz	
Trumoo Lowfat Chocolate Milk	\$2.99	xxxxxxxx x	128 Oz	

	available: store1	available: store2	available: store3
Gatorade Lemon-Lime	Yes	Yes	Yes
Prego Tomato Sauce	Yes	No	Yes
Trumoo Lowfat Chocolate Milk	No	No	Yes

HBase Architecture | Write Path

1. Client creates a row to put.
2. Client checks with meta* for which RegionServer hosts this row.
3. Row is written into write-ahead log (WAL).
4. Row is written to MemStore.

HBase Architecture | Write Path



HBase Architecture | Write Path

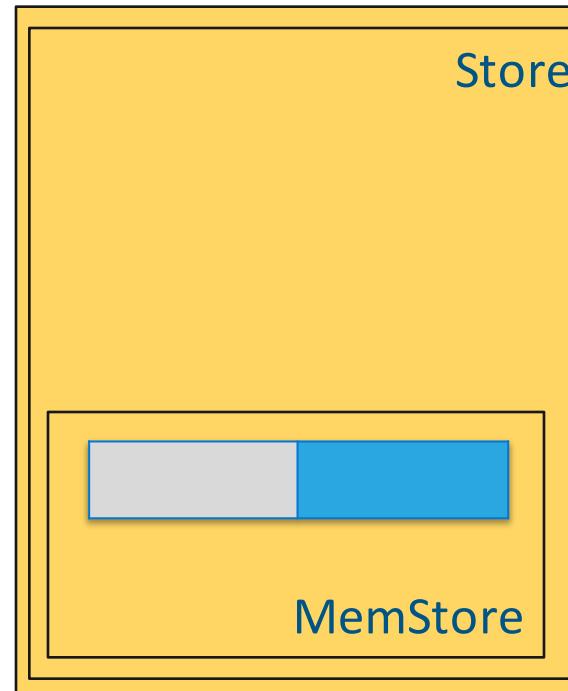
Put



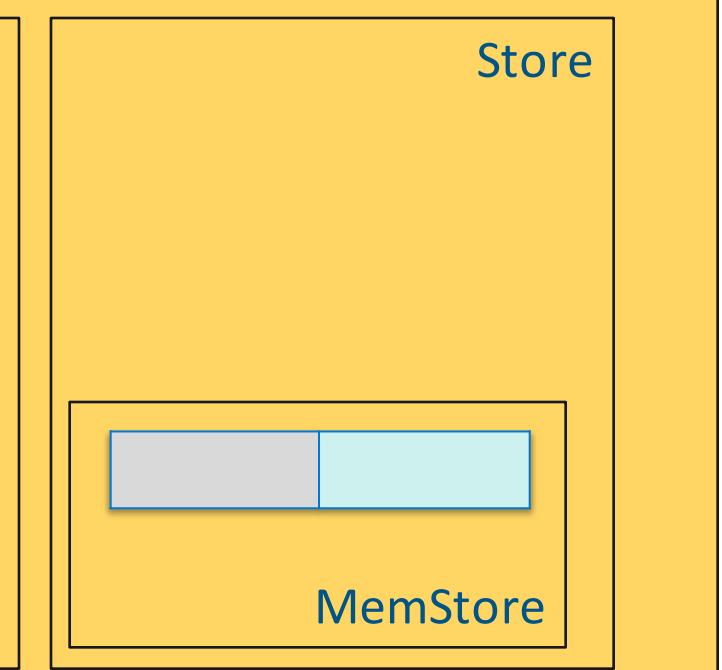
WAL



Store



Store



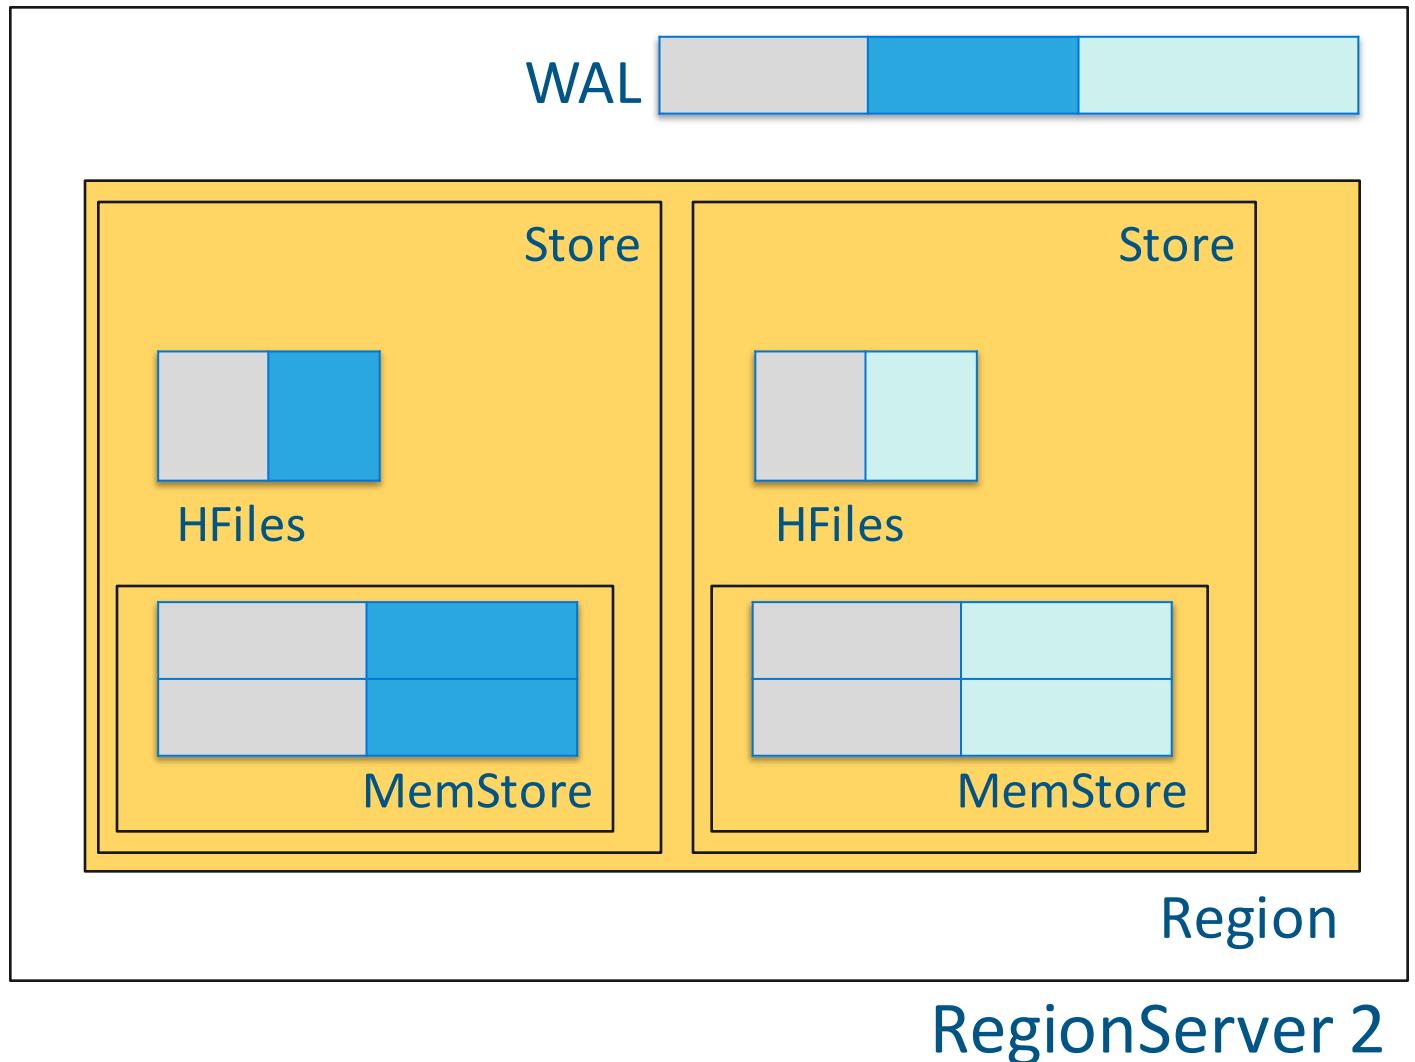
Region

RegionServer 2

HBase Architecture | Write Path

- When MemStore gets full or a flush is triggered, contents of MemStore are flushed to disk.
 - HFiles are created.

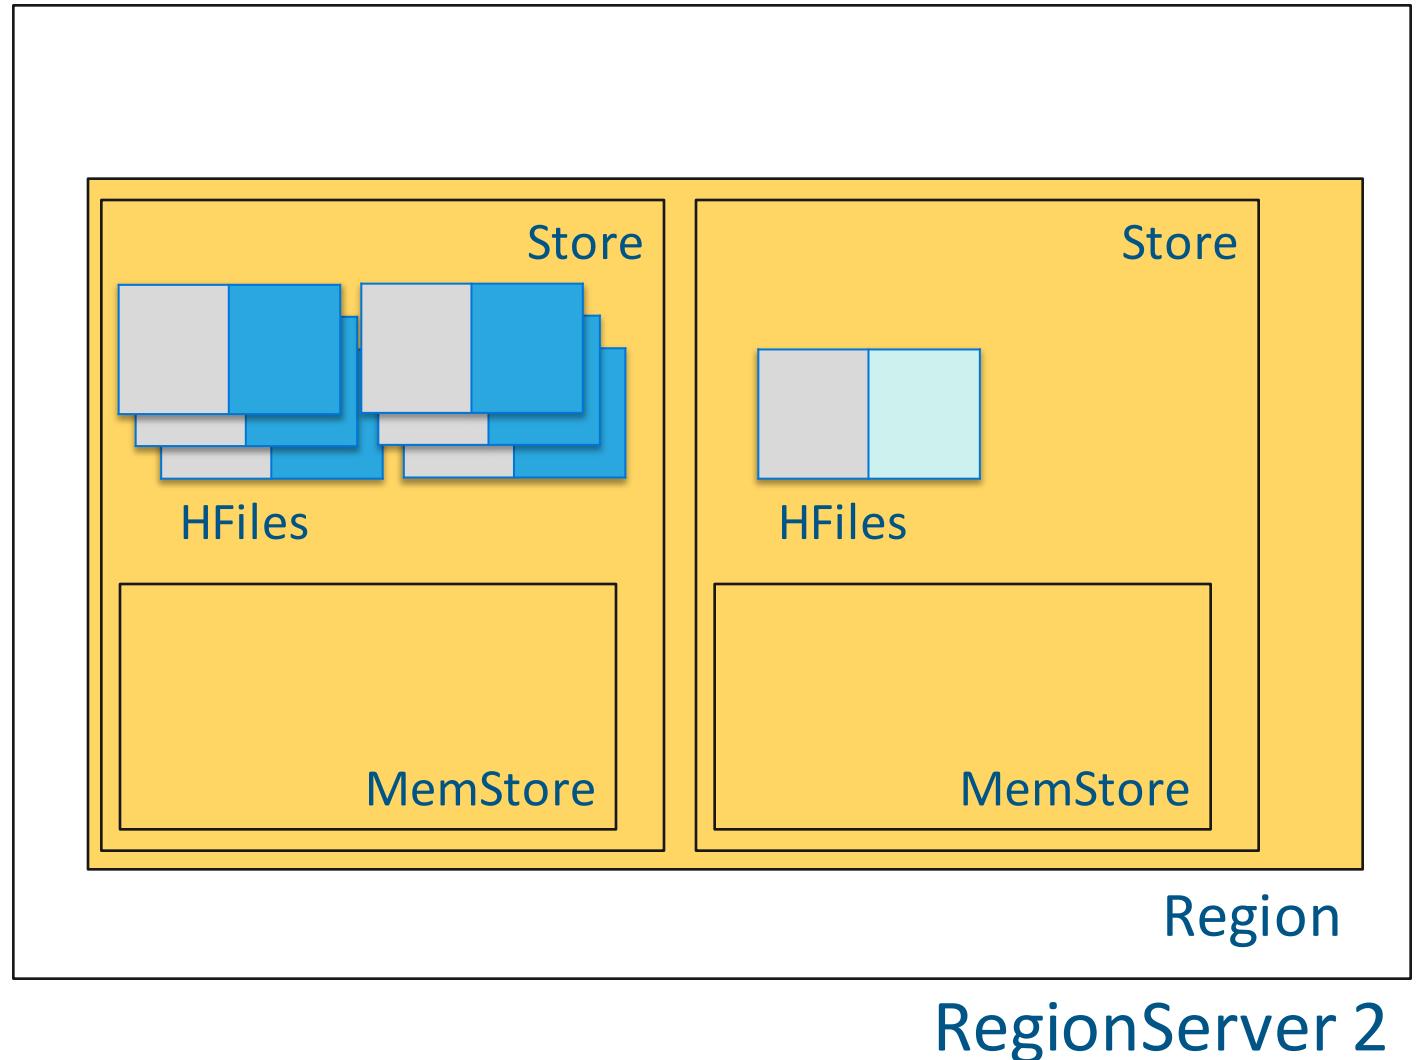
HBase Architecture | Write Path



HBase Architecture | Write Path

- Each subsequent write repeats this process.
 - Write to WAL.
 - Write to MemStore.
 - Flush when MemStore fills or a flush is triggered.
 - Create HFiles.
- Lots of HFiles in a Region mean lots of disk seeks on read.
 - Might be better to combine (compact) HFiles.

HBase Architecture



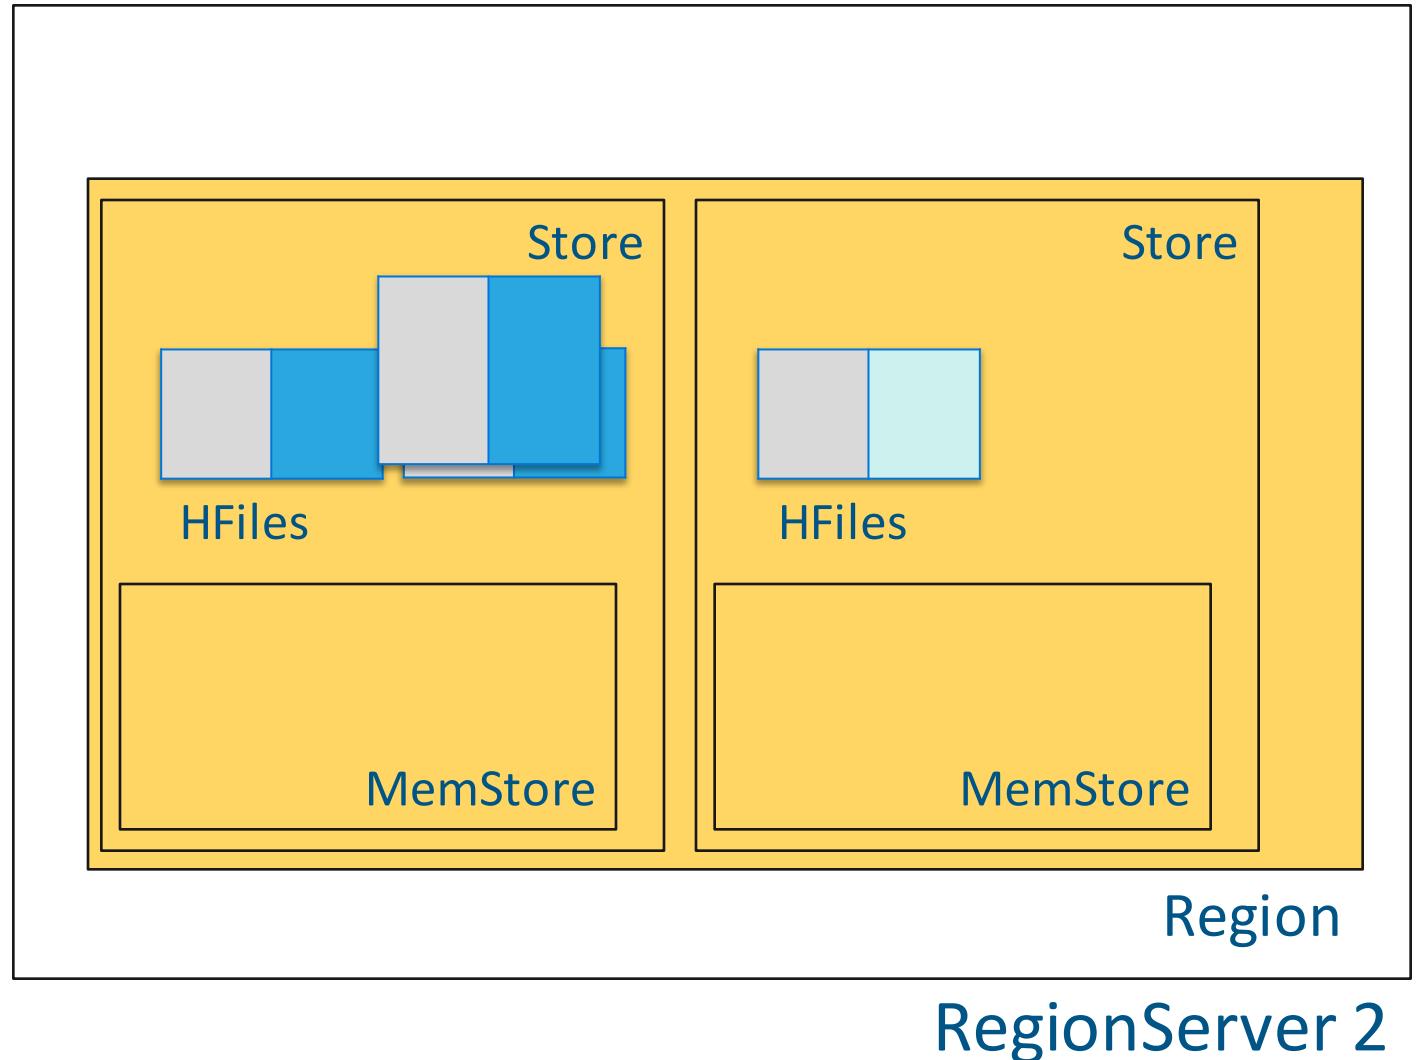
HBase Architecture | Compactions

- Minor compactions
 - Merge some HFiles (in a given Store).
- Major compactions
 - Merge all HFiles (in a given Store).
 - Take care of other HBase housekeeping tasks.

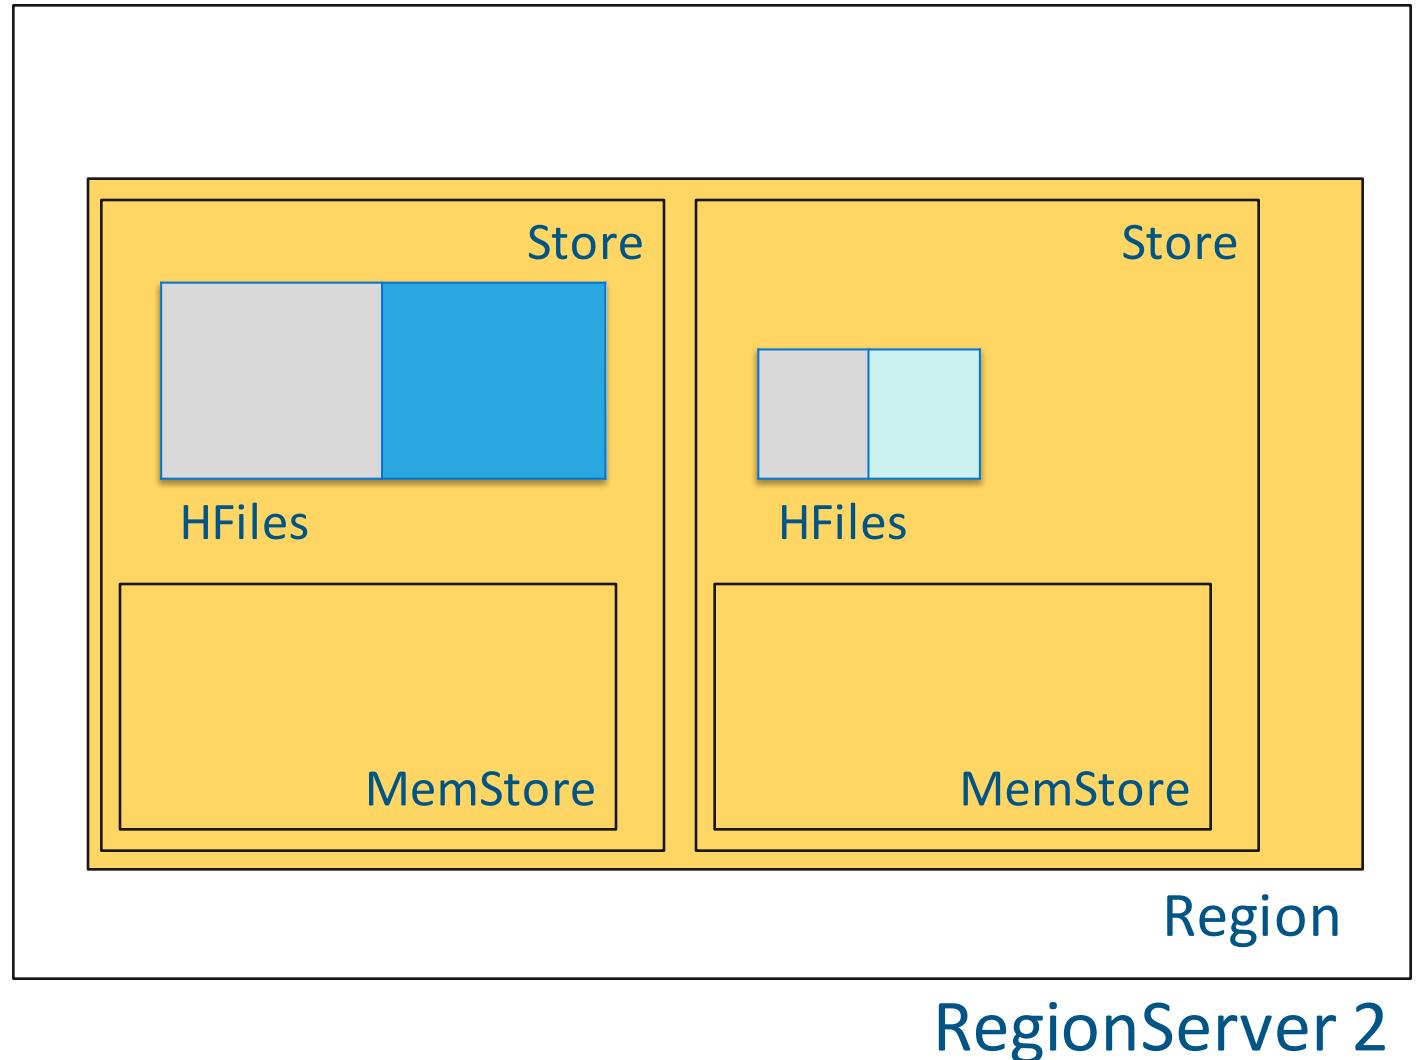
HBase Architecture | Compaction



HBase Architecture | Minor compaction



HBase Architecture | Major compaction



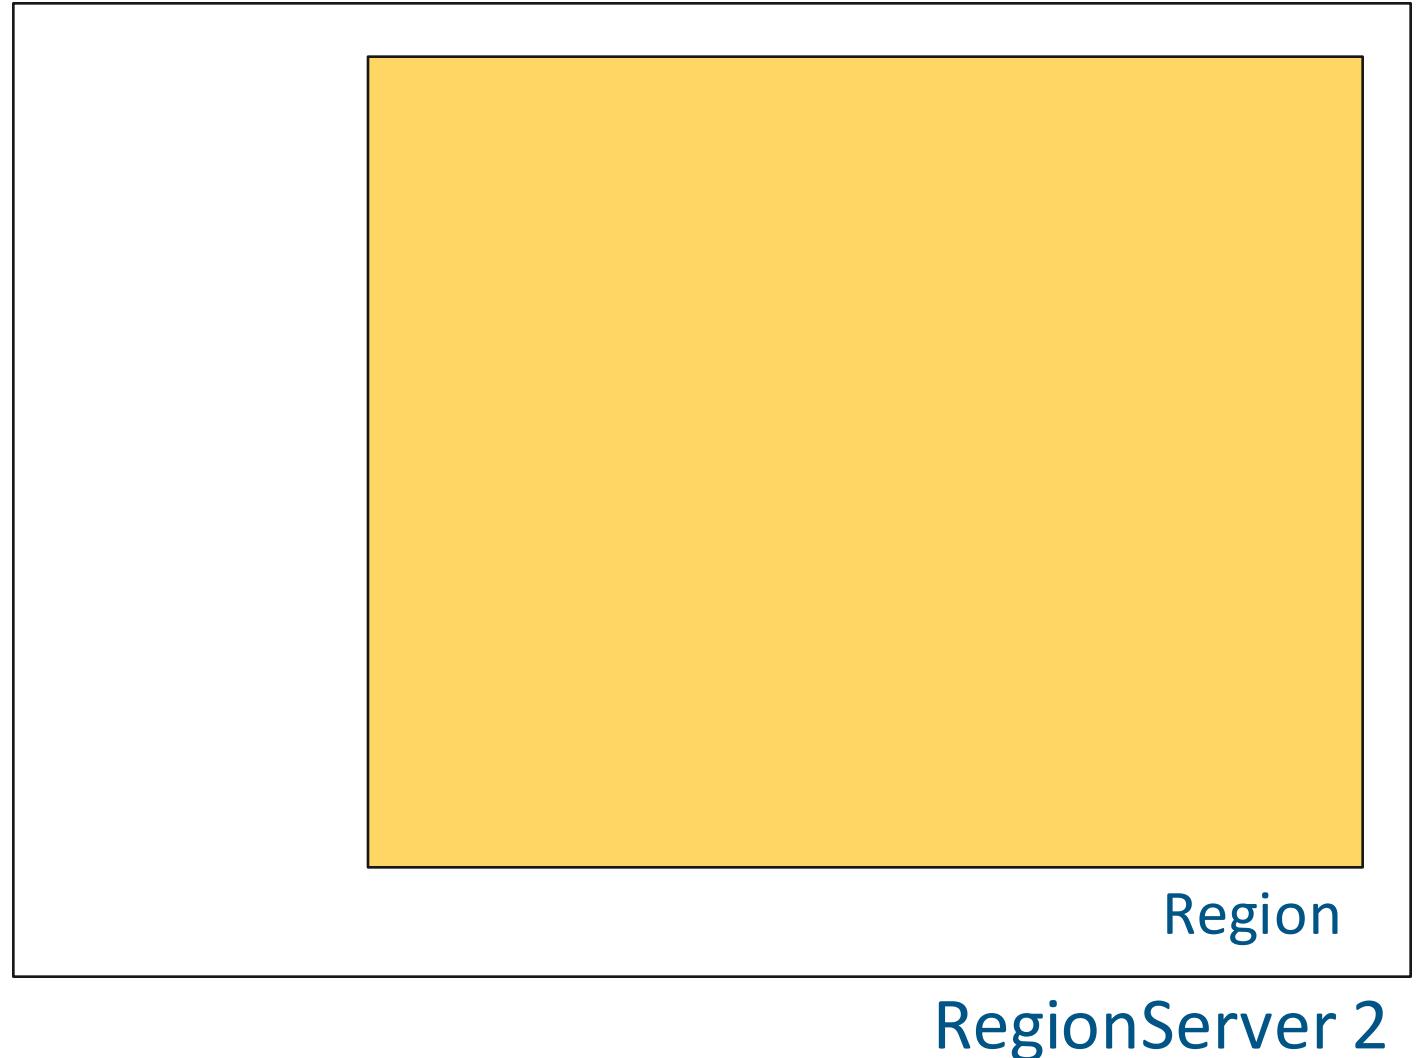
HBase Architecture | Compactions

- Minor compactions
 - Controlled by policy (pluggable).
- Major compactions
 - Automatic (by time) or manually triggered.
 - Tend to be run during off-peak times.

HBase Architecture | Splits

- Eventually, Regions become imbalanced.
 - Some grow to be huge, others remain small.
 - Leads to disparate load across RegionServers.
- In these cases, HBase can split a Region into two.
- Each Region is then available to be moved to a different RegionServer, if necessary.

HBase Architecture | Splits



HBase Architecture | Splits

RegionServer 3: Yeah!
Pick me!

Master: RegionServer 2
is really busy... Maybe
another RegionServer
can handle one of its
Regions?

Region

Region

RegionServer 2

HBase Architecture | APIs

- Conventional write path can be accessed through multiple APIs:
 - Java API
 - Most full-featured.
 - REST API
 - Easily accessible.
 - Thrift API
 - Support for many languages (e.g. C, C++, Perl, Ruby, Python).

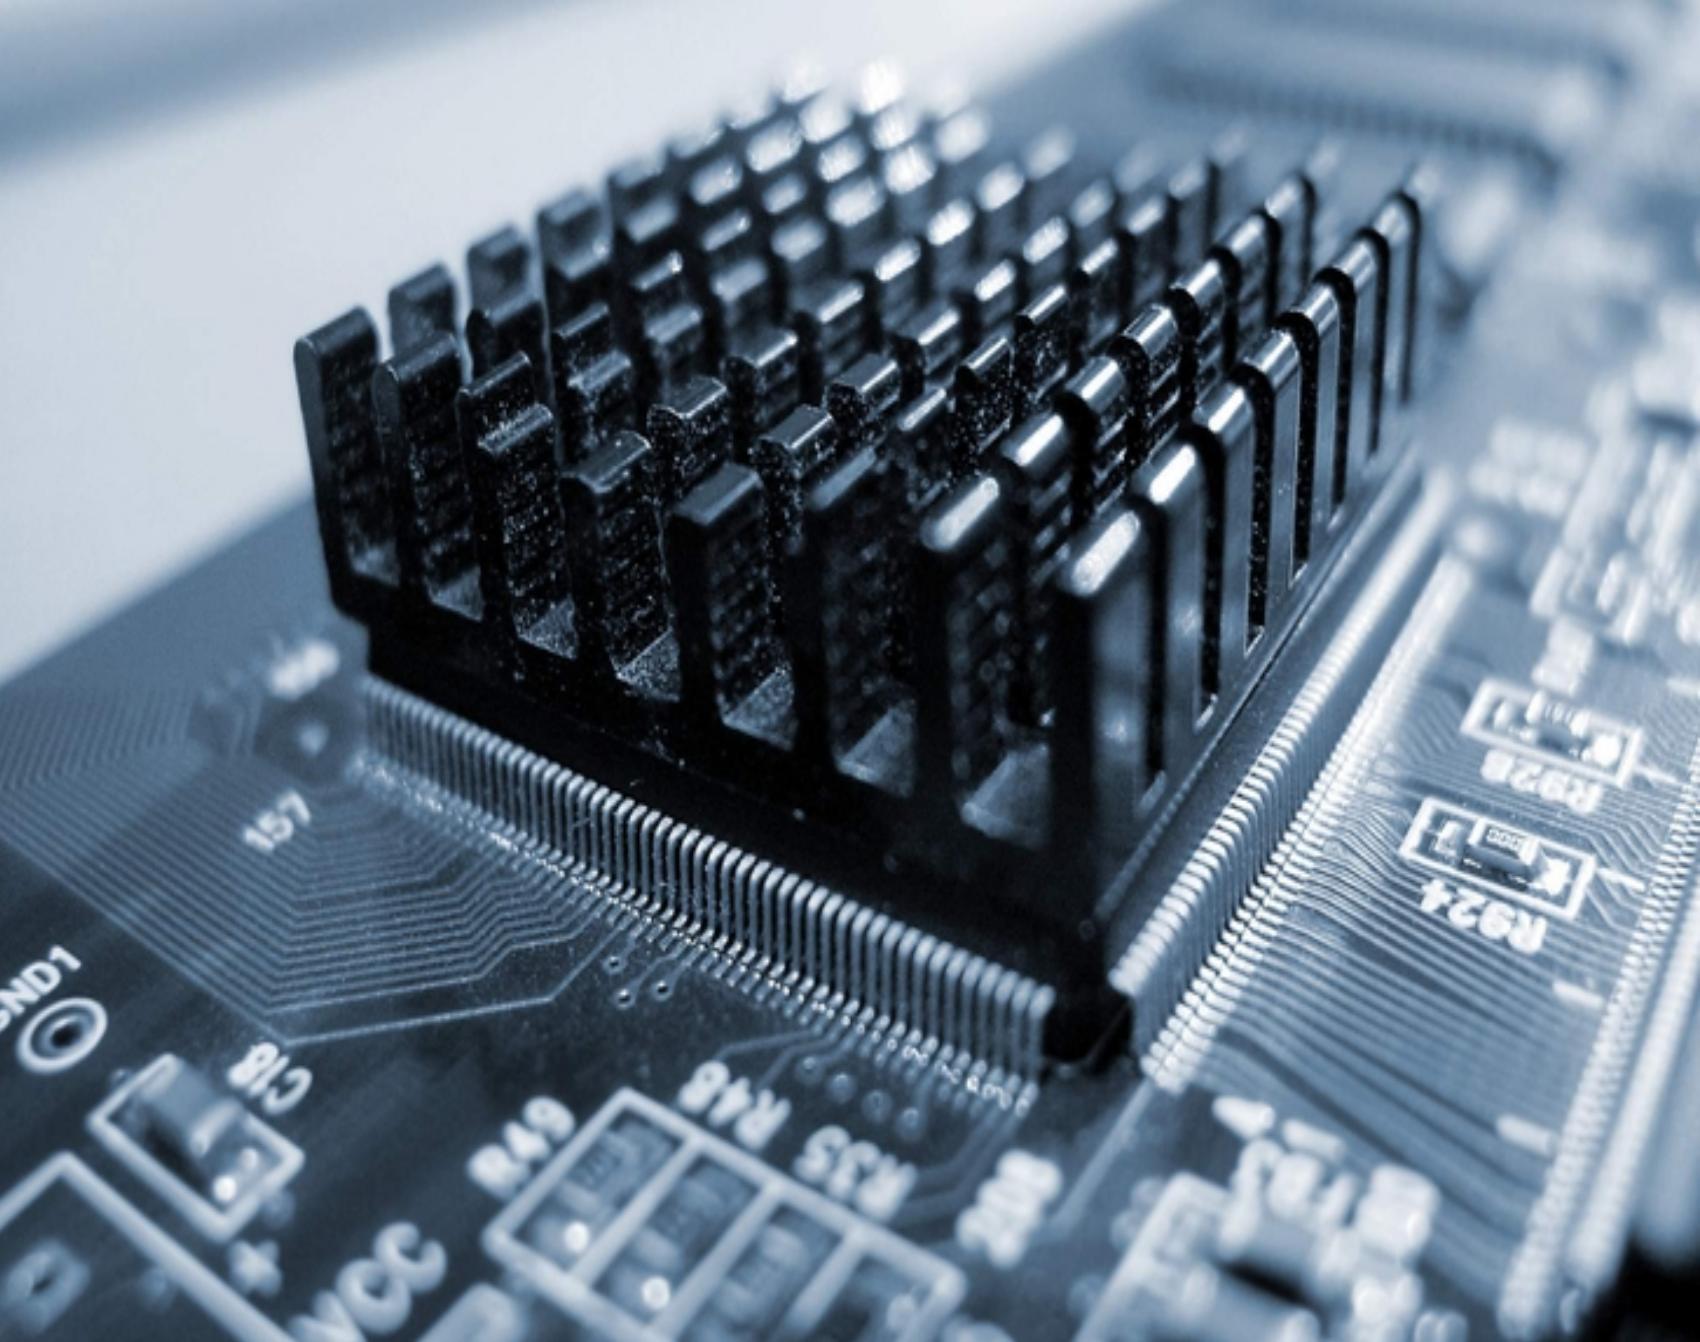
HBase Architecture | APIs

- This write path is durable, but if you're importing a lot of data, it can be problematic...
 - Every put goes into WAL, which means disk seeks. Lots of puts mean lots of disk seeks.
 - Lots of data into MemStores means lots of flushing to disk.
 - Lots of flushing to disk might mean lots of compactions.

HBase Architecture | Bulk Loading

- Bypass conventional write path.
 - Extract data from source.
 - Transform data into HFiles (done with MapReduce job) directly.
 - Tell RegionServers to serve these HFiles.

Enough of Architecture



What's up next, Doc?



- Break
- What have we learned from the users
- How can you benefit from that information



cloudera
Break

Apache HBase “Nascar” Slide



Apache HBase “Nascar” Slide



Apache HBase “Nascar” Slide



Apache HBase “Nascar” Slide



Apache HBase “Nascar” Slide



Apache HBase “Nascar” Slide



Apache HBase “Nascar” Slide





What have we learned from
all these users?



There are some patterns which repeat often.

Just like a lego block, maybe you can fit one directly in your system!



Know your ...

Data

- Entity Data
- Time-centric Event Data

Use of data

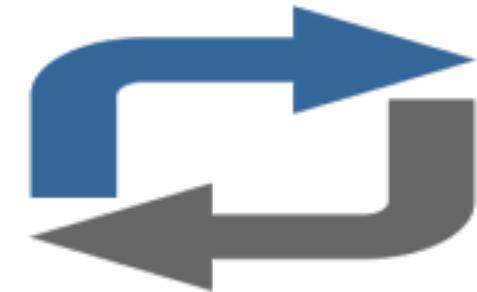
- Operational
- Analytical

How it goes in and out

- Real-time vs Batch
- Random vs Sequential



USE



Know your data ...

There are primarily two kinds of big data workloads. They have different storage requirements.

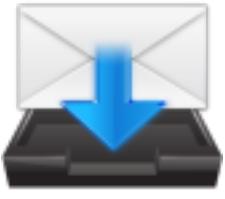
- Entity centric data
- Time centric event data



Entity centric data



Users



Accounts



Location



Clicks and
Metrics



Sensor Data

- Scales up with # of entities
- Billions of distinct entities

Time centric event data



Stock Ticker Data

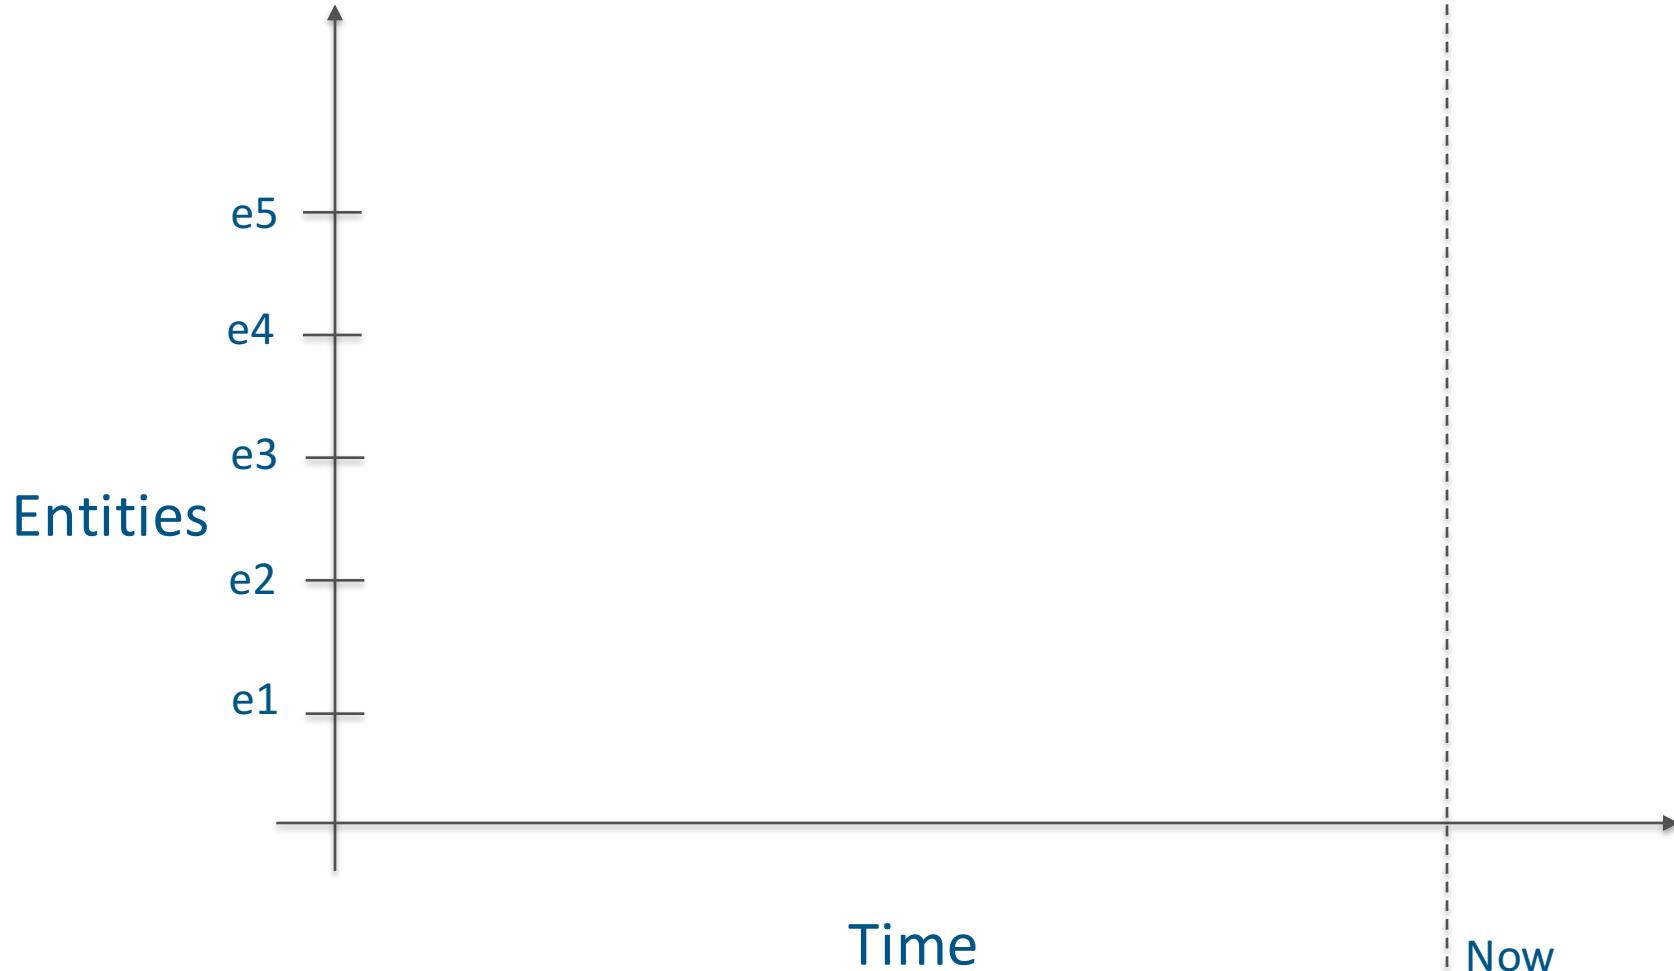


Monitoring applications

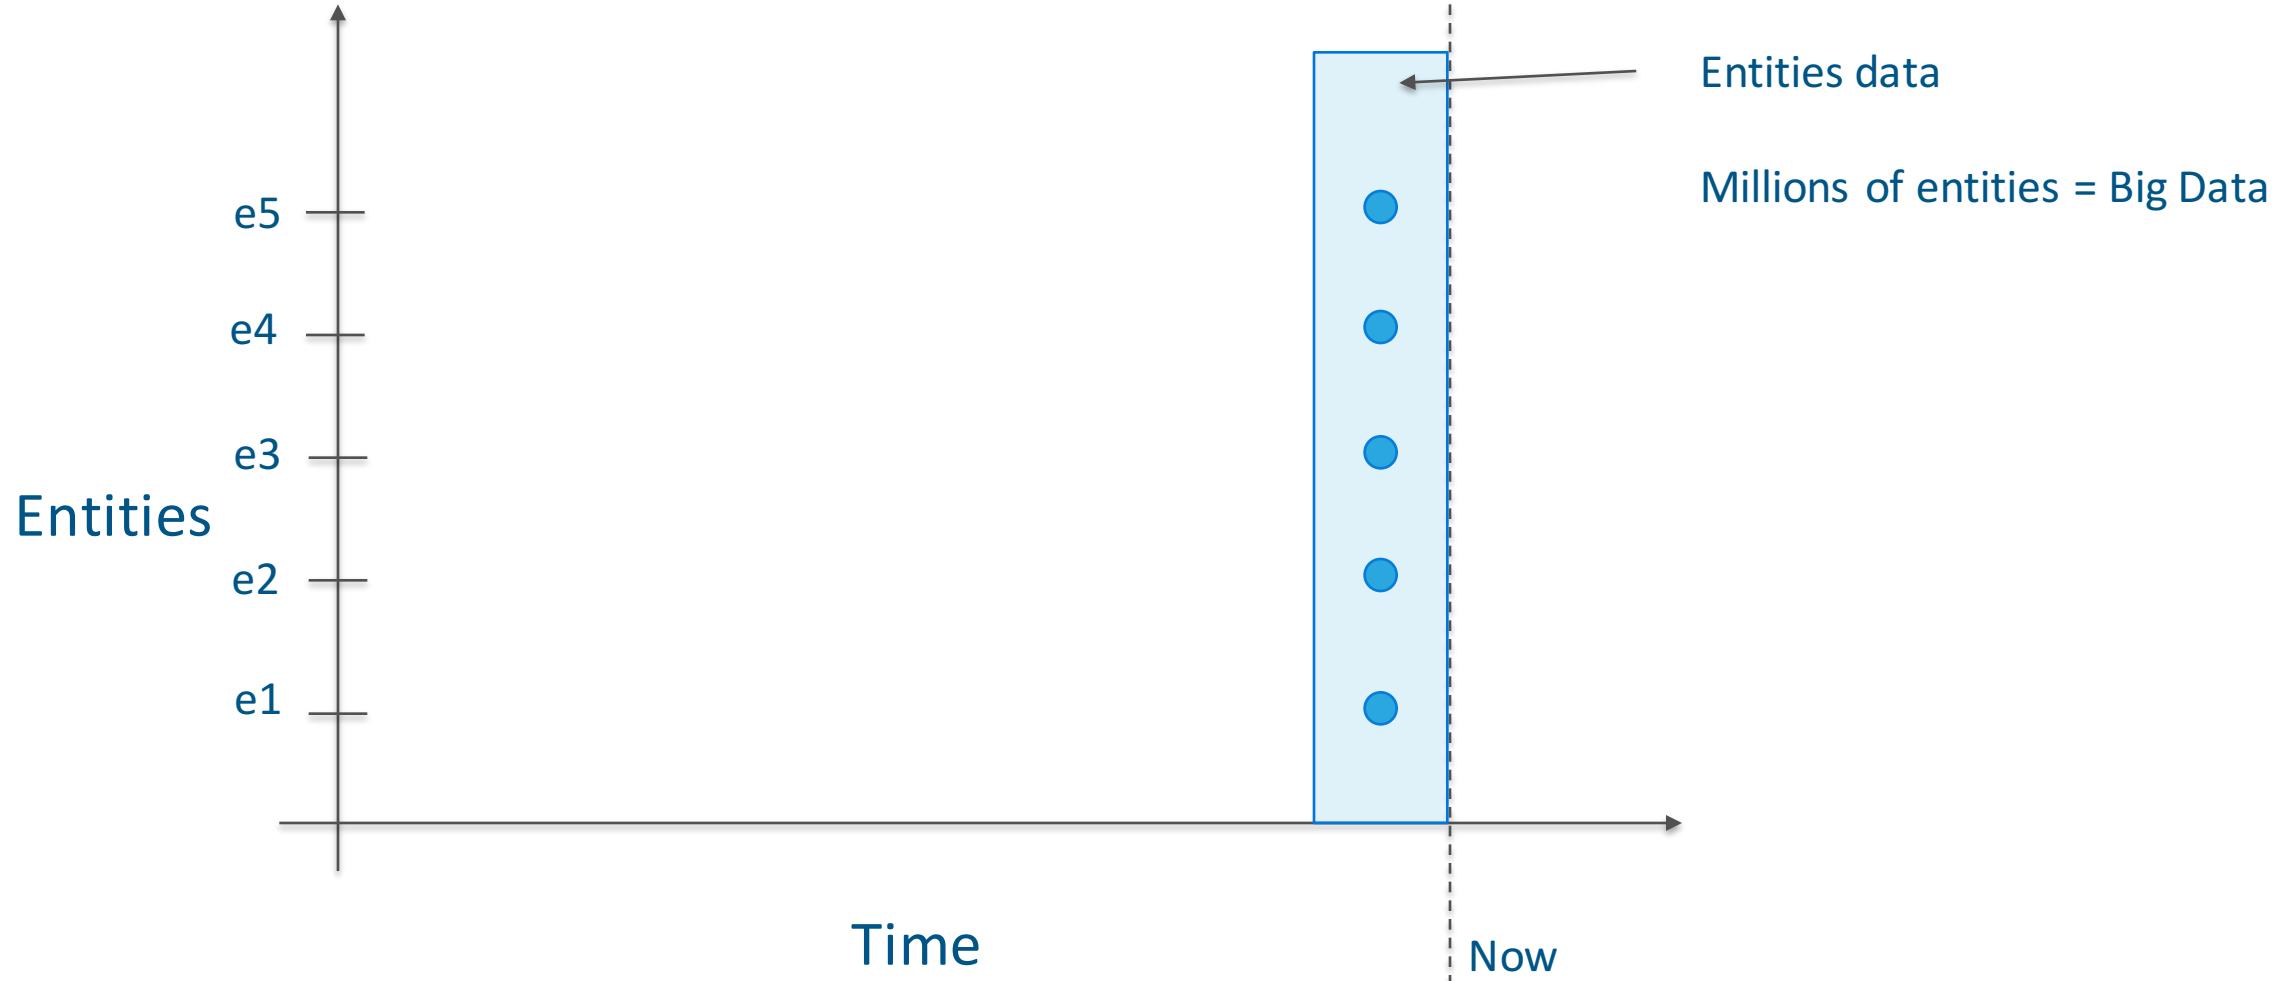


Periodic Sensor Data

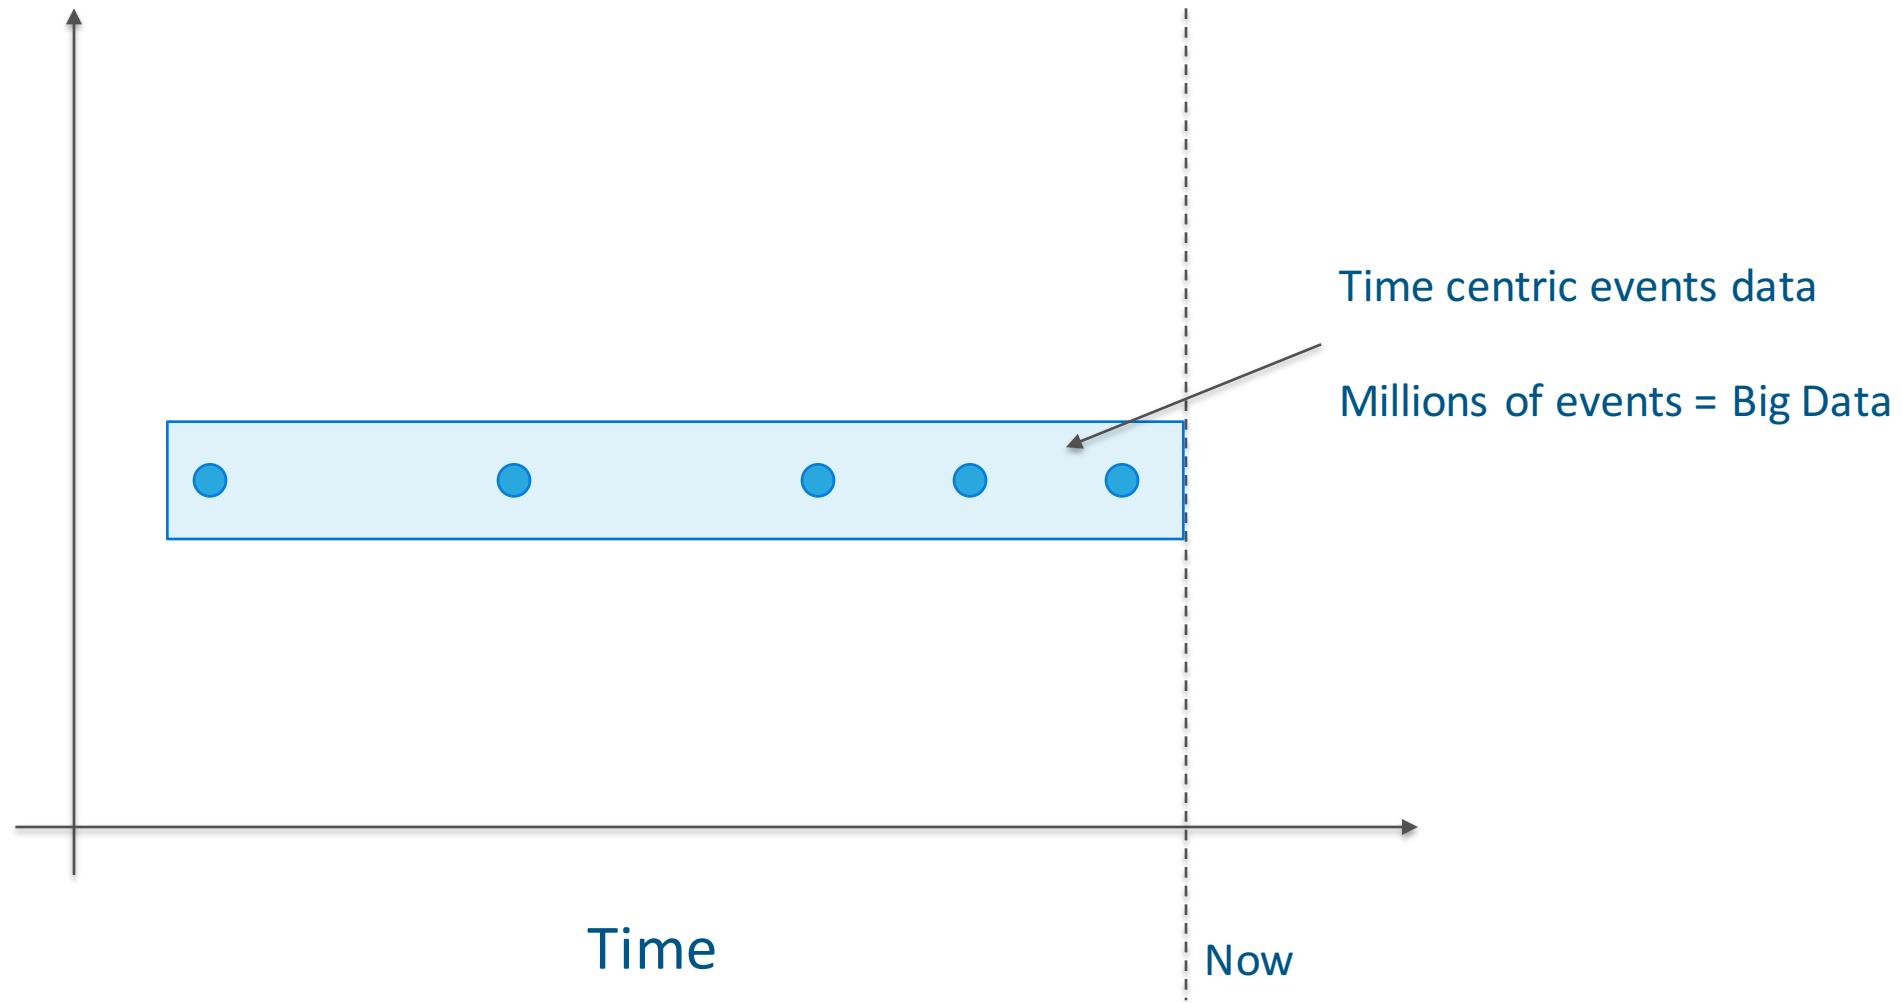
- Time-series data points over a period
- Scales up due to finer grained intervals, retention policies, and the passage of time



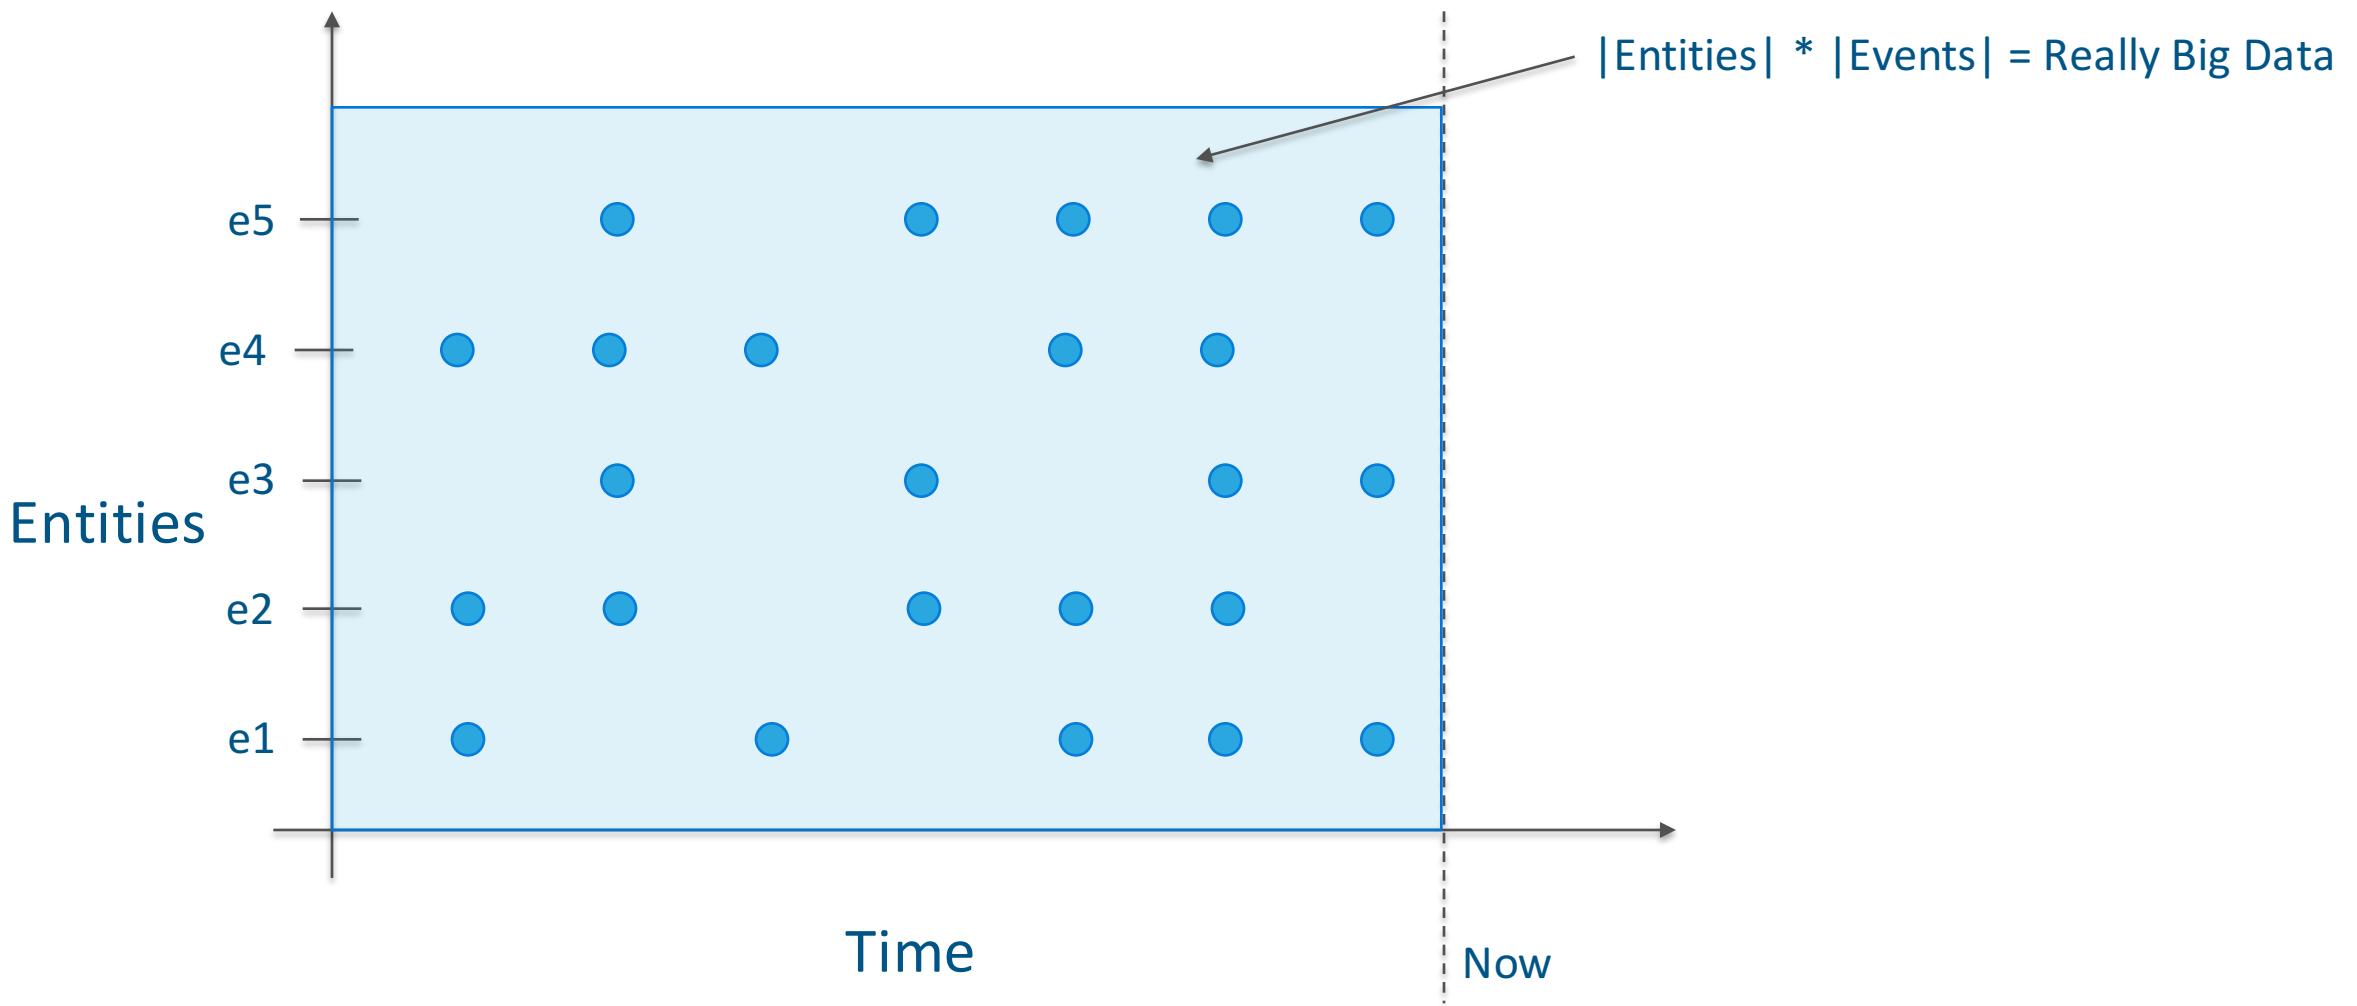
Entities data



Time-centric events data



Time-centric events about Entities



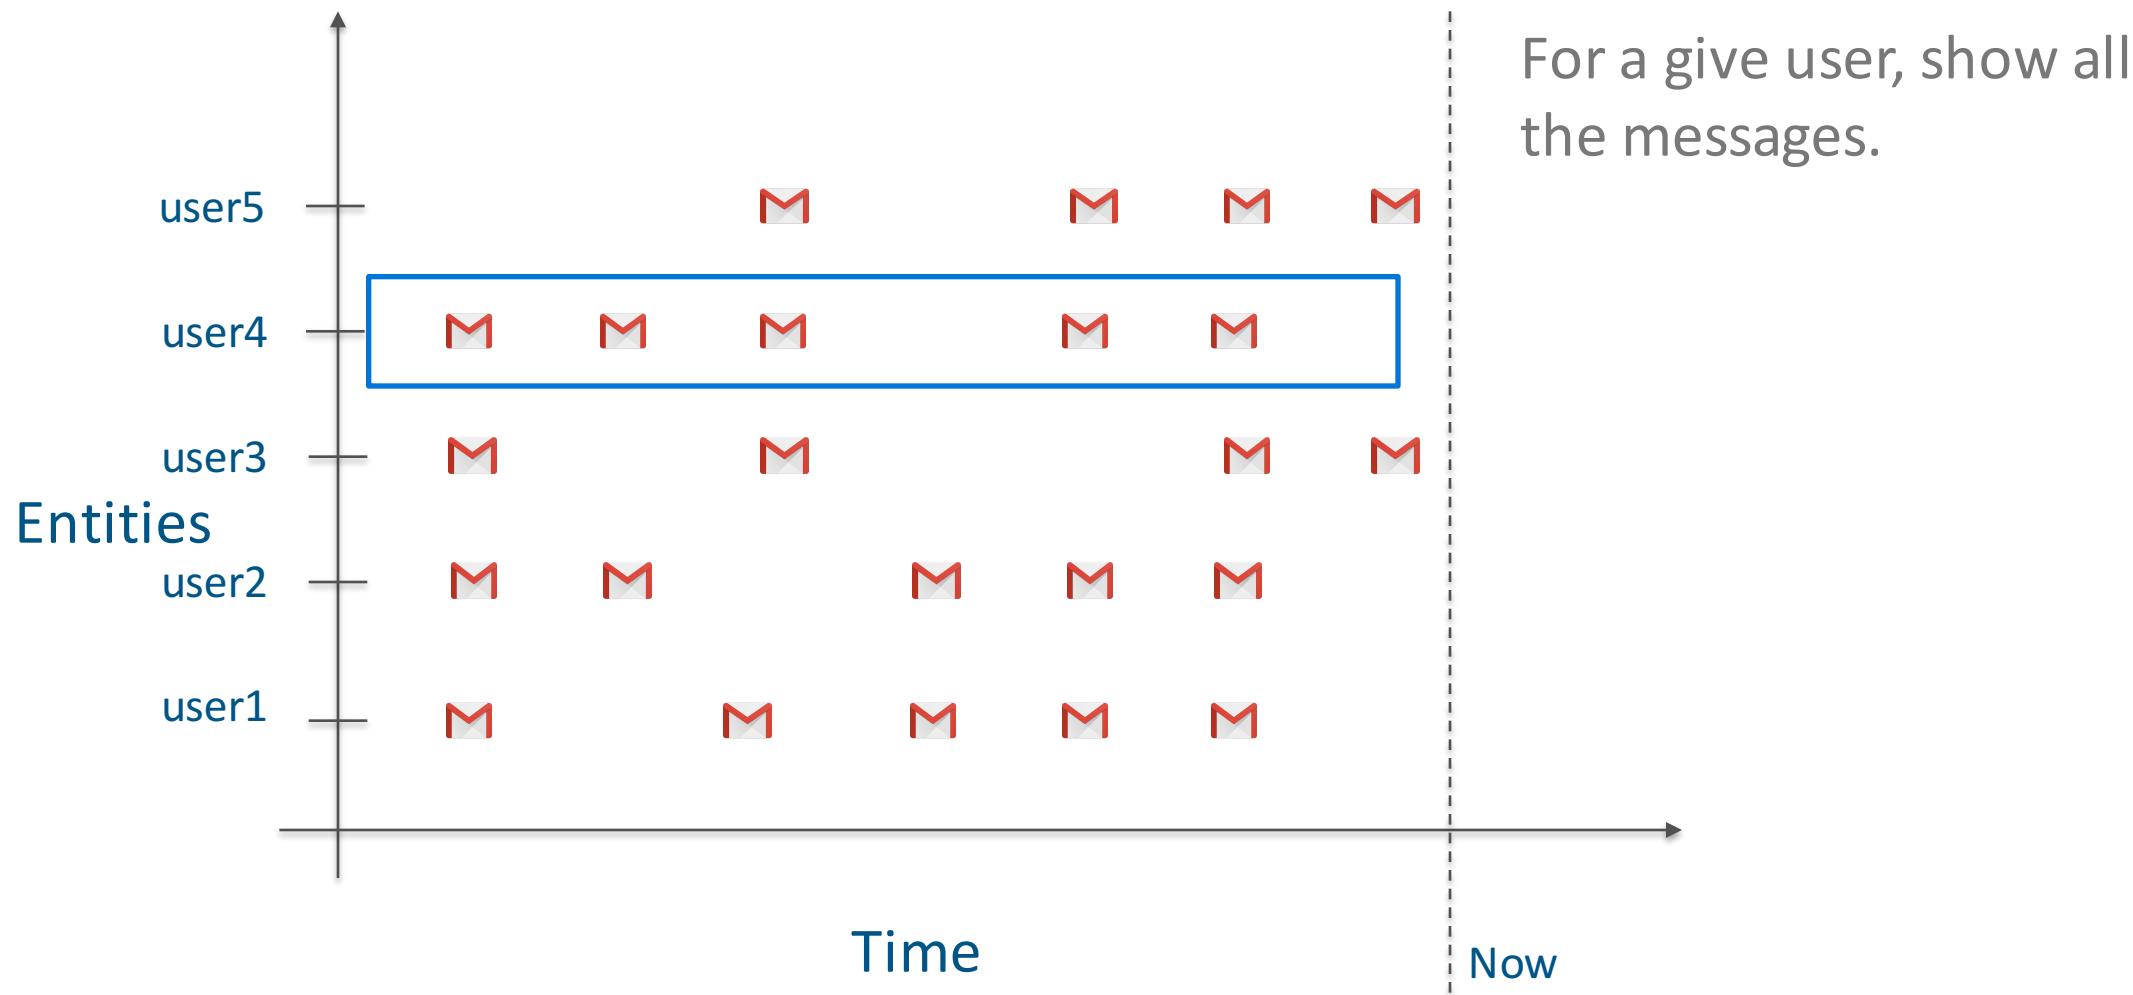
What questions do you ask?

- Do you focus in on entity first?

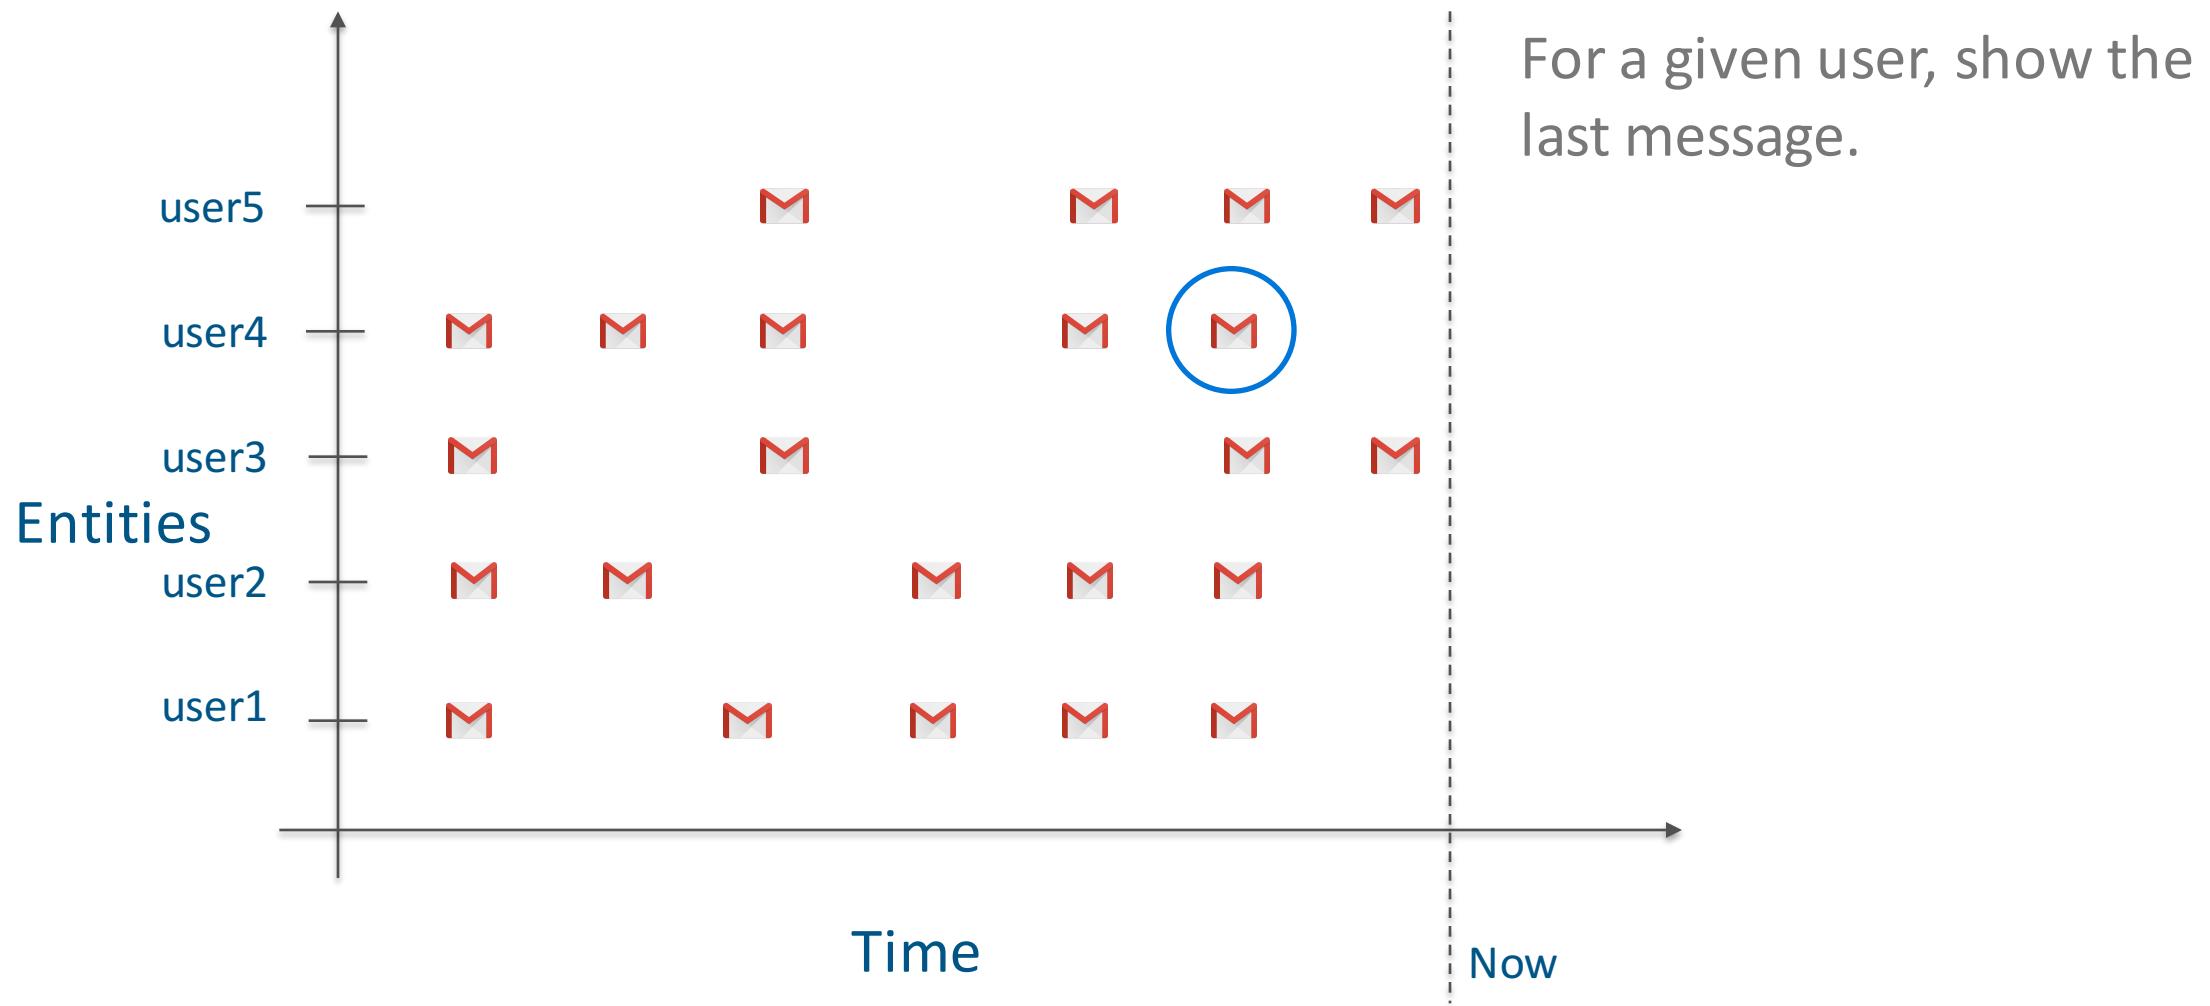
OR

- Do you focus in on time ranges first?
- Your answer will help you determine where and how to store your data.

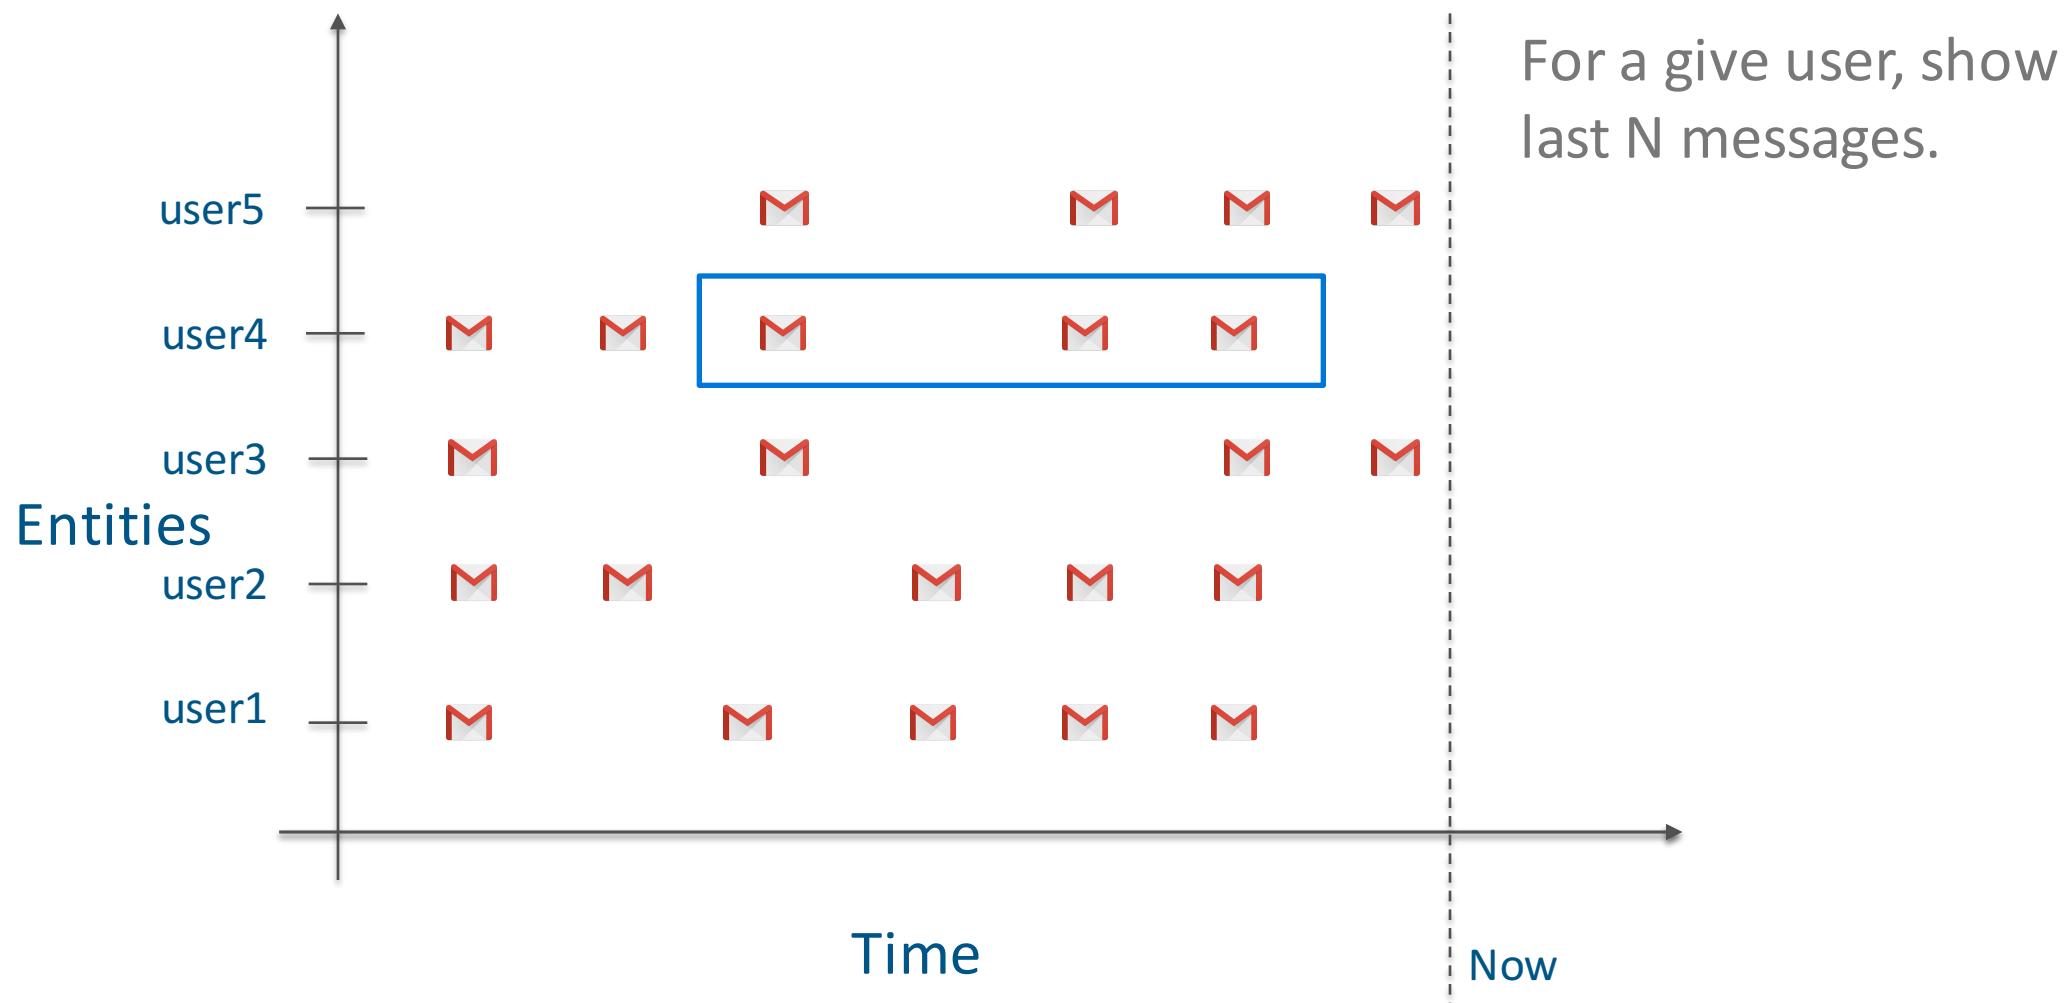
Entity first questions...



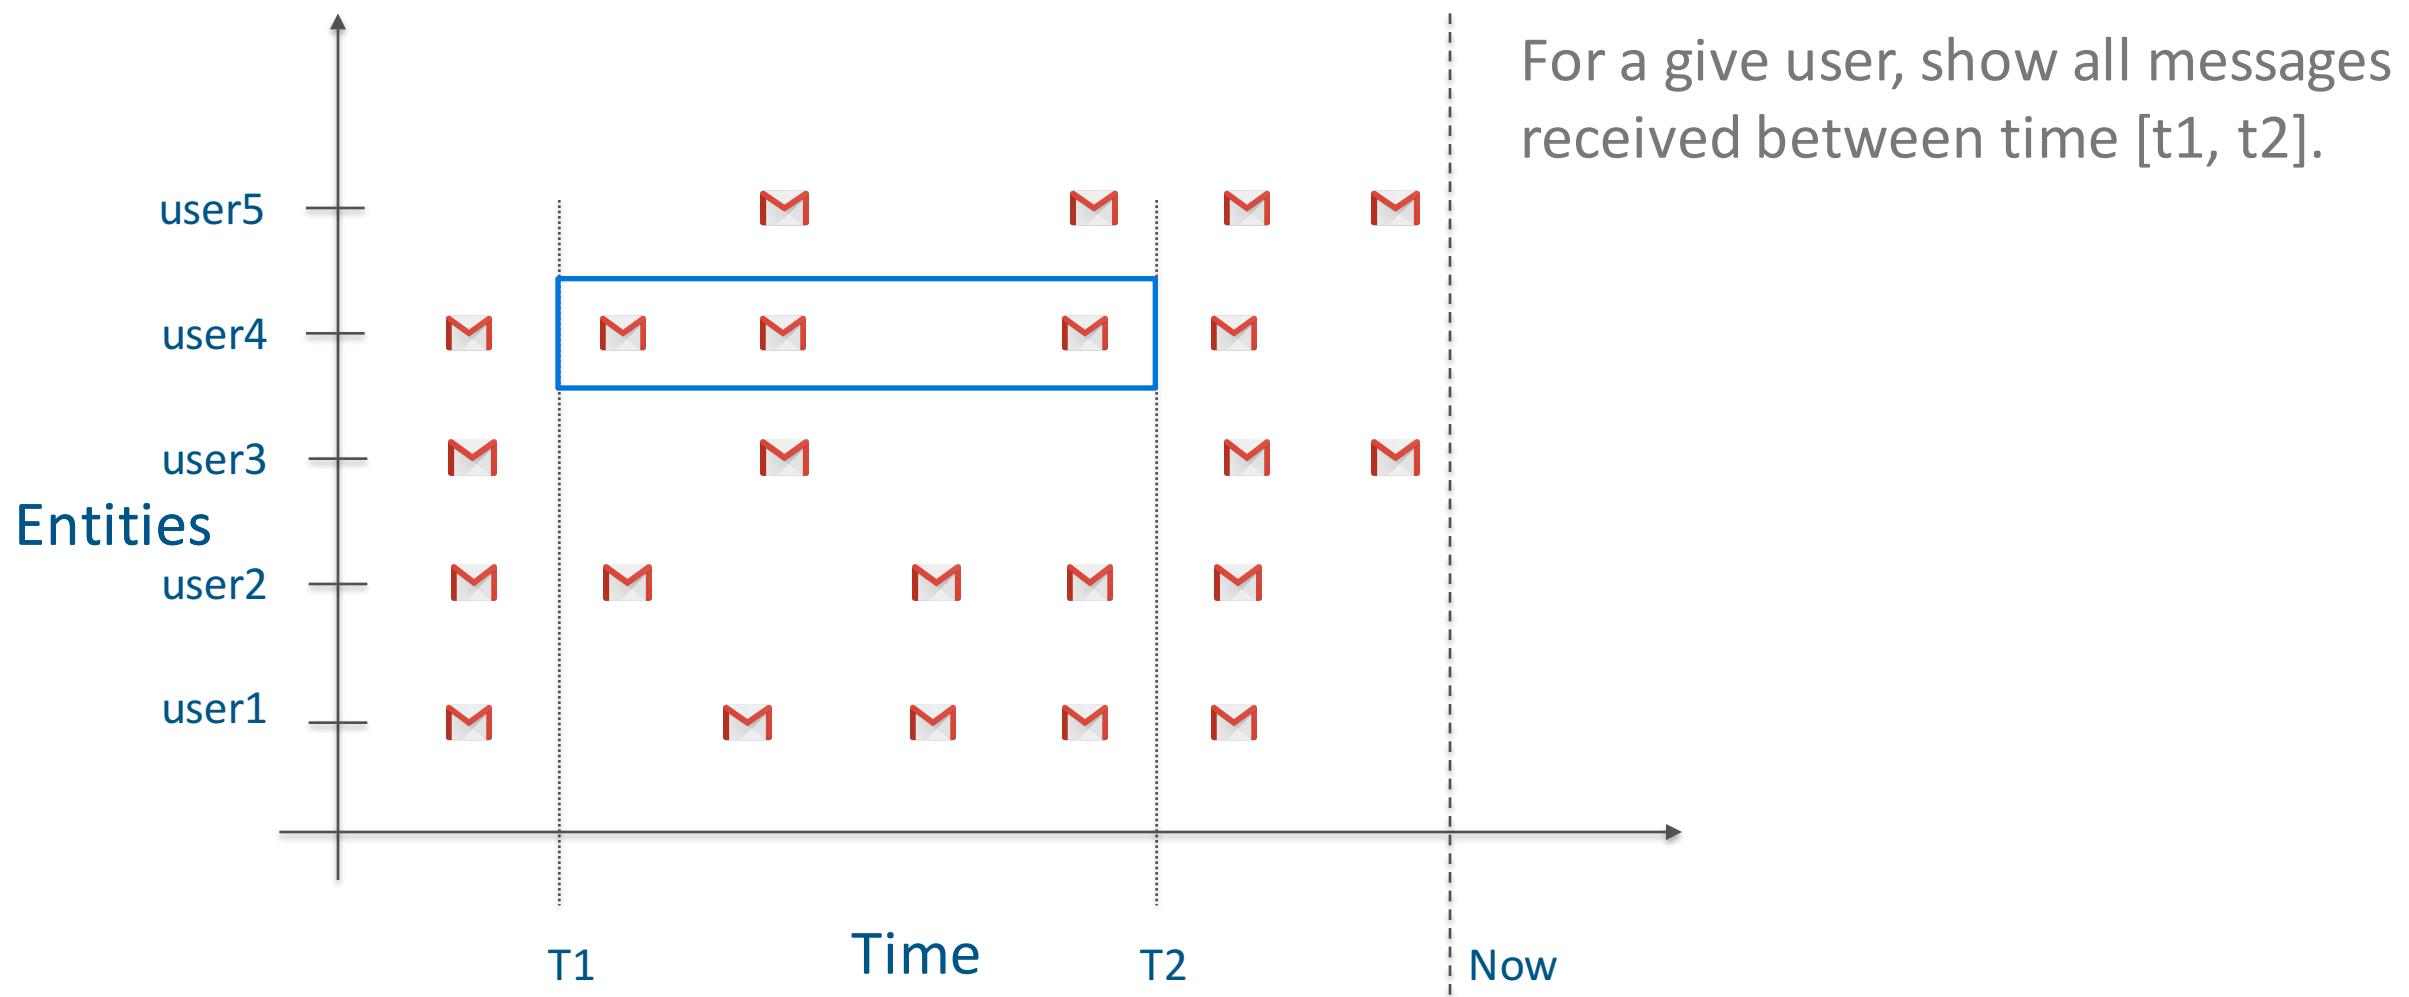
Entity first questions...



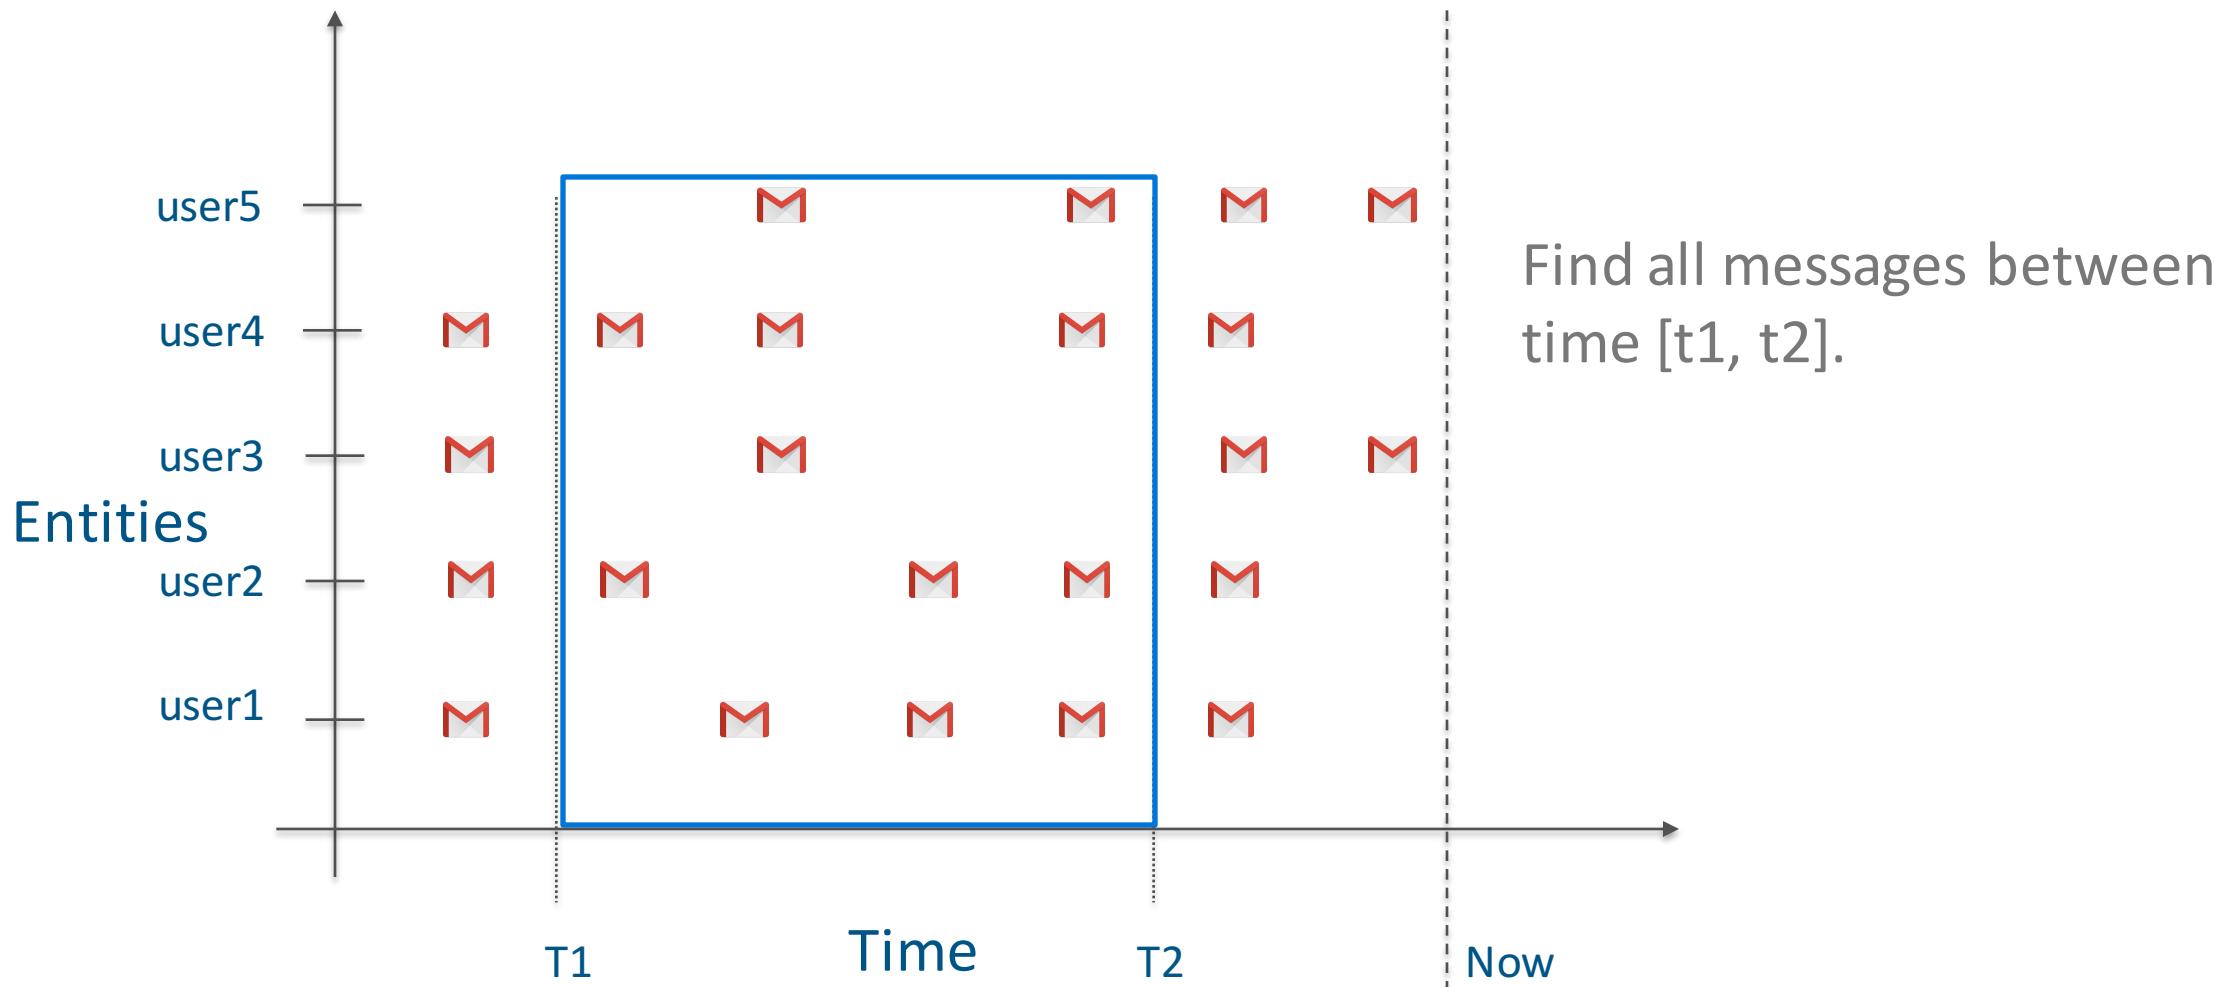
Entity first questions...



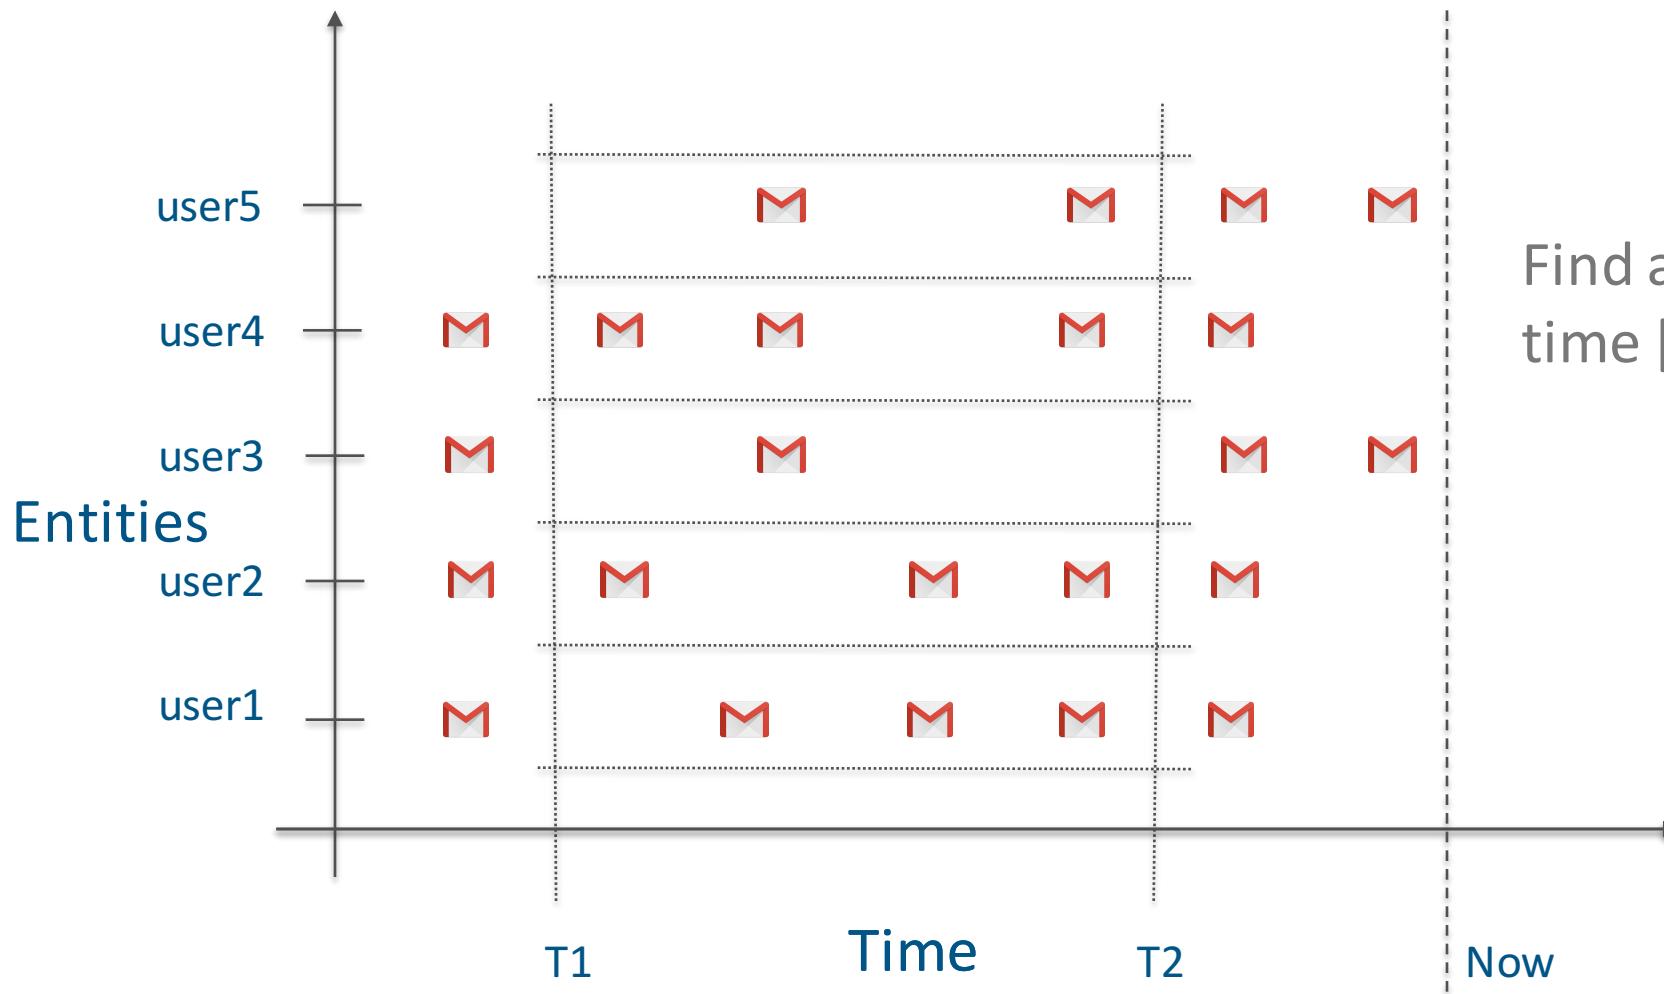
Entity first questions...



Time centric event first questions...

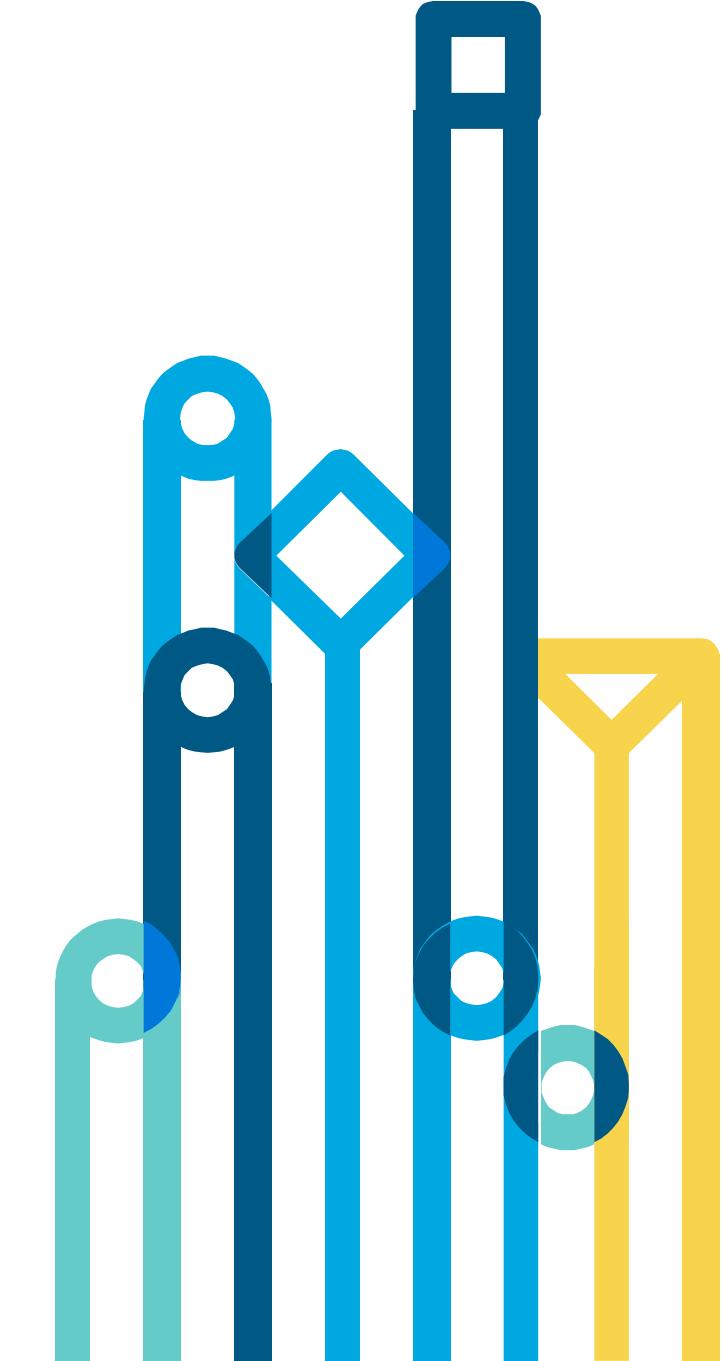


Time centric event first questions...



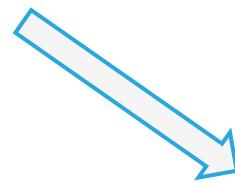
Find all messages between
time $[t_1, t_2]$ for all users.

How does the data get in and out
of HBase?

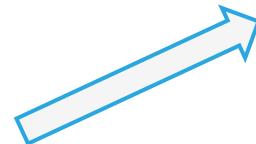


Getting data in...

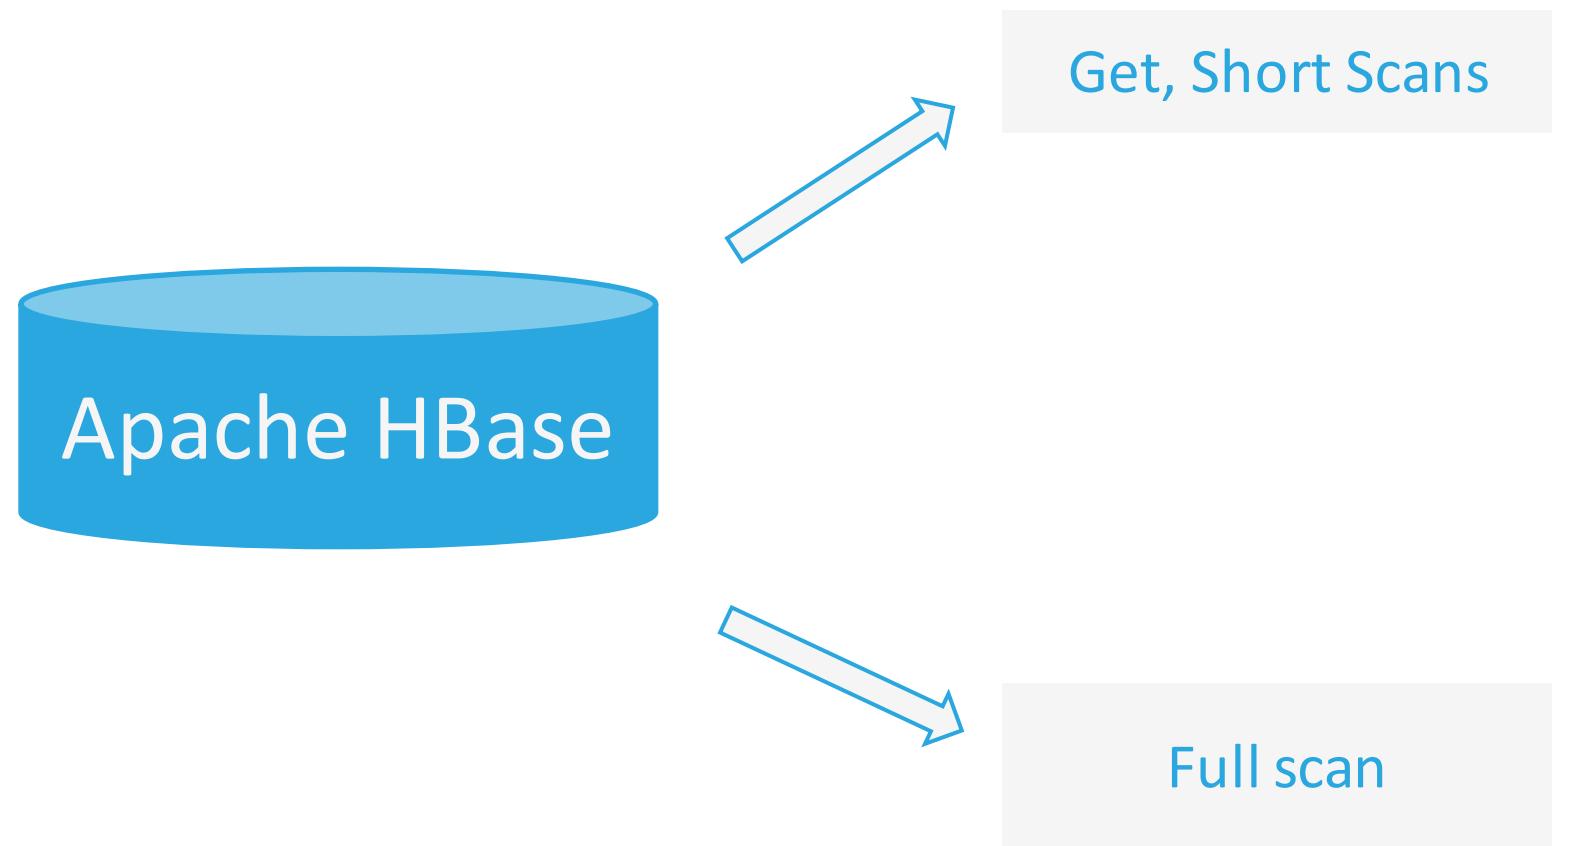
Put, Incr, Append



Bulk Import



Getting data out...

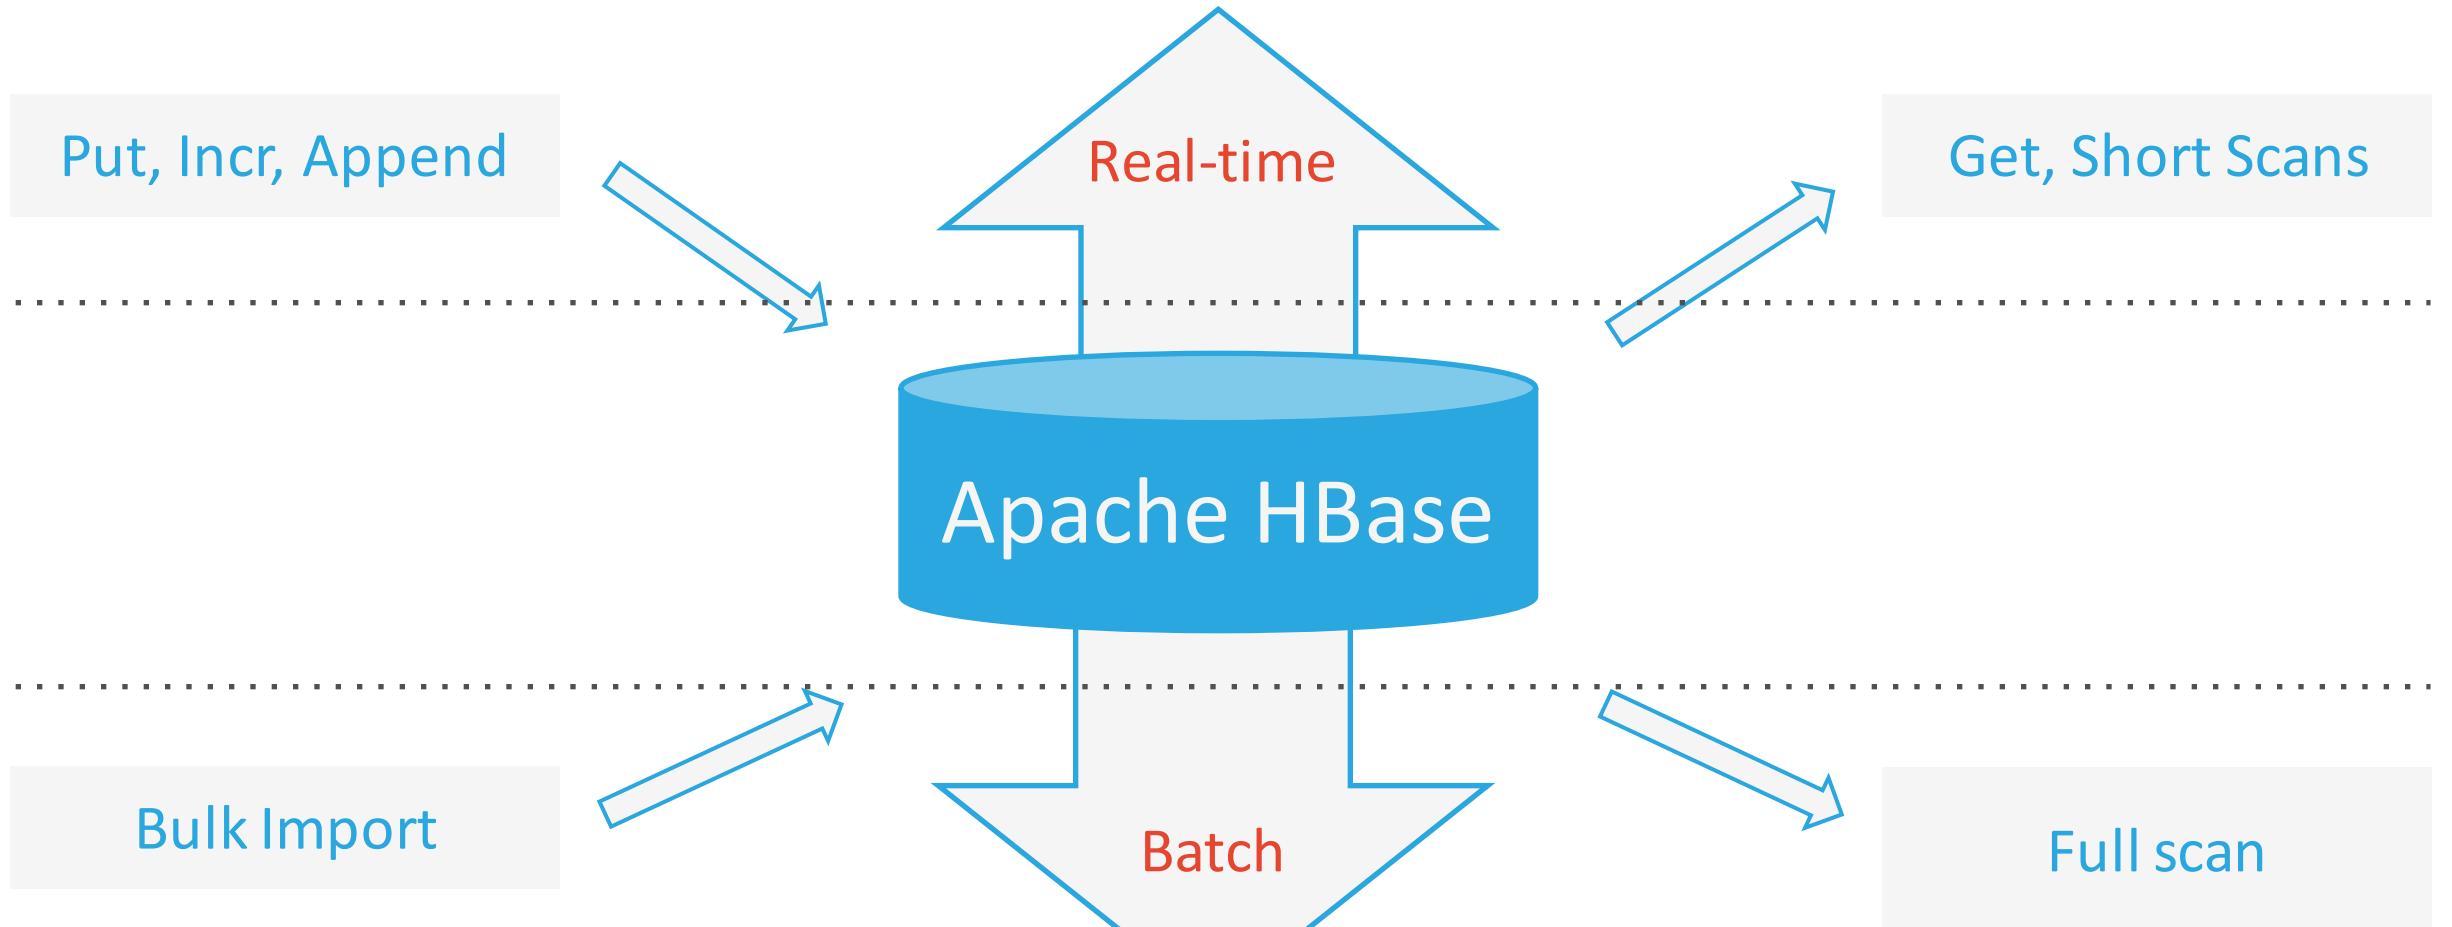


So, what's the best way?

Depends on your use case

Bottom-line: Disk I/O takes times.

- Limited disk read-write heads in a cluster
- Use the I/O bandwidth of your cluster efficiently

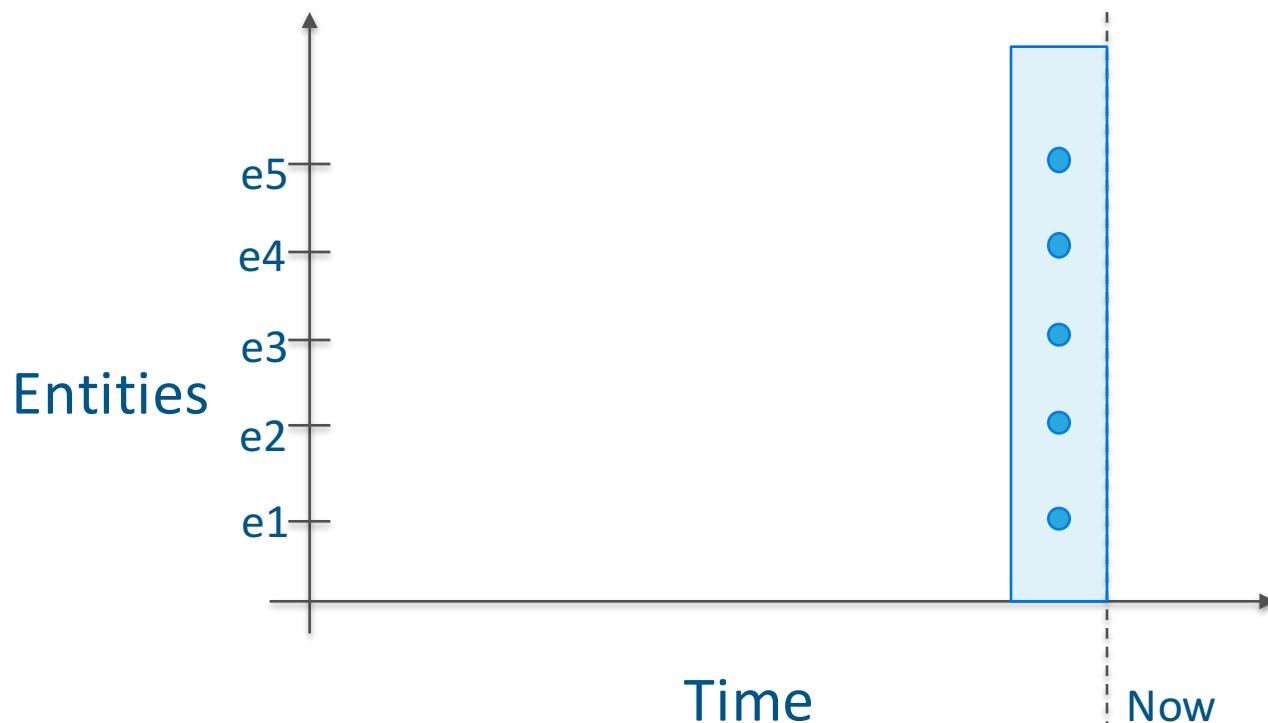


Let's dive into use case ...



Simple Entities

- Purely entity data, no relation between entities
 - Often from many different sources
 - Could be a well-done de-normalized RDBMS port



Simple Entities : Schema

- Row per entity
- Row key => entity ID, or hash of entity ID
- Column => Property / field, possibly timestamp

Simple Entities : Example

OCLC : Online Computer Library Center

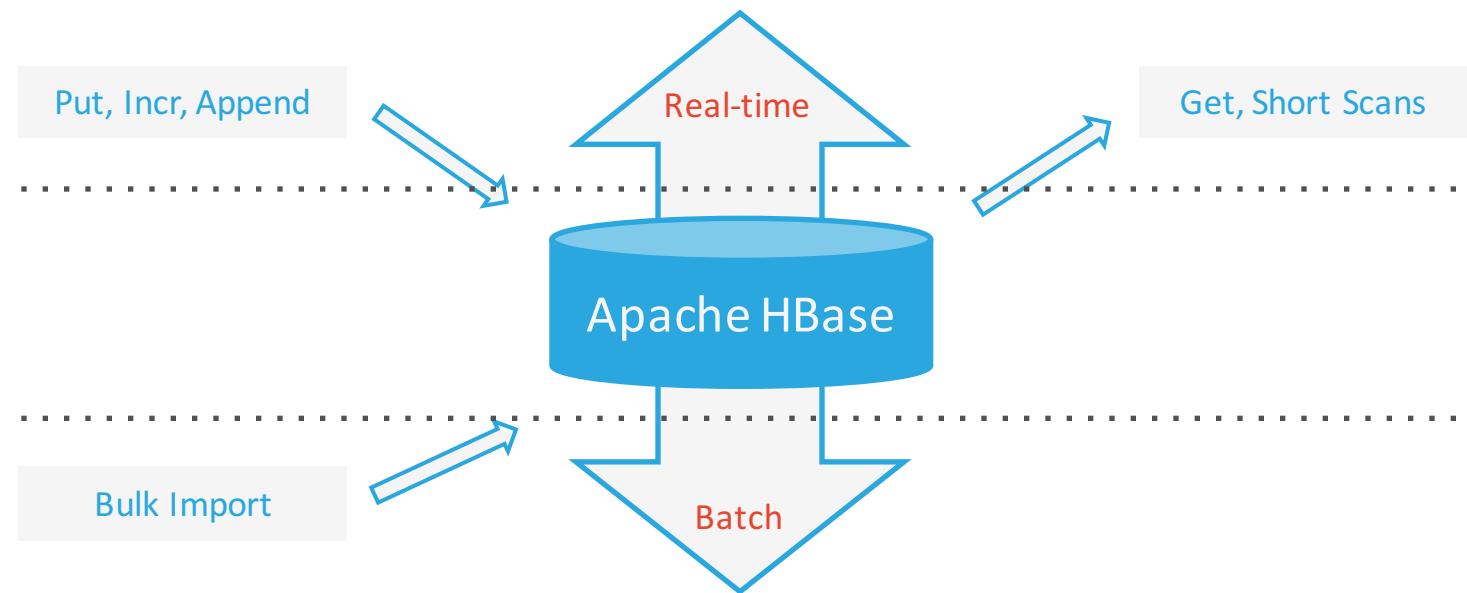


Workloads:

- Lookup books → Real time read
- Add new book one at a time, update information about existing books, issue books → Real-time write
- New library joins the group, import its data → Batch write

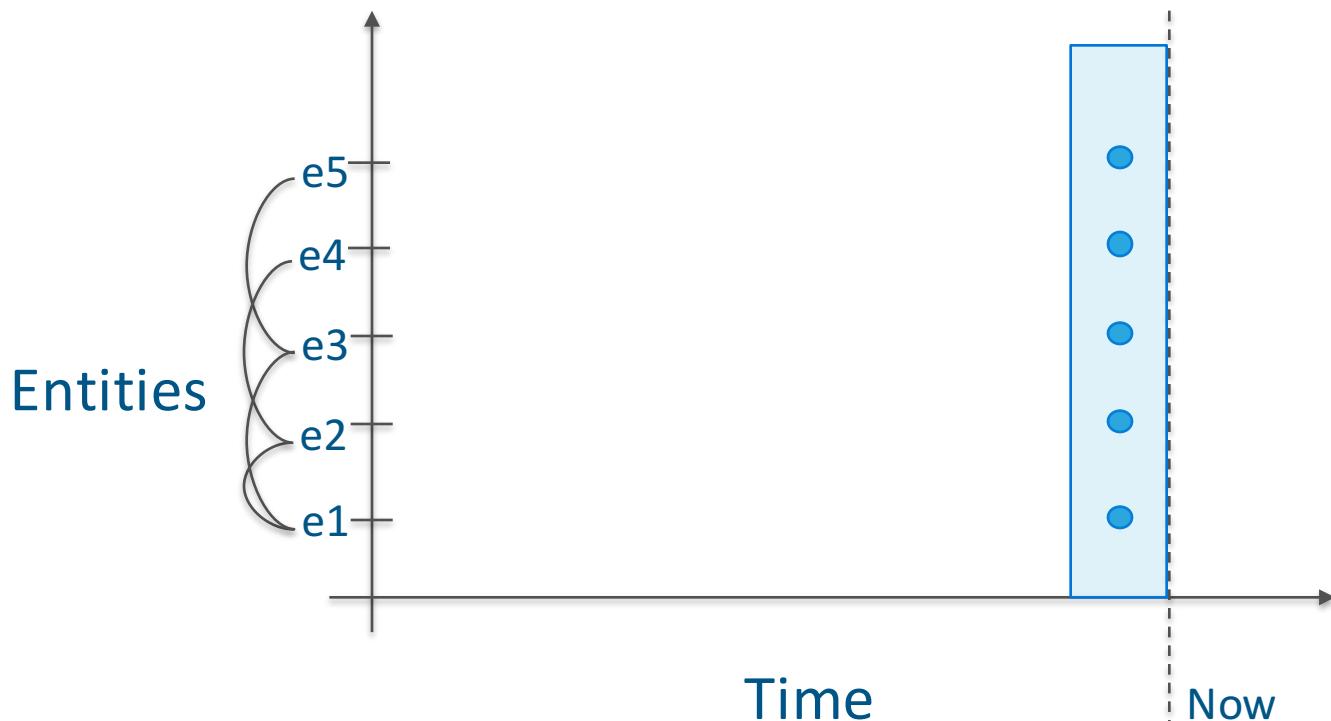
Simple Entities : Access Pattern

- Access Patterns
 - Writes : Batch / Real-time
 - Reads: Real-time



Linked Entities (Graph Data)

- Entity are linked to form a graph



Linked Entities (Graph Data) : Schema

- Row per Node (Entity)
- Row key => Node ID (Entity ID)
- Column => “Relationship:OtherNodeID”
- Value => Meta data about relationship

Linked Entities (Graph Data) : Example

Social Network (Facebook)

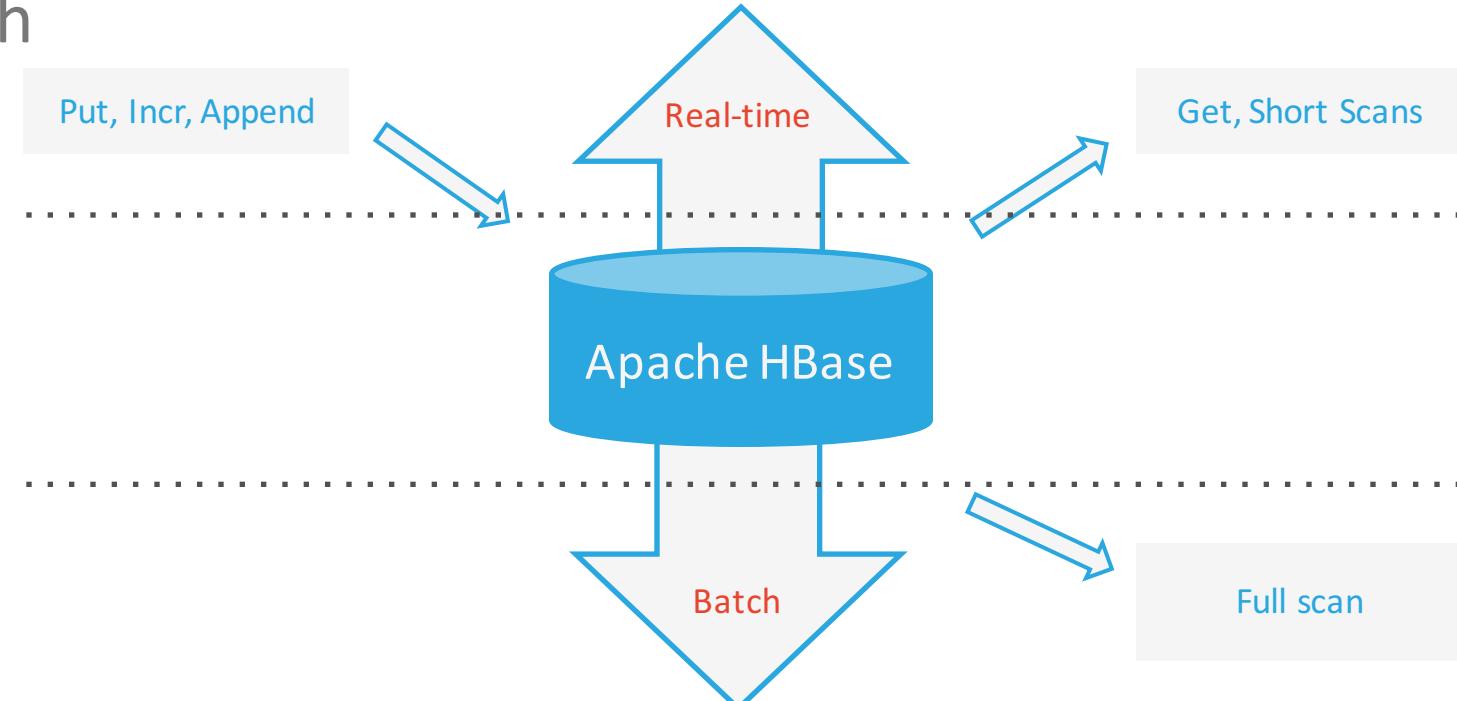


Workloads:

- Get any info about a user → Real time read
- Update any info about a user → Real time write
- Limited graph analysis (based on immediate friends) → Batch read

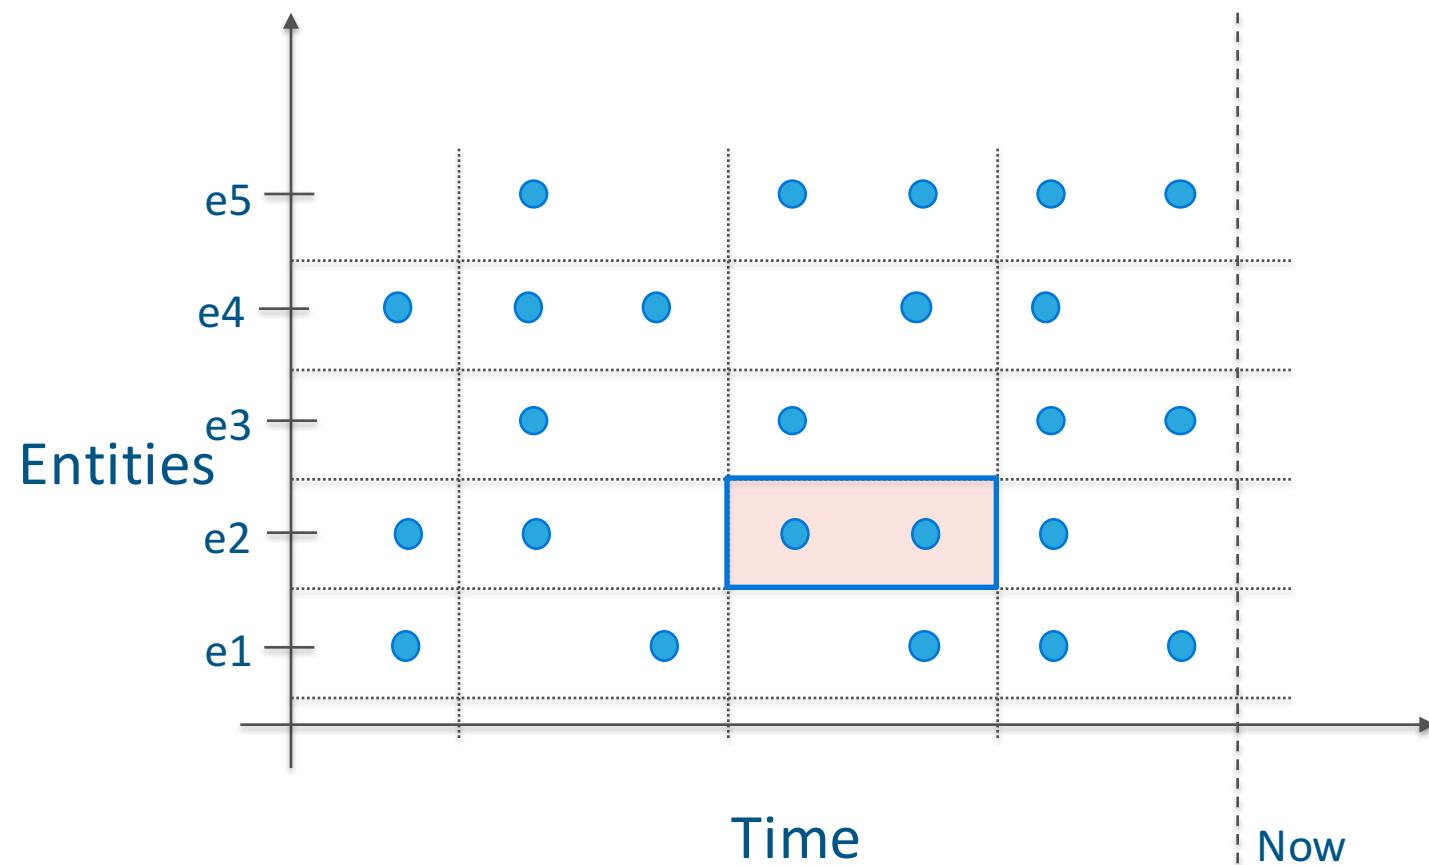
Linked Entities (Graph Data) : Access Pattern

- Access Patterns
 - Reads: Real-time or Batch
 - Writes: Real-time



Time-coupled entities

- Events about entities in time centric
- **Focus on entities first**



Time-coupled entities : Schema

- Row = Entity's events in a time slice
- Row key = Entity ID + (time / k)
- Column Qualifier = timestamp

Time-coupled entities: Example

Messaging service



Primary Workload

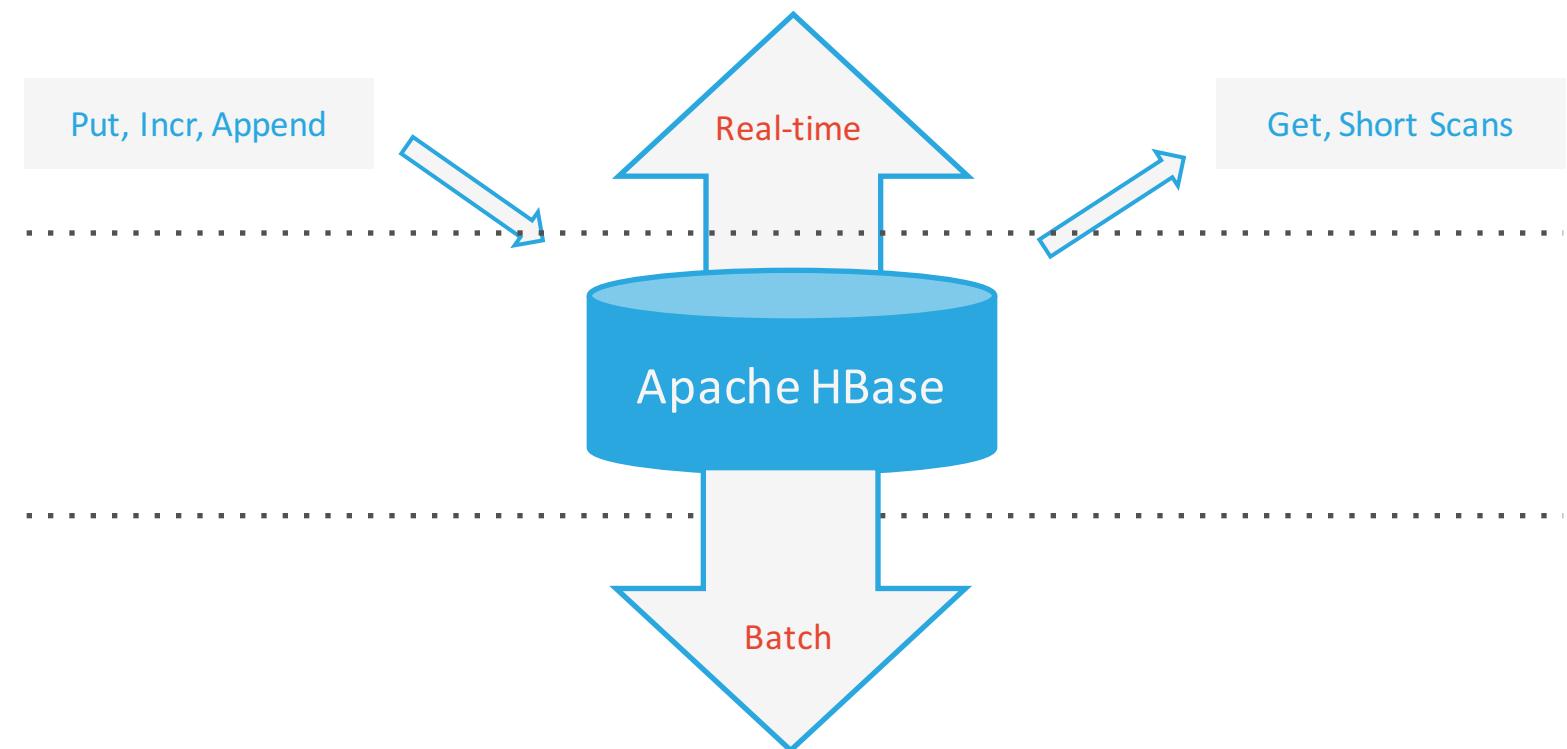
- Sending a message, update metadata (read, star, move, delete) →

Real-time write

- Reading a message, get last N messages → Real-time read

Time-coupled entities : Access Pattern

- Access Pattern
 - Writes: Real-time
 - Reads: Real-time



HBase is great!

But not for everything ...

Current HBase weak spots

- HBase architecture can handle a lot
 - Engineering tradeoffs optimize for some use cases
 - HBase can still do things it is not optimal for
 - Other systems are fundamentally more efficient for some workloads
- Just because it is not good today, doesn't mean it can't be better tomorrow!

A not so good use case: Large Blob Store

- Saving large objects >50 MB per cell
- Examples
 - Raw video storage in HBase
- Problems:
 - **Write amplification** when re-optimizing data for read (compactions on large unchanging data)
 - New: Medium Object (MOB) supported (lots of 100KB-10MB cells)

Another not good use case: Analytic archive

- Store data chronologically, time as primary index
 - Row key = timestamp
 - Real time writes
 - Column-centric aggregations over all rows
- Schema
 - Row key: timestamp
 - Column qualifiers: properties with data or counters
- Example
 - Machine logs organized **by timestamp** (causes write hot-spotting)

Summary

- HBase is used widely across industry
- Few patterns learnt from these users
- Understanding
 - Data : Entity and time-centric events
 - Questions you ask from your data
 - How does data gets in and out
- When not to use HBase



Scalable time series database

Time-Series

Data points for entities over time



OpenTSDB

- **Store trillions of data points**
- **Millisecond precision**
- **Keep raw data forever**
- **Scales to millions of writes per sec**
- **Generate graphs from GUI**

OpenTSDB

- **Store trillions of data points**
- **Millisecond precision**
- **Keep raw data forever**
- **Scales to millions of writes per sec**
- **Generate graphs from GUI**

OpenTSDB : Use Cases

- System Monitoring
 - Servers
 - Network
- Sensor Data
- Stock market data

OpenTSDB : Example

OVH

- Large cloud/hosting provider
- Monitor everything: networking, temperature, voltage, application performance, resource utilization, customer-facing metrics, etc.
- 35 servers, 100k writes/s, 25TB raw data



Yahoo!

- Monitoring application performance and statistics
- 15 servers, 280k writes/s

The Yahoo! logo in its signature purple color.

cloudera

Source: <http://www.slideshare.net/HBaseCon/ecosystem-session-6>

OpenTSDB : Datapoints

- In OpenTSDB, there are
 - Metric
 - Timestamp
 - Value
 - Tags (key-value pairs) : to identify the entity

OpenTSDB : Datapoints example

- E.g. 10 servers handling requests web requests
- Metric: num_requests_per_second
- Tags: “host=web-server-1”, “host=web-server-2”, and so on
- Example data points
 - num_requests_per_second 1439828251 50 host=web-server-1
 - num_requests_per_second 1439828251 72 host=web-server-2
 - num_requests_per_second 1439828252 30 host=web-server-3
 - ...so on

OpenTSDB : How it works

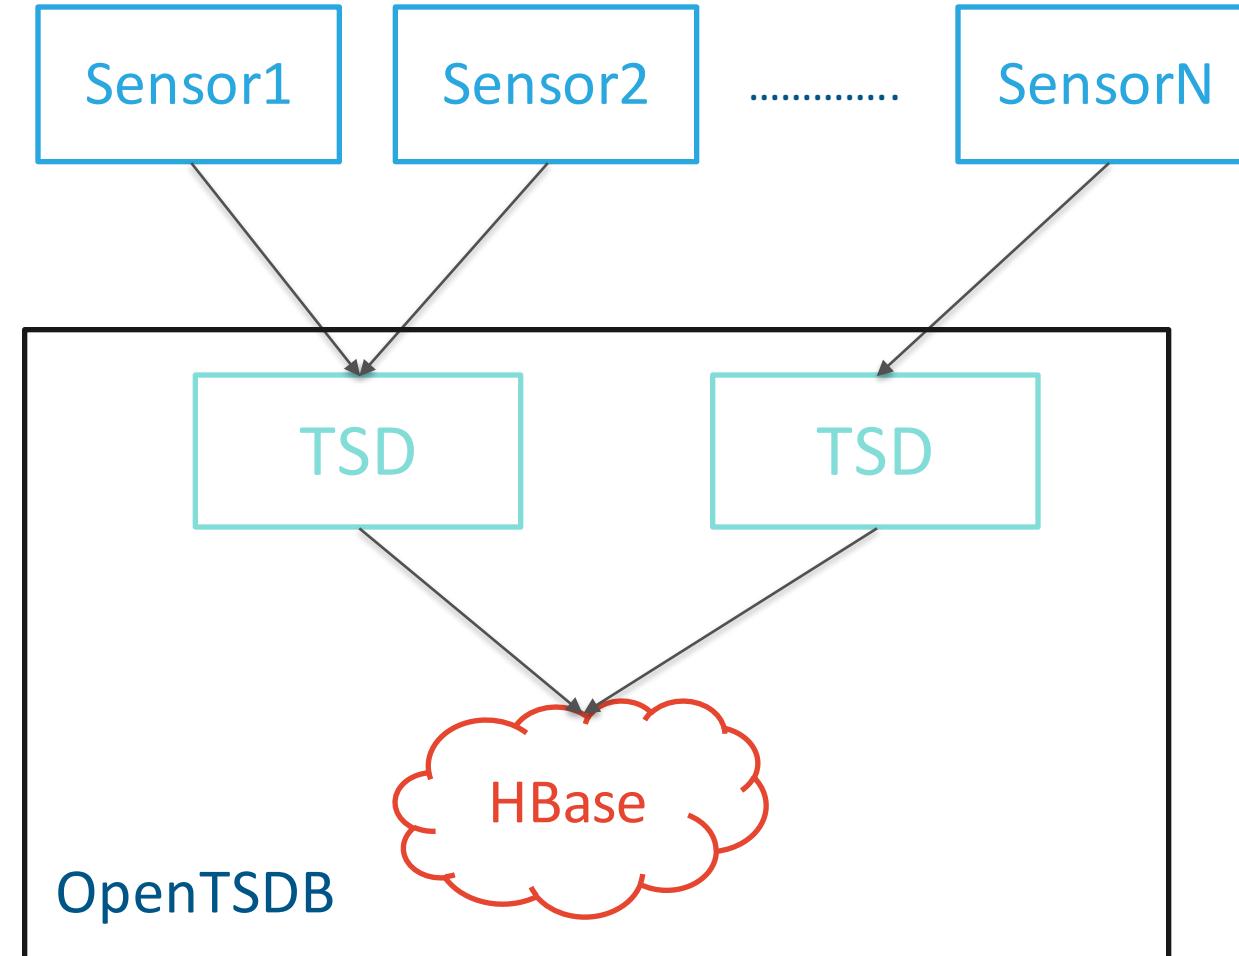


Image source: <http://opentsdb.net/overview.html>

OpenTSDB : Writing data

- Telnet
 - put <metric> <timestamp> <value> <tagk1=tagv1[tagk2=tagv2 ...tagkN=tagvN]>
 - Example: put num_requests_per_second 1439828251 50 host=web-server-1
- HTTP API
 - <host>:<port>/api/put
 - JSON objects containing data points
- Bulk Import
 - Using 'import' CLI utility

OpenTSDB : Reading data

- OpenTSDB GUI
 - Select metrics and tags to generate graphs
- HTTP API
 - <host>:<port>/api/query

OpenTSDB : Storing data – row key

- Row key is a concatenation of metric, timestamp and tags
 - `num_requests_per_second1439827200host=web-server-1`
- Since data is stored in sorted order, chunking happens in this order
 1. Metric
 - Enables fast scan of all time series for a metric
 2. Time
 - Normalized on 1 hour boundaries
 - All data points for an hour are stored in a single row
 3. Tags

OpenTSDB : Storing data – column

- Offset from timestamp in row key
- Example
 - num_requests_per_second 1439828251 50 host=web-server-1
 - num_requests_per_second 1439828251 72 host=web-server-2
 - num_requests_per_second 1439828252 30 host=web-server-3

Row key	Data:1051	Data:1052
num_requests_per_second1439827200host=web-server-1	50	
num_requests_per_second1439827200host=web-server-2	72	
num_requests_per_second1439827200host=web-server-3		30

OpenTSDB : GUI

Graph Stats Logs Version

From: 2014/09/28-10:00:00 To: (now) Autoreload WxH: 835x400

transfer_time +

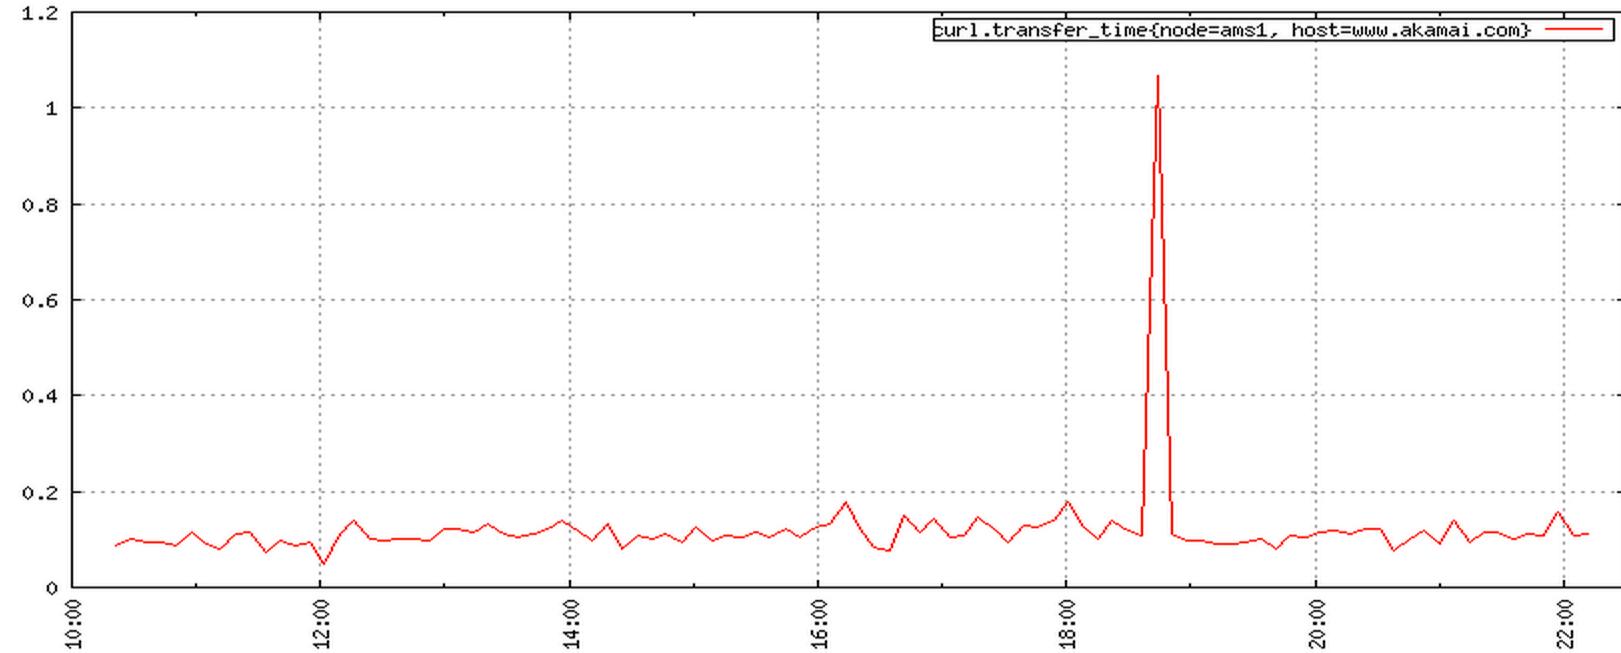
Metric: curl.transfer_time Rate Rate Ctr Right Axis
Tags node: ams1 Rate Ctr Max: Rate Ctr Reset:
host: www.akamai.com Aggregator: avg
 Downsample
avg 1m

Axes Key Style

Key location:

- Horizontal layout
- Box
- No key (overrides others)

5656 points retrieved, 655 points plotted in 46ms.





- High performance relational database layer over HBase for low-latency applications
- JDBC API

Phoenix : Use Case

Scalability of HBase

+

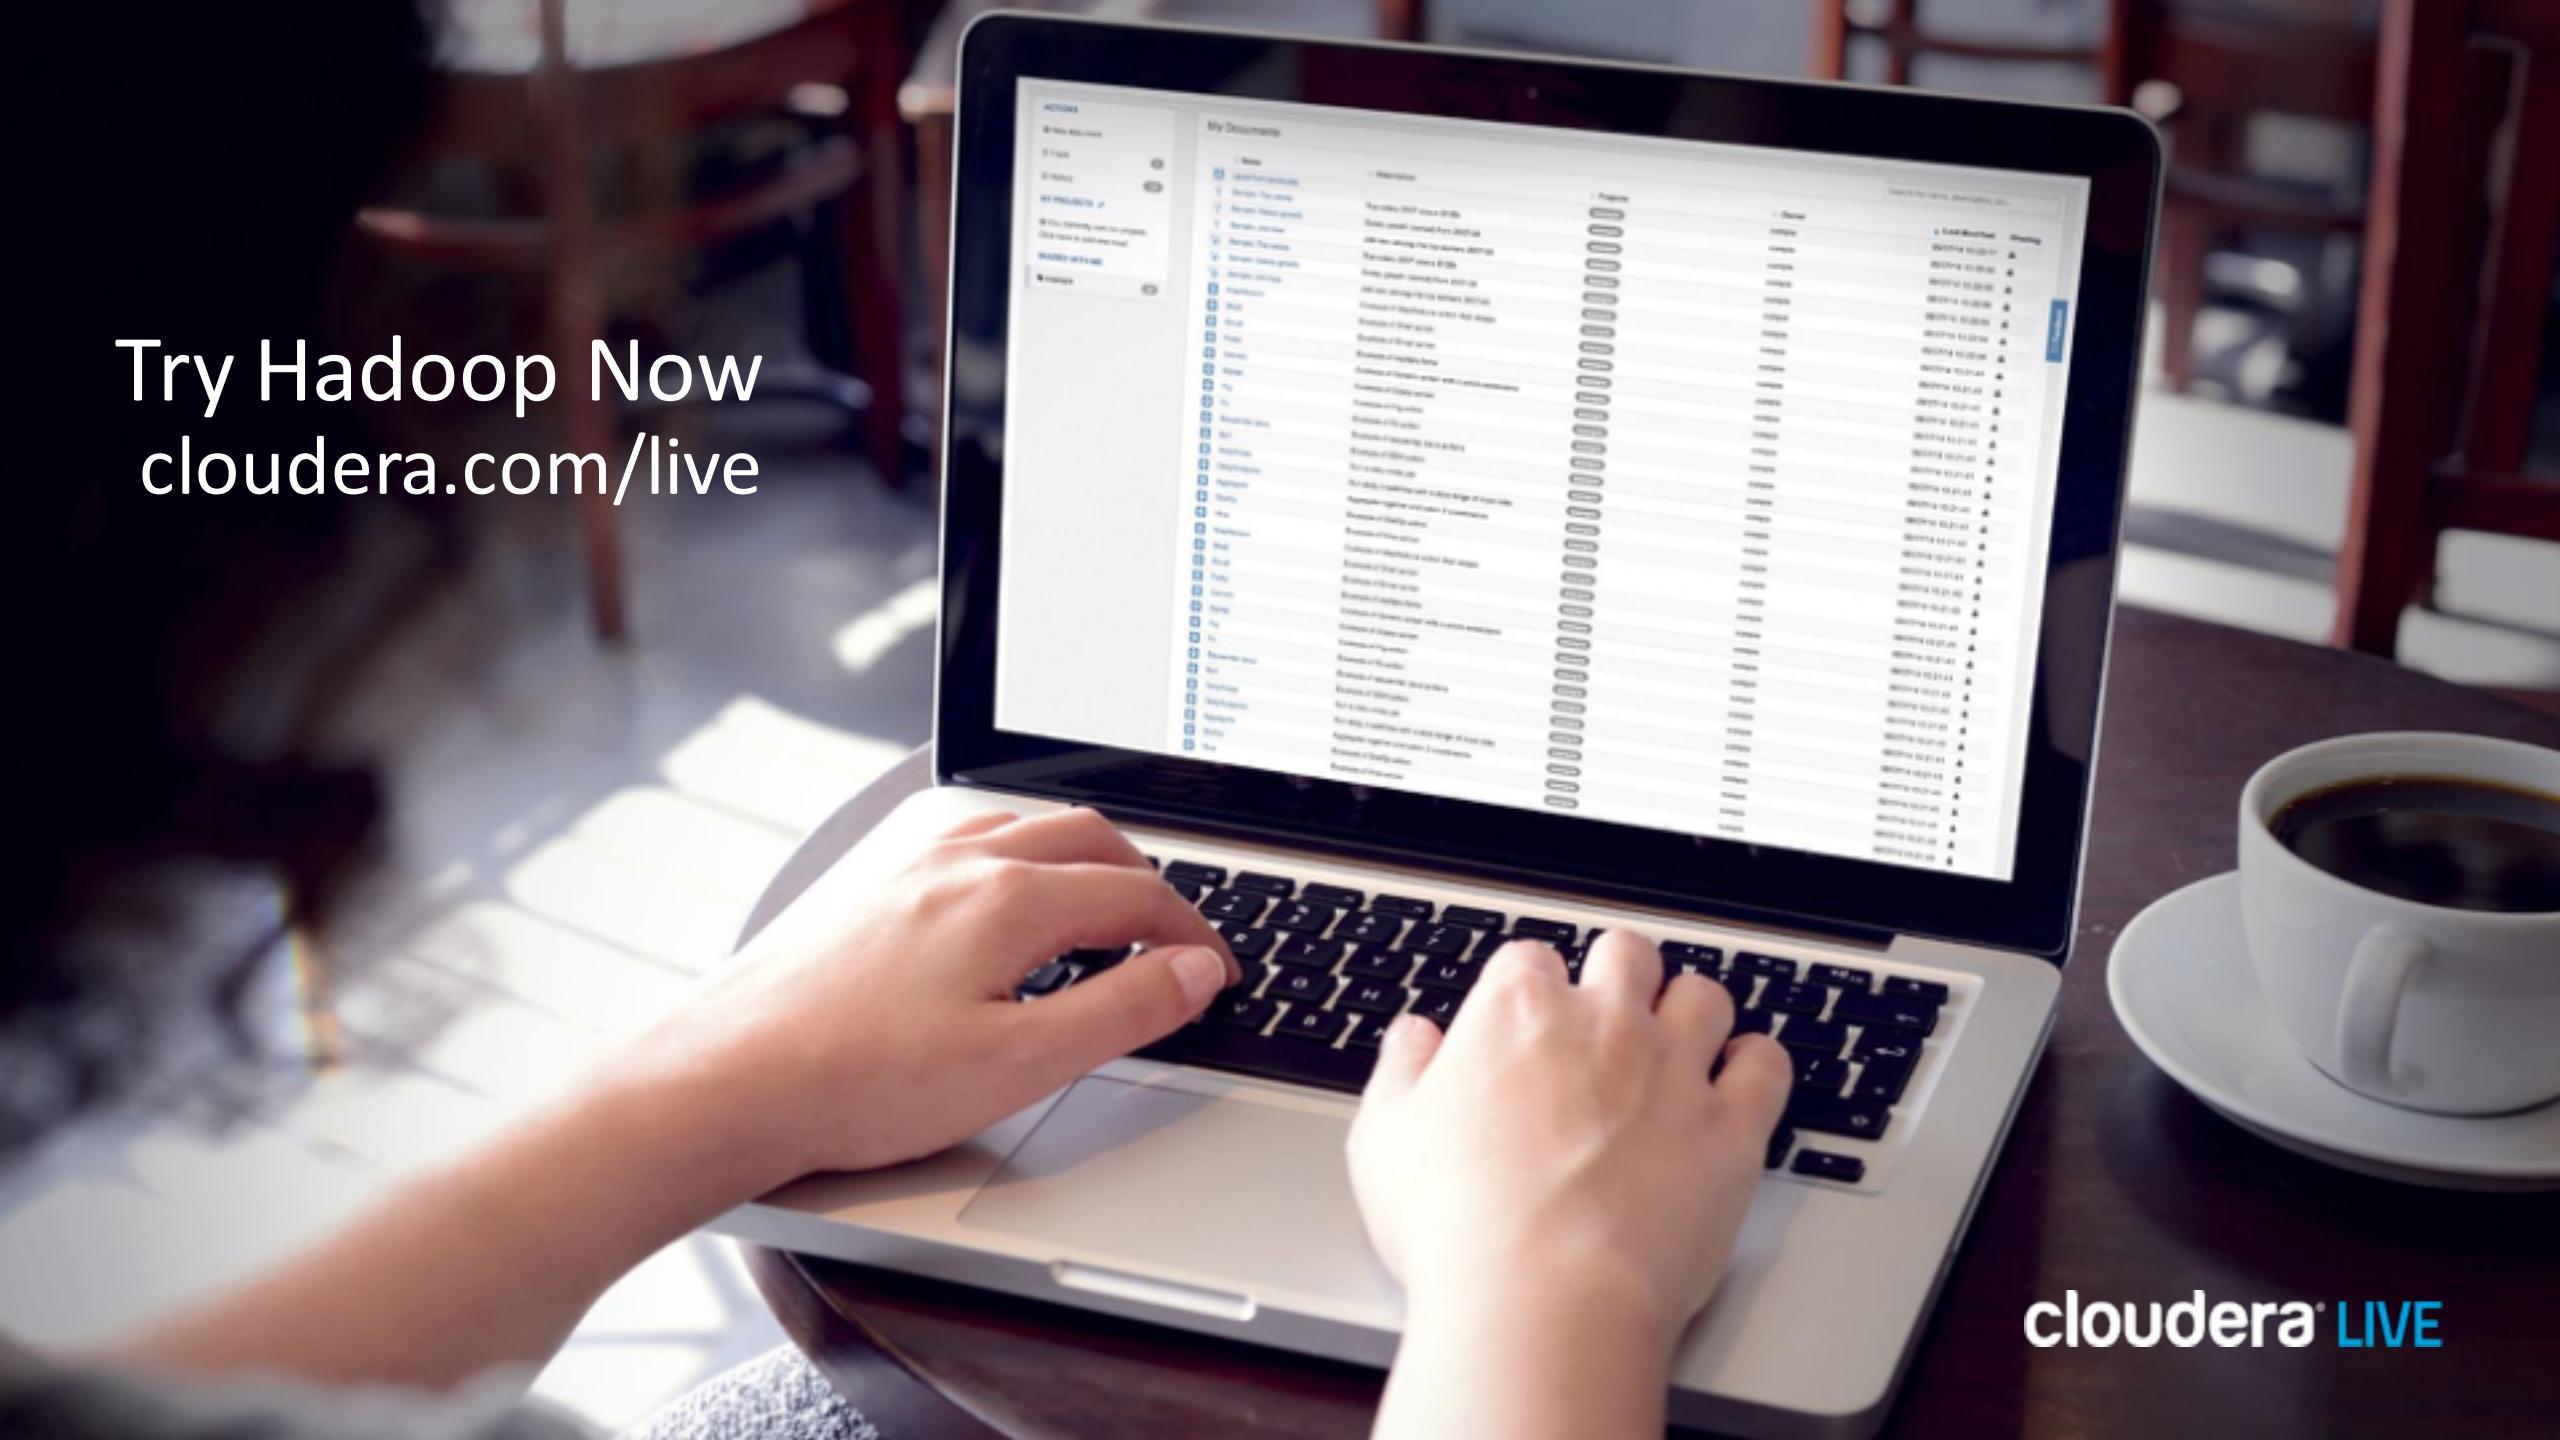
SQL interface access

Phoenix

- Provides typed access to data
- Provides secondary indexes
- Compiles SQL queries to native HBase scans
- Executes scans parallelly
- Directly uses HBase API, server-side hooks and custom filters
- Brings computation to the data
 - Pushes where clause to server-side filter
 - Executes aggregate queries using server-side hooks

That's it folks!



A person is sitting at a desk, working on a laptop. The laptop screen displays a web-based interface for managing a Hadoop file system. The interface shows a sidebar with 'Actions' (New Document, Edit, History, My Products, etc.) and a main area titled 'My Documents' with sections for 'Rows', 'Projects', and 'Tables'. The 'Rows' section lists numerous files and documents with names like 'Test table 1000 rows 100B', 'Data points recorded from 2007-08', 'Data points recorded from 2007-09', 'Predictions 2007 rows 100B', 'Data points recorded from 2008', 'Predictions 2008 rows 100B', and many others. A large table titled 'Data Points' is also visible on the right side of the screen.

Try Hadoop Now
cloudera.com/live

cloudera **LIVE**

Hello, Cloudera Customers and Users!

These community forums are intended for developers and admins using Cloudera's Apache Hadoop-based platform to build your applications. Welcome your suggestions and feedback [here](#).

Join this community to get a 40% discount for O'Reilly Media print books, and 50% for e-books and videos (bundles not included) -- as well as

To participate in upstream open source projects, use their respective upstream mailing lists.

Ask a Question

Type your question here...

Continue

Community

News (2 items)

Title	Posts
Community Guidelines & News Latest Post – This community is now mobile-friendly	5
Release Announcements Latest Post – Announcing: New Cloudera ODBC drivers for Impala a...	40

Get community help or provide feedback

cloudera.com/community

Sources

- A Survey of HBase Application Archetypes
 - Lars George, Jon Hsieh
 - <http://www.slideshare.net/HBaseCon/case-studies-session-7>
- OpenTSDB 2.0
 - Benoit Sigoure, Chris Larsen
 - <http://www.slideshare.net/HBaseCon/ecosystem-session-6>
- Hadoop and HBase: Motivations, Use cases and Trade-offs
 - Jon Hsieh
- Phoenix
 - <https://phoenix.apache.org>



cloudera
Questions ?